

Training Loss Convergence per Dataset (mean $\pm$ std, 3 seeds)

