

Supplementary Material:

FL-EHDS: A Privacy-Preserving Federated Learning Framework for the European Health Data Space

Fabio Liberti

Department of Computer Science
Universitas Mercatorum, Rome, Italy
fabio.liberti@studenti.unimercatorum.it
ORCID: 0000-0003-3019-5411

Abstract—This document provides supplementary material for the FL-EHDS paper, including complete algorithm pseudocode for all framework components, extended experimental figures, detailed algorithm comparison analysis with effect sizes (rank-biserial correlation), 10-seed statistical validation, advanced FL paradigm descriptions, infrastructure component specifications, extended EHDS interoperability details, clinical imaging experiment configurations with Chest X-ray baseline results, 10-seed Wilcoxon validation on imaging (90 experiments confirming Ditto’s superiority and HPFL’s modality-dependent reversal), Byzantine robustness on medical imaging (108 experiments with Krum defense), local epochs sweep, scalability analysis up to K=100 clients, noise-as-regularization analysis for small-cohort studies, Article 71 opt-out impact with policy recommendations, an extended threat model with attack taxonomy and defense mapping, a detailed framework positioning analysis, EHDS governance hypothesis testing (720 experiments across 3 hypotheses validating compound compliance, non-uniform opt-out fairness, and algorithm governance resilience on Heart Disease and Cardiovascular datasets), and thesis robustness validation (516 experiments confirming hyperparameter insensitivity, data efficiency thresholds, and personalization persistence under compound EHDS stress). The open-source reference implementation (~40K lines, 159 modules) is available at <https://github.com/FabioLiberti/FL-EHDS-FLICS2026>.

CONTENTS

Appendix A: Systematic Review and Datasets

A-A	PRISMA Flow Diagram	4
A-B	Dataset Landscape	4

Appendix B: Framework Architecture and Design

B-A	Algorithm Pseudocode	4
B-A.1	FedAvg with EHDS Compliance	4
B-A.2	Gaussian Differential Privacy Mechanism	4
B-A.3	HDAB Permit Validation	4
B-A.4	Secure Aggregation Protocol	5
B-A.5	FedProx for Non-IID Data	7
B-A.6	Article 71 Opt-Out Registry Protocol	7
B-A.7	FHIR R4 Preprocessing Pipeline	8

B-A.8	Privacy Budget Accountant	8
B-B	Detailed Architecture Description	9
B-B.1	Layer 1: Governance	9
B-B.2	Layer 2: FL Orchestration	10
B-B.3	Layer 3: Data Holders	10
B-B.4	Deployment Readiness Assessment	11
B-C	Infrastructure Components	12
B-C.1	Communication Layer	12
B-C.2	Serialization	12
B-C.3	Caching Layer	12
B-C.4	Orchestration	12
B-C.5	Monitoring	12
B-C.6	Model Watermarking	13
B-C.7	Cross-Silo Enhancements	13
B-D	Extended EHDS Interoperability	13
B-D.1	OMOP Common Data Model	13
B-D.2	IHE Integration Profiles	13
B-D.3	Cross-Border Data Exchange	13
B-D.4	Interoperability Architecture	13
B-E	Extended Threat Model and Security Analysis	14
B-E.1	Adversary Model	14
B-E.2	Protected Assets	15
B-E.3	Out-of-Scope Threats and Boundary Conditions	15
B-F	Framework Positioning: FL-EHDS vs. Existing Platforms	15
B-G	Cross-Border Governance Challenges: Extended Analysis	16
B-G.1	The Privacy Budget Conflict	16
B-G.2	Controller/Processor Ambiguity	17
B-G.3	Gradient Data Classification	17
B-G.4	Implications for FL-EHDS Design	17

Appendix C: Advanced FL Paradigms

C-A	Vertical Federated Learning	18
C-B	Byzantine-Resilient Aggregation	18

C-C	Continual Federated Learning	18	E-D.6	Statistical Significance	33
C-D	Multi-Task Federated Learning	18	E-D.7	Article 71 Opt-Out Impact	33
C-E	Hierarchical Federated Learning	18	E-D.8	Experimental Configuration Summary	34
C-F	Personalized Federated Learning	18	E-D.9	Deep MLP Algorithm Differentiation	35
C-G	Asynchronous Federated Learning	19	E-D.10	Local Epochs Sweep	36
C-H	Fairness-Aware Federated Learning	19	E-D.11	Scalability Analysis	36
Appendix D: Experimental Methodology		19	E-D.12	Confusion Matrix Analysis	38
D-A	Systematic Cascading Experiment Registry	19	E-D.13	RDP Composition Tightness	40
D-A.1	Cascading Methodology	20	E-D.14	Top- k Communication Efficiency	40
D-A.2	Campaign Summary	20	E-D.15	Differential Privacy Robustness at Scale	40
D-A.3	Clinical Dataset Portfolio	20	E-D.16	Combined Scalability and Differential Privacy	42
D-A.4	Cascade Descriptions and Key Findings	20	E-D.17	Byzantine Resilience under Differential Privacy	42
D-A.5	Standalone Validation Experiments	22	E-D.18	Extended Tabular Experiment Figures	42
D-A.6	Research Questions Index	23	E-E	EHDS-Specific Experiments	42
D-A.7	Reproducibility Infrastructure	23	E-E.1	Cross-Border Heterogeneous DP	42
Appendix E: Tabular Experiment Results		23	E-E.2	Dynamic Article 71 Opt-Out	42
E-A	Supplementary Experimental Figures	23	E-E.3	Membership Inference Attack (MIA)	49
E-A.1	Hospital Data Distribution	23	E-E.4	Partial Client Participation	49
E-A.2	Per-Client Training Time	23	E-E.5	Privacy Budget Exhaustion	49
E-A.3	Client Participation Matrix	23	E-F	Advanced FL Paradigm Experimental Validation	49
E-A.4	Gradient Norm Evolution	23	E-F.1	Fairness-Aware Aggregation	56
E-A.5	Communication Cost Analysis	23	E-F.2	Continual Federated Learning	56
E-A.6	Learning Rate Sensitivity	23	E-F.3	Extended Personalization Comparison	57
E-A.7	Batch Size Impact	23	E-F.4	Federated Unlearning (Right to Be Forgotten)	57
E-A.8	Per-Client Accuracy Trajectories	23	E-G	Client Selection and Gradient Inversion Privacy Validation	57
E-B	Extended Algorithm Comparison	23	E-G.1	Client Selection Strategy Comparison	57
E-B.1	Algorithms Evaluated	23	E-G.2	Gradient Inversion Privacy Analysis	57
E-B.2	Non-IID Configuration	23	E-H	Extended Federated Learning Validation	57
E-B.3	Convergence at Different Heterogeneity Levels	26	E-H.1	Dataset Coverage Expansion	57
E-B.4	Final Accuracy vs. Heterogeneity	26	E-H.2	Hierarchical FL Topologies	60
E-B.5	Complete 17-Algorithm Non-IID Evaluation (210 Experiments)	26	E-H.3	Asynchronous FL Strategies	60
E-B.6	Convergence Speed	30	E-H.4	Multi-Task FL	60
E-B.7	Evidence-Based Algorithm Recommendation Matrix	30	E-H.5	Extended Byzantine Robustness	60
E-C	Extended Results from Main Paper	30	E-H.6	Shapley Values and Client Valuation	60
E-C.1	Heart Disease and Diabetes Algorithm Comparison	30	E-H.7	Cascade 7: Trustworthiness Detailed Results	66
E-C.2	Per-Hospital Heterogeneity Analysis	30	E-H.8	Cascade 10: Detailed Results	66
E-D	Extended Tabular Experiment Results	30	E-H.9	Main Paper Extended Content (v3)	71
E-D.1	Heterogeneity Impact	30			
E-D.2	Client Scaling	30			
E-D.3	Learning Rate Sensitivity	30			
E-D.4	Privacy-Utility Tradeoff (Differential Privacy)	30			
E-D.5	Fairness Analysis	33			

	E-H.10	Extended Standalone Experiment Results	74
E-I		Computational Cost	76
E-J		Thesis Robustness Validation	76
	E-J.1	Phase A: Ditto Hyperparameter Sensitivity	77
	E-J.2	Phase B: Personalization vs. Data Fraction	77
	E-J.3	Phase C: Compound EHDS Stress Frontier	78
	E-J.4	Summary of Thesis Robustness Claims	78
Appendix F: Medical Imaging Results			78
F-A		Clinical Imaging: Extended Details . . .	78
	F-A.1	Datasets	78
	F-A.2	Model Architectures	79
	F-A.3	Experimental Configuration .	79
	F-A.4	Reproducibility	79
	F-A.5	Chest X-ray Results	79
F-B		Brain Tumor and Skin Cancer Results .	80
F-C		Statistical Significance and Effect Sizes	80
F-D		10-Seed Wilcoxon Signed-Rank Tests on Imaging	80
F-E		Extended Algorithm Comparison on Imaging (5 Seeds)	81
F-F		Imaging Convergence Analysis	82
F-G		Per-Client Fairness on Imaging	82
F-H		Diagnostic Equity Analysis	83
	F-H.1	Definition	83
	F-H.2	DEI Results	84
	F-H.3	Key Findings	84
	F-H.4	EHDS Policy Implications .	84
F-I		Byzantine Robustness on Medical Imaging	84
Appendix G: EHDS Governance Hypothesis Testing			86
G-A		Extended Governance Validation (E1–E4)	86
G-B		Hypothesis H1: Compound EHDS Compliance	86
G-C		Hypothesis H2: Non-Uniform Article 71 Opt-Out Fairness	86
G-D		Hypothesis H3: Algorithm Governance Resilience	87
G-E		Cross-Dataset Comparison and Small-Sample Analysis	88
G-F		Governance Validation Summary	88
References			89

APPENDIX A SYSTEMATIC REVIEW AND DATASETS

A. PRISMA Flow Diagram

The systematic review underlying FL-EHDS follows the PRISMA 2020 guidelines [?]. Five academic databases (PubMed, IEEE Xplore, Scopus, Web of Science, arXiv) were searched using Boolean queries combining federated learning, health data spaces, and EHDS-specific terminology, covering publications from January 2022 to March 2026. The search yielded 847 initial records; after automated and manual deduplication (264 duplicates removed), 583 unique records underwent title-and-abstract screening against pre-defined inclusion criteria: (i) primary focus on federated learning applied to health data; (ii) relevance to European regulatory frameworks or cross-border clinical data sharing; (iii) peer-reviewed publication or recognized institutional origin. A total of 47 studies met all criteria and were included in the qualitative synthesis. Supplementary searches of institutional websites (n=15), regulatory bodies and standardization organizations (n=8), and backward/forward citation tracking (n=2) were conducted but did not contribute additional studies beyond those already identified. Figure S-1 presents the complete PRISMA flow diagram.

B. Dataset Landscape

The FL-EHDS framework supports 19 healthcare datasets spanning four modalities. Table S-I provides a comprehensive overview. This diversity enables evaluation across multiple EHDS-relevant dimensions: data scale (120–253K samples), feature dimensionality (9–30 tabular, high-dimensional imaging), task complexity (binary to 5-class), partition strategies (natural hospital-based and synthetic Dirichlet), and interoperability standards (CSV, FHIR R4, OMOP-CDM). Experimentally evaluated datasets in the main paper are marked with $\checkmark$.

APPENDIX B FRAMEWORK ARCHITECTURE AND DESIGN

A. Algorithm Pseudocode

This section provides formal algorithmic descriptions of all FL-EHDS framework components. Each algorithm is presented with: (1) a contextual explanation of *why* the component is needed in the EHDS regulatory context; (2) the formal pseudocode; and (3) practical considerations for deployment. The algorithms are organized following the data flow through the three-layer architecture: governance validation (Layer 1), privacy-preserving aggregation (Layer 2), and local data processing (Layer 3).

Reading guide: Algorithms S1–S4 form the core FL-EHDS training pipeline. Algorithms S5–S6 address EHDS-specific challenges (non-IID data and citizen opt-out). Algorithms S7–S8 handle data preprocessing and privacy budget management.

1) *FedAvg with EHDS Compliance:* Algorithm S1 presents the core federated averaging procedure adapted for EHDS regulatory requirements, operating in a client-server architecture where the central aggregator coordinates training across

distributed hospital nodes within a Secure Processing Environment.

Key Design Decisions:

- **ValidatePermit:** Before each round, the HDAB-issued permit is verified against temporal bounds and Article 53 permitted purposes.
- **SelectParticipants:** Configurable client selection—full participation or sampling for large federations.
- **FilterOptedOut:** Records from citizens who exercised Article 71 opt-out rights are excluded *before* gradient computation.
- **Weighted Aggregation:** Gradients weighted by local dataset size (n_h), following original FedAvg [13].
- **ClipGradient:** L2-norm clipping bounds individual contributions, providing sensitivity bounds for DP.

Relationship to subsequent components: The ClipGradient operation in Algorithm S1 establishes a bounded sensitivity C for each client’s contribution. This bound is the prerequisite for Algorithm S2 (Gaussian DP), which calibrates noise proportional to C . Meanwhile, ValidatePermit invokes Algorithm S3 (Permit Validation) and FilterOptedOut invokes Algorithm S6 (Opt-Out Filtering).

2) *Gaussian Differential Privacy Mechanism:* Algorithm S2 implements the Gaussian mechanism for differential privacy, applied at the aggregation server after receiving clipped gradients.

Mathematical Foundation: The noise scale $\sigma = C \cdot \sqrt{2 \ln(1.25/\delta)}/\epsilon$ guarantees (ϵ, δ) -DP. The cumulative privacy expenditure is tracked using Rényi DP (RDP) [26] composition, providing 5–6 $\times$ tighter bounds than naive composition.


Practical Considerations:

- $\epsilon = 10$: moderate noise, <2pp accuracy drop on PTB-XL and Cardiovascular (see Table S-XIV in Section E-D)
- $\epsilon = 1$: strong privacy; personalized methods (Ditto, HPFL) retain 87–89% on PTB-XL while FedAvg collapses to 52%
- The ϵ selection must be negotiated with HDABs during permit approval

Integration with privacy budget: The ϵ consumed by Algorithm S2 in each round is tracked by Algorithm S8 (RDP Privacy Budget Accountant). If the cumulative budget exceeds the threshold approved in the HDAB data permit, training is automatically terminated. The next algorithm (S3) formalizes the permit validation that authorizes each round.

3) *HDAB Permit Validation:* Algorithm S3 ensures all FL operations comply with the data permit issued by HDABs. Under EHDS Article 53, secondary use of health data is only lawful for specifically enumerated purposes (scientific research, public health surveillance, AI training for health). The permit validation module is invoked at the beginning of every FL round—not just at training start—to guarantee continuous compliance even if a permit is revoked mid-study or its temporal validity expires. Each validation event

PRISMA 2020 flow diagram for new systematic reviews which included searches of databases, registers and other sources



*Consider, if feasible to do so, reporting the number of records identified from each database or register searched (rather than the total number across all databases/registers).
**If automation tools were used, indicate how many records were excluded by a human and how many were excluded by automation tools.

Source: Page MJ, et al. BMJ 2021;372:n71. doi: 10.1136/bmj.n71.

This work is licensed under CC BY 4.0. To view a copy of this license, visit <https://creativecommons.org/licenses/by/4.0/>

Fig. S-1. PRISMA 2020 flow diagram for the systematic review. Database searches across PubMed, IEEE Xplore, Scopus, Web of Science, and arXiv identified 847 records; after deduplication (264 removed) and screening, 47 studies met inclusion criteria (2022–2026, FL/EHDS focus, peer-reviewed or recognized institutional origin). Additional records from institutional websites (n=15), organisations (n=8), and citation searching (n=2) were assessed but did not contribute to the final inclusion set. Adapted from Page et al. (BMJ 2021;372:n71), licensed under CC BY 4.0.

is logged as a GDPR Article 30 processing record, creating an immutable audit trail that regulators can inspect.

EHDS Governance Role: This algorithm is the enforcement point between the Governance Layer (Layer 1) and the FL Orchestration Layer (Layer 2). Without it, a permit expiring at round 15 of a 20-round study would allow unauthorized data processing for rounds 16–20.

Failure modes: Each exception type triggers a different response: `PermitExpiredError` terminates the entire study; `PurposeMismatchError` indicates a configuration error requiring researcher intervention; `UnauthorizedCategoryError` may allow continued training on the authorized subset of categories. All failure events are logged for regulatory audit. Once a round passes permit validation, the aggregation server collects client gradients using the secure protocol described next.

4) *Secure Aggregation Protocol:* Even though FL prevents raw data sharing, the gradient updates themselves can leak patient information: Zhu et al. [21] demonstrate that gradients can be inverted to reconstruct training images. In the EHDS

context, where gradients encode patterns from sensitive health records across 27 Member States, this is an unacceptable privacy risk. Secure aggregation addresses this by ensuring the SPE aggregation server can compute $\sum_k \Delta_k$ without ever observing any individual hospital’s gradient Δ_k .

Algorithm S4 implements this using Shamir’s secret sharing and pairwise masking, ensuring the server observes only the aggregate gradient.

Protocol Phases: (1) Each client splits gradients into K shares using (t, K) -threshold Shamir secret sharing; (2) Clients add pairwise random masks negotiated via ECDH key exchange; (3) The server computes the sum—masks cancel out and only the true aggregate remains.

Defense-in-depth: Secure aggregation (Algorithm S4) combined with differential privacy (Algorithm S2) provides layered protection: even if the aggregation server is compromised, it learns only the noisy aggregate—never individual hospital contributions. The combination addresses the unresolved GDPR question of whether model gradients constitute “per-

TABLE S-I
FL-EHDS DATASET LANDSCAPE: COMPLETE FRAMEWORK COVERAGE

Dataset	Type	Samples	Feat.	Classes	FL Partition	EHDS Category (Art. 33)	EHDS Level	Eval.
<i>A. Tabular — Clinical EHR</i>								
Diabetes 130-US	Tabular	101,766	22	2	Dirichlet ($\alpha=0.5$)	EHR, ICD-9, medications	L2: FHIR-mappable	✓
Heart Disease UCI	Tabular	920	13	2	Natural (4 hospitals)	Vitals, ECG, lab results	L2: FHIR-mappable	✓
PTB-XL ECG [†]	Tabular	21,799	9	5	Natural (52 sites)	SCP-ECG (EN 1064), diagnostics	L2: FHIR-mappable	✓
Cardiovascular Disease	Tabular	70,000	11	2	Dirichlet ($\alpha=0.5$)	Vitals, lab, risk factors	L2: FHIR-mappable	✓
Breast Cancer Wisconsin	Tabular	569	30	2	Dirichlet ($\alpha=0.5$)	Pathology (FNA cytology)	L2: FHIR-mappable	✓
Stroke Prediction	Tabular	5,110	10	2	Dirichlet	Cardiovascular risk factors	L2: FHIR-mappable	—
CDC Diabetes BRFSS	Tabular	253,680	21	2	Dirichlet	Population health survey	L2: FHIR-mappable	—
CKD UCI	Tabular	400	24	2	Dirichlet	Renal panel, comorbidities	L2: FHIR-mappable	—
Cirrhosis Mayo	Tabular	418	18	2	Dirichlet	Hepatology, drug trial	L2: FHIR-mappable	—
<i>B. Tabular — FHIR-Native</i>								
Synthea FHIR R4	FHIR	1,180	14	2	Hospital profile	Patient, Condition, Encounter	L1: FHIR-native	qual.
SMART Bulk FHIR	FHIR	120	12	2	Single export	NDJSON Bulk Data (Art. 46)	L1: FHIR-native	qual.
<i>C. Generated Pipelines (in-memory)</i>								
FHIR R4 Synthetic	Gen.	config.	10	2	Hospital profile	Generated FHIR bundles	L1: FHIR-native	qual.
OMOP-CDM Harmonized	Gen.	config.	36	2	Cross-border	Vocabulary harmonization	L3: OMOP	qual.
<i>D. Medical Imaging</i>								
Chest X-ray	Imaging	5,856	—	2	Dirichlet ($\alpha=0.5$)	Radiology (DICOM)	L4: Imaging	✓
Brain Tumor MRI	Imaging	7,023	—	4	Dirichlet ($\alpha=0.5$)	Neuro-imaging (DICOM)	L4: Imaging	✓
Skin Cancer	Imaging	3,297	—	2	Dirichlet ($\alpha=0.5$)	Dermatology (DICOM)	L4: Imaging	✓
Diabetic Retinopathy	Imaging	35,126	—	5	Dirichlet	Ophthalmology (DICOM)	L4: Imaging	—
Brain Tumor MRI (alt.)	Imaging	3,264	—	4	Dirichlet	Neuro-imaging (DICOM)	L4: Imaging	—
ISIC Skin Lesions	Imaging	2,357	—	9	Dirichlet	Dermatology (DICOM)	L4: Imaging	—

EHDS Levels: L1 = FHIR-native (Art. 46 compliant); L2 = FHIR-mappable (standard clinical features with FHIR mapping in metadata); L3 = OMOP-CDM harmonized (cross-border vocabulary alignment, Art. 50); L4 = Medical imaging (DICOM, Art. 33 “medical images”).

[†]PTB-XL: European-origin dataset (PTB, Berlin, Germany) with SCP-ECG coding (EN 1064). 52 recording sites enable natural hospital-based FL partitioning—the strongest EHDS benchmark in the framework.

Eval.: ✓ = quantitative experimental evaluation (P1.2); qual. = qualitative pipeline validation; — = supported but not evaluated in current paper. **config.** = sample count depends on generation parameters.

Algorithm S1: FL-EHDS FedAvg Training

Input: Hospitals $\mathcal{H} = \{h_1, \dots, h_K\}$, permit P , rounds T

Output: Global model $\theta^{(T)}$

Server executes:

Initialize $\theta^{(0)}$

for round $t = 1$ to T **do**

 // Governance check (Layer 1)

if not `ValidatePermit( $P, t$ )` **then abort**

$\mathcal{H}_t \leftarrow \text{SelectParticipants}(\mathcal{H})$

for each hospital $h \in \mathcal{H}_t$ **in parallel do**

$\Delta_h^{(t)}, n_h \leftarrow \text{LocalTrain}(h, \theta^{(t-1)})$

 // Aggregation with privacy (Layer 2)

$\theta^{(t)} \leftarrow \theta^{(t-1)} + \frac{1}{\sum_h n_h} \sum_{h \in \mathcal{H}_t} n_h \cdot \Delta_h^{(t)}$

`LogCompliance( $t, \mathcal{H}_t$ )`

return $\theta^{(T)}$

LocalTrain(h, θ) at hospital h :

 // Opt-out filtering (Article 71)

$\mathcal{D}_h \leftarrow \text{FilterOptedOut}(\mathcal{D}_h, \text{OptOutRegistry})$

$\theta_h \leftarrow \theta$

for epoch $e = 1$ to E **do**

for batch $\mathcal{B} \in \mathcal{D}_h$ **do**

$\theta_h \leftarrow \theta_h - \eta \nabla \mathcal{L}(\theta_h; \mathcal{B})$

$\Delta_h \leftarrow \theta_h - \theta$

 // Privacy protection (Layer 3)

$\Delta_h \leftarrow \text{ClipGradient}(\Delta_h, C)$

return $\Delta_h, |\mathcal{D}_h|$

Algorithm S2: Gaussian DP Mechanism**Input:** Gradient Δ , sensitivity C , privacy budget ϵ , δ **Output:** Noisy gradient $\hat{\Delta}$

```

// Compute noise scale from Gaussian mechanism
 $\sigma \leftarrow C \cdot \sqrt{2 \ln(1.25/\delta)}/\epsilon$ 
// Add calibrated Gaussian noise to each parameter
for each parameter  $w \in \Delta$  do
   $\tilde{w} \leftarrow w + \mathcal{N}(0, \sigma^2)$ 
// Track cumulative privacy expenditure
PrivacyAccountant.spend( $\epsilon$ )
if PrivacyAccountant.budget_exhausted() then
  raise PrivacyBudgetExhaustedError
return  $\hat{\Delta}$ 

```

Algorithm S3: Data Permit Validation**Input:** Permit P , round t , requested categories $\mathcal{C}$ **Output:** Boolean validity

```

// Check temporal validity
if CurrentTime() >  $P$ .valid_until then
  raise PermitExpiredError
// Check purpose alignment (Article 53)
if  $P$ .purpose  $\notin$  AllowedPurposes then
  raise PurposeMismatchError
// Check data category authorization
for each category  $c \in \mathcal{C}$  do
  if  $c \notin P$ .authorized_categories then
    raise UnauthorizedCategoryError
// Log access for GDPR Article 30
AuditTrail.log(permit= $P$ , round= $t$ , categories= $\mathcal{C}$ )
return True

```

sonal data”: with both mechanisms active, the information available to any single party is provably bounded. The following algorithms address how local training handles EHDS-specific data challenges.

5) *FedProx for Non-IID Data*: Algorithm S5 extends FedAvg with a proximal term that penalizes local model divergence from the global model [14]. In EHDS cross-border federations, data heterogeneity is a structural feature, not an exception: hospitals in different Member States serve distinct demographics, follow national clinical guidelines, and use different diagnostic thresholds. For instance, heart disease prevalence ranges from 39.2% in Rome to 62.6% in Amsterdam in our experimental setting. Without drift control, local models can diverge so far from the global consensus that aggregation produces a deteriorated global model. The proximal term $\frac{\mu}{2} \|\theta_h - \theta\|^2$ acts as a regularizer that keeps each hospital’s local update within a controlled distance of the global model, balancing personalization with collaboration.

When to use in EHDS: Recommended for federations with moderate non-IID conditions and when client dropout is expected (hospitals may temporarily disconnect). FedProx tolerates partial participation better than FedAvg because the proximal term stabilizes local updates even with fewer training epochs.

Parameter Selection: $\mu = 0$ reduces to FedAvg; $\mu \in [0.01, 0.1]$ provides stable convergence; $\mu > 1$ may prevent

local adaptation. The choice of μ should be documented in the data permit application so that the HDAB can assess the expected privacy-utility trade-off.

Before any local training begins (whether with FedAvg, FedProx, or any other algorithm), the framework must enforce citizen opt-out rights. The following algorithm ensures this compliance.

6) *Article 71 Opt-Out Registry Protocol*: Algorithm S6 implements the citizen opt-out mechanism mandated by EHDS Article 71. This article grants every EU citizen the right to object to secondary use of their electronic health data—a fundamental right that must be enforced *before* any gradient computation occurs. The algorithm queries the national opt-out registry maintained by each Member State and removes matching records from the local training dataset.

Granularity levels: (1) *Blanket opt-out*—citizen refuses all secondary use; (2) *Purpose-specific*—e.g., permitting scientific research but blocking commercial analytics; (3) *Category-specific*—e.g., allowing demographics but blocking genomic data. This granularity reflects the EHDS principle that citizens should have meaningful control, not merely a binary yes/no choice.

EHDS Governance Role: Opt-out filtering operates at Layer 3 (Data Holders) before local training. Registry lookups use LRU caching with configurable TTL to minimize latency (<10ms per round) while ensuring timely propagation of new

Algorithm S4: Secure Aggregation (Pairwise Masking)**Input:** Client gradients $\{\Delta_1, \dots, \Delta_K\}$, threshold t **Output:** Aggregated gradient Δ_{agg}

// Phase 1: ECDH key exchange + Shamir sharing

for each client k **do** $\text{shares}_k \leftarrow \text{ShamirShare}(\Delta_k, t, K)$ Distribute shares_k to other clients

// Phase 2: Add pairwise random masks

for each client k **do** $\hat{\Delta}_k \leftarrow \Delta_k + \sum_{j < k} r_{jk} - \sum_{j > k} r_{kj}$

// Phase 3: Server reconstructs aggregate

 $\Delta_{agg} \leftarrow \sum_{k=1}^K \hat{\Delta}_k$ // Masks cancel: $\sum_k \sum_{j < k} r_{jk} - \sum_k \sum_{j > k} r_{kj} = 0$ **if** ActiveClients $< t$ **then** **raise** SecureAggregationError**return** Δ_{agg} **Algorithm S5: FedProx Local Update****Input:** Local data $\mathcal{D}_h$, global model θ , proximal weight μ **Output:** Local update Δ_h $\theta_h \leftarrow \theta$ **for** epoch $e = 1$ to E **do** **for** batch $\mathcal{B} \in \mathcal{D}_h$ **do** $g \leftarrow \nabla \mathcal{L}(\theta_h; \mathcal{B})$ // Proximal term: $\nabla \frac{\mu}{2} \|\theta_h - \theta\|^2$ $g \leftarrow g + \mu(\theta_h - \theta)$ $\theta_h \leftarrow \theta_h - \eta \cdot g$ $\Delta_h \leftarrow \theta_h - \theta$ **return** Δ_h

opt-out decisions. All filtering statistics are logged for GDPR Article 30 audit compliance.

Impact on model quality: High opt-out rates reduce training data volume, potentially degrading model performance—particularly for underrepresented subpopulations. The audit log captures filtering statistics to quantify this impact and support transparency reporting. Once opted-out records are excluded, the remaining data must be harmonized into a consistent format before local model training can proceed.

7) *FHIR R4 Preprocessing Pipeline:* Algorithm S7 standardizes heterogeneous EHR data into FHIR R4 format for ML consumption. This preprocessing step is critical in the EHDS context because only 34% of European healthcare providers currently achieve full FHIR R4 compliance [7]. The remaining 66% use legacy formats (HL7v2, CDA, proprietary CSV exports) that must be harmonized before FL training can proceed on a consistent feature space.

Four-stage pipeline: (1) *Format detection* automatically identifies the source format; (2) *Terminology mapping* converts local codes to international standards (ICD-10 for diagnoses, ATC for medications, LOINC for laboratory results); (3) *FHIR transformation* produces validated FHIR R4 bundles using the six Article 33 data categories (Patient Summary, E-Prescription, Laboratory Results, Medical Imaging, Hospital Discharge, Rare Disease); (4) *Tensor extraction* converts structured FHIR resources into numerical tensors ready for model training.

EHDS Relevance: Without this harmonization step, hospitals in different Member States would produce incompatible feature spaces, making federated aggregation meaningless. The pipeline ensures that a gradient computed in a Finnish hospital is semantically compatible with one from an Italian hospital.

Validation requirements: The FHIR validation step rejects records with missing mandatory fields or invalid terminology codes, ensuring data quality before model training. Rejected records are logged (without patient-identifiable content) for audit purposes. With harmonized data ready for training, the final core component manages the overall privacy budget across the entire study.

8) *Privacy Budget Accountant:* Algorithm S8 tracks cumulative privacy expenditure across FL rounds using Rényi Differential Privacy (RDP) moment accounting [26]. In the EHDS governance model, the total privacy budget ε_{total} is a parameter of the data permit: the researcher specifies the desired budget in the permit application, and the HDAB evaluates whether the proposed budget provides sufficient privacy protection for the requested data categories and population size.

Why RDP accounting: Naive DP composition (adding ε per round) yields loose bounds: 20 rounds at $\varepsilon=0.5$ each would consume $\varepsilon=10$ total. RDP provides 5–6 $\times$ tighter bounds [26], [36], meaning the same 20 rounds can achieve the same privacy guarantee with significantly less noise—and therefore better model utility.

Algorithm S6: Article 71 Opt-Out Filtering
Input: Local dataset $\mathcal{D}_h$, purpose p , categories $\mathcal{C}$
Output: Filtered dataset $\mathcal{D}'_h$
// Synchronize with national opt-out registry
 $\text{OptOutRecords} \leftarrow \text{FetchOptOutRegistry}(\text{MemberState})$
 $\mathcal{D}'_h \leftarrow \emptyset$
for each record $r \in \mathcal{D}_h$ **do**
 $\text{citizen_id} \leftarrow r.\text{pseudonymized_id}$
 $\text{opted_out} \leftarrow \text{False}$
 // Check purpose-specific opt-out
 if $(\text{citizen_id}, p) \in \text{OptOutRecords}$ **then**
 $\text{opted_out} \leftarrow \text{True}$
 // Check category-specific opt-out
 for each $c \in \mathcal{C}$ **do**
 if $(\text{citizen_id}, c) \in \text{OptOutRecords}$ **then**
 $\text{opted_out} \leftarrow \text{True}$
 if not opted_out **then**
 $\mathcal{D}'_h \leftarrow \mathcal{D}'_h \cup \{r\}$
// Log filtering statistics for audit
 $\text{AuditLog.record}(\text{total}=|\mathcal{D}_h|, \text{filtered}=|\mathcal{D}'_h|)$
return $\mathcal{D}'_h$

Algorithm S7: FHIR R4 Preprocessing
Input: Raw EHR records $\mathcal{R}$, feature specification $\mathcal{F}$
Output: Training tensors (X, y)
 $\text{format} \leftarrow \text{DetectFormat}(\mathcal{R})$ *// HL7v2, CDA, CSV*
 $\text{parser} \leftarrow \text{GetParser}(\text{format})$
 $\text{records} \leftarrow \text{parser.parse}(\mathcal{R})$
// Map to standard terminologies
for each $r \in \text{records}$ **do**
 $r.\text{diagnoses} \leftarrow \text{MapToICD10}(r.\text{diagnoses})$
 $r.\text{medications} \leftarrow \text{MapToATC}(r.\text{medications})$
 $r.\text{labs} \leftarrow \text{MapToLOINC}(r.\text{labs})$
 $\text{fhir_bundle} \leftarrow \text{ToFHIR}(\text{records})$
 $\text{ValidateFHIR}(\text{fhir_bundle})$
 $X \leftarrow \text{ExtractFeatures}(\text{fhir_bundle}, \mathcal{F})$
 $X \leftarrow \text{StandardScaler.fit_transform}(X)$
 $y \leftarrow \text{ExtractLabels}(\text{fhir_bundle})$
return (X, y)

Hard budget enforcement: When the cumulative expenditure approaches ε_{total} , the accountant raises a `BudgetExhaustedError` that terminates training. This prevents “privacy bankruptcy”—a situation where continued training would violate the privacy guarantee approved in the data permit. The per-round allocation strategy distributes remaining budget uniformly across remaining rounds, adapting dynamically if training converges faster than expected.

B. Detailed Architecture Description

This section provides a comprehensive technical specification of the FL-EHDS three-layer architecture, detailing all modules, services, protocols, and standards implemented in the reference framework. Figure 2 in the main paper provides the high-level view; Figure S-2 presents the detailed component-level diagram; the description below enumerates every component with its specific technical parameters.

1) *Layer 1: Governance:* Four principal modules comprise the governance layer:

1.1 HDAB Integration (per Member State). Each Health Data Access Body instance implements: OAuth2/mTLS authentication with bearer token management (refresh tokens, scopes: `permits:read`, `permits:write`); a permit store with SQLite persistence backend; configurable HDAB strictness level (scale 1–5, where DE=5, ES=2); and national privacy constraints including per-jurisdiction differential privacy bounds ($\varepsilon_{\max}$: DE=1.0, FR=3.0, IT=5.0).

1.2 Data Permit Manager (Article 53). Manages the complete permit lifecycle: PENDING → APPROVED → ACTIVE → EXPIRED/REVOKED. Validates permitted purposes (scientific research, public health surveillance, AI system development, personalized medicine, official statistics) against authorized data categories (EHR, lab results, imaging, genomic, registry, ECG, pathology). Implements expiry verification and strict mode enforcement for continuous compliance.

1.3 Opt-Out Registry (Article 71). Supports three granularity levels: record-level, patient-level, and dataset-level opt-out.

Algorithm S8: RDP Privacy Budget Accountant
Input: Total budget $(\epsilon_{total}, \delta_{total})$, rounds T
Output: Per-round budget allocation
 $\lambda \leftarrow [0] \times \text{MAX_ORDER}$ // Rényi moments
rounds_completed $\leftarrow 0$
function AllocateRound():
 $\epsilon_{spent} \leftarrow \text{ComputeEpsilon}(\lambda, \delta_{total})$
 $\epsilon_{remaining} \leftarrow \epsilon_{total} - \epsilon_{spent}$
if $\epsilon_{remaining} < \epsilon_{min}$ **then**
 raise BudgetExhaustedError
 $\epsilon_t \leftarrow \epsilon_{remaining} / (T - \text{rounds_completed})$
return ϵ_t
function RecordRound(σ, q):
for order = 1 to MAX_ORDER **do**
 $\lambda[\text{order}] += \text{ComputeMoment}(\text{order}, \sigma, q)$
rounds_completed $+= 1$

Registry lookups use LRU caching (max 100K records, TTL 10 min) with 300-second synchronization intervals to national registries. Configurable actions on opt-out match: exclude, anonymize, or error. Pre-training filtering ensures only opted-in data enters gradient computation.

1.4 Cross-Border Coordinator (Articles 46, 50). Implements multi-HDAB synchronization protocols across 10 pre-configured EU country profiles (DE, FR, IT, ES, NL, SE, PL, AT, BE, PT). Each profile specifies per-country constraints: ϵ_{max} for differential privacy, data retention periods (365–1,825 days), opt-out rates (3–12%), and network latency characteristics (15–60ms based on geographic distance). The fee model (Article 42) implements cost-recovery calculation: base fee + per-record + per-round + per-MB charges.

Inter-layer interface: Data Permit Authorization (OAuth2/mTLS, purpose-validated, Articles 53+71) connects Layer 1 to Layer 2.

2) *Layer 2: FL Orchestration:* Five modules operate within the Secure Processing Environment (SPE), conforming to HealthData@EU infrastructure requirements:

2.1 Aggregation Engine. Implements 17 FL algorithms spanning six categories: *Baseline* (FedAvg, FedProx with $\mu=0.01$, SCAFFOLD, FedNova, FedDyn); *Adaptive* (FedAdam with server_lr=0.1, FedYogi with $\beta_2=0.99$, FedAdagrad); *Personalization* (Ditto with $\lambda=0.1$, Per-FedAvg via MAML, pFedMe, MOON); *Non-IID* (FedLC for logit calibration, FedDecorr); *SAM* (FedSAM, FedSpeed); *Advanced* (FedExp ICLR’23, FedLESAM ICML’24, HPFL ICLR’25). Client selection strategies: random, performance-based, fairness-aware, latency-aware. Early stopping: patience=10, min_delta=0.001. Model checkpointing with atomic saves.

2.2 Privacy Protection Module. Three sub-modules provide defense-in-depth:

- *DP-SGD:* Gaussian mechanism with Rényi DP (RDP) accounting. Configurable ϵ -budget (default 1.0), $\delta=10^{-5}$. Gradient clipping: max_norm=1.0, type=L2/L ∞ . RDP composition provides 5–6 $\times$ tighter bounds than naive composition over 100+ rounds.

- *Secure Aggregation:* Pairwise masking protocol with ECDH key exchange (SECP384R1 curve). Alternative protocols: Shamir secret sharing (threshold reconstruction), homomorphic encryption (CKKS scheme via TenSEAL). Dropout threshold: 50%.
- *Byzantine Resilience:* Six defense rules (Krum, Multi-Krum, Trimmed Mean, Coordinate-wise Median, Bulyan, FLTrust). Anomaly detection at 3σ threshold. TEE integration for SGX/TrustZone hardware attestation.

2.3 Compliance Module. GDPR audit logging (Article 30) in JSON format with 7-year retention (2,555 days). Per-round purpose limitation check (Article 53). Opt-out enforcement: pre-training filtering combined with per-round verification. Anonymity assessment: k -anonymity, l -diversity checks. Automated EHDS compliance report with verification scoring.

2.4 Communication Infrastructure. gRPC bidirectional streaming for model updates (primary protocol). WebSocket for real-time monitoring and events. Compression: GZIP, LZ4, ZSTD, Snappy. Model chunking (ResNet-18: 44.7 MB/round $\rightarrow$ 8.9 MB with Top- k 1%). Security: TLS/mTLS with ECDHE key exchange. Retry: 3 attempts with $2\times$ exponential backoff. Message types: MODEL_UPDATE, GRADIENT_UPDATE, HEARTBEAT, PERMIT_VALIDATION, CONSENT_CHECK, AUDIT_LOG.

2.5 Monitoring & Metrics. Tracks training convergence (accuracy, loss, F1, AUC per round), communication cost (MB/round), privacy budget consumption (ϵ spent per round), participation statistics (samples, opt-outs), and resource utilization (CPU, memory, GPU). Structured logging via Structlog (JSON format) with Streamlit dashboard integration.

Inter-module flows: Privacy processing (DP-SGD, gradient clipping, secure masking) connects Aggregation $\rightarrow$ Privacy. Audit trail (GDPR Article 30, JSON logging, 7-year retention) connects Privacy $\rightarrow$ Compliance. Transport layer (gRPC streaming, TLS/mTLS, GZIP compression) connects Communication $\leftrightarrow$ all modules.

3) *Layer 3: Data Holders:* Each data holder institution (hospitals, disease registries, research centers across 27 Member States) implements a uniform component stack:




Fig. S-2. FL-EHDS detailed architecture with component-level specifications. Layer 1 (Governance) includes four modules: HDAB Integration with OAuth2/mTLS authentication, Data Permit Manager (Article 53) with lifecycle management, Opt-Out Registry (Article 71) with LRU-cached filtering, and Cross-Border Coordinator (Articles 46, 50) with per-country privacy constraints. Layer 2 (FL Orchestration) operates within the Secure Processing Environment (SPE) with five modules: Aggregation Engine (17 FL algorithms), Privacy Protection (DP-SGD, Secure Aggregation, Byzantine Resilience), Compliance Module (GDPR audit logging), Communication Infrastructure (gRPC/WebSocket), and Monitoring & Metrics. Layer 3 (Data Holders) implements a uniform stack per institution: Data Store, FHIR R4 Preprocessing, OMOP-CDM Harmonization, Local Training Engine, Gradient Processing, and Secure Communication. Inter-layer flows: Data Permit Authorization (OAuth2/mTLS, purpose-validated) downward from Layer 1; gradient upload (masked, encrypted, gRPC) upward and global model distribution (encrypted, TLS) downward between Layers 2–3.

3.1 Data Store. Institutional health data in heterogeneous formats: EHR (FHIR R4), vitals (LOINC-coded), diagnoses (ICD-10 variants per country: ICD-10-GM for Germany, CIM-10 for France, ICD-9-CM legacy for Italy), ECG records (SCP-ECG standard EN 1064), and medical imaging (DICOM).

3.2 FHIR R4 Preprocessing Engine. Processes six FHIR resource types: Patient, Observation, Condition, Medication-Request, Procedure, DiagnosticReport. Terminology mapping to international standards: SNOMED-CT (clinical concepts), LOINC (laboratory codes), ICD-10 (diagnoses), ATC (medications). Extracts 36 standardized features with normalization, encoding, and missing value imputation.

3.3 OMOP-CDM Harmonizer. Maps local vocabulary codes to standard OMOP Concept IDs. Populates OMOP tables: Person, ConditionOccurrence, Measurement, DrugExposure, ProcedureOccurrence, VisitOccurrence. Per-country vocabulary mapping ensures cross-border semantic compatibility.

3.4 Local Training Engine. Adaptive training with configurable parameters: optimizer (Adam, lr=0.001), batch size (32–256, dynamically adjusted), local epochs (3–5), dropout (0.3), L2 regularization. Device support: CUDA (GPU), MPS (Apple Silicon), CPU fallback for resource-constrained institutions.

3.5 Gradient Processing. Three-stage pipeline: (1) gradient

clipping (max_norm=1.0, L2 norm); (2) local DP noise addition (Gaussian, calibrated to ϵ); (3) pairwise masking (ECDH shared keys with counterpart clients).

3.6 Secure Communication. AES-256-GCM symmetric encryption for gradient payloads. ECDHE key exchange (SECP384R1 curve). mTLS mutual authentication with certificate-based identity. Nonce generation for replay attack prevention.

Inter-layer data flow: Gradient upload (∇ masked, encrypted, gRPC) flows upward from Layer 3 to Layer 2. Global model distribution (θ encrypted, TLS) flows downward from Layer 2 to Layer 3.

Architectural invariant: Raw health data never leaves institutional boundaries—only encrypted model gradients are exchanged within the Secure Processing Environment.

4) Deployment Readiness Assessment: Table S-II provides a transparent assessment of each FL-EHDS component’s production readiness, distinguishing between fully implemented modules, simulation backends, and components requiring external integration.

TABLE S-II
FL-EHDS COMPONENT READINESS FOR EHDS PRODUCTION DEPLOYMENT. ✓: PRODUCTION-READY. ~: SIMULATION BACKEND (FUNCTIONAL, REQUIRES BINDING TO EXTERNAL SERVICES). ○: REQUIRES EXTERNAL INFRASTRUCTURE.

Component	Status	Binding Req.
<i>Layer 2: FL Orchestration</i>		
17 FL algorithms	✓	None
DP-SGD + RDP accounting	✓	None
Secure aggregation	✓	None
Byzantine resilience	✓	None
Monitoring dashboard	✓	None
<i>Layer 3: Data Holders</i>		
Local training engine	✓	None
FHIR R4 preprocessing	✓	FHIR server URL
OMOP-CDM harmonization	✓	CDM schema
Secure communication	✓	TLS certificates
<i>Layer 1: Governance</i>		
HDAB permit lifecycle	~	HDAB REST/gRPC
Art. 71 opt-out registry	~	National registry API
Cross-border coordinator	~	Multi-HDAB endpoints
Art. 53 purpose limitation	~	Permit schema
GDPR audit trail	✓	None

Simulation backends implement complete functional logic (OAuth2/mTLS auth, permit CRUD, LRU-cached registry lookups) and require only endpoint configuration—not architectural changes—for production binding. HDAB services expected 2027–2029.

C. Infrastructure Components

Deploying FL across 27 EU Member States requires production-grade infrastructure: reliable communication channels between hospitals and the SPE aggregator, efficient serialization of gradient tensors, distributed coordination for concurrent studies, and comprehensive monitoring with EHDS-specific alerting. This section describes the infrastructure components implemented in the reference framework, each designed to operate within the constraints of cross-border healthcare networks (firewalls, bandwidth limitations, regulatory requirements).

1) *Communication Layer*: The communication layer must bridge heterogeneous network environments: high-bandwidth data center connections between national HDABs, moderate hospital-to-aggregator links, and potentially bandwidth-constrained rural clinics. The framework supports two transport protocols selectable per deployment, with configurable compression and retry policies.

gRPC: Bidirectional streaming, Protocol Buffers (30% bandwidth reduction vs. JSON), HTTP/2 multiplexing. Ideal for data center deployments.

WebSocket: Browser-compatible, firewall-friendly (standard HTTP upgrade), event-driven. Ideal for edge deployments and browser-based participation.

Selection criteria: gRPC is recommended for production EHDS deployments where both endpoints support HTTP/2 (typical for hospital-to-national aggregator links). WebSocket is preferred when traffic must traverse web application firewalls or when browser-based dashboards participate directly in federation monitoring.

2) *Serialization*: **Binary Format**: Tensor metadata + raw binary, 30% smaller than JSON, 15% smaller than pickle,

cross-platform (Python, C++, Java).

Delta Serialization: Transmits only changed parameters, sparse encoding, up to 90% bandwidth reduction for fine-tuning.

EHDS-Compliant: Embeds permit ID, timestamp, provenance; cryptographic signatures; GDPR Article 30 audit fields.

3) *Caching Layer*: In production EHDS deployments, multiple FL studies may run concurrently on overlapping data holders. A distributed locking mechanism prevents race conditions during gradient aggregation—ensuring that two concurrent studies do not interfere with each other’s model updates. Algorithm S14 implements Redis-based distributed locking with TTL-based automatic release, preventing deadlocks if a server node fails mid-aggregation.

Redis-based caching: model checkpoints, client states, real-time metrics. Features: LRU/LFU/TTL eviction, distributed locking, automatic serialization, cache warming.

4) *Orchestration*: **Kubernetes**: Deploys FL clients/aggregators as pods, HPA for elastic scaling, ConfigMaps for hyperparameters, Secrets for HDAB API keys.

Ray: Actor-based FL, automatic fault tolerance, Ray Tune for federated HPO, Object Store for gradient sharing.

Auto-Scaling: Reactive (queue depth/latency), Predictive (ML-based forecasting), Scheduled (time-based patterns).

5) *Monitoring*: **Prometheus Metrics**: Counters (rounds_total, permits_validated), Gauges (active_clients, privacy_budget_remaining), Histograms (round_duration, communication_latency), Summaries (gradient_norm_quantiles).

Grafana Dashboards: FL training progress, client health, latency heatmaps, privacy budget consumption, EHDS com-

Communication Manager Configuration

```
transport: gRPC | WebSocket
compression: gzip | lz4 | zstd | none
chunk_size: 1MB
retry_policy:
  max_retries: 3
  backoff: exponential
  base_delay: 1s
connection_pool:
  max_connections: 100
  idle_timeout: 300s
```

Algorithm S14: Distributed Lock for Aggregation

Input: Lock name, TTL, client ID

Output: Lock acquired (boolean)

acquired $\leftarrow$ Redis.SET(lock_name, client_id, NX, EX=TTL)

if acquired **then**

 PerformAggregation()

if Redis.GET(lock_name) == client_id **then**

 Redis.DEL(lock_name)

return acquired

pliance status.

Alerting: Privacy budget exhaustion, client dropout threshold, model divergence, permit expiration.

6) *Model Watermarking:* IP protection for FL models trained on EHDS data: **Spread Spectrum** (frequency domain, robust to fine-tuning), **LSB** (low-order weight bits), **Backdoor-based** (input-output ownership proof), **Passport Layers** (dedicated ownership encoding).

7) *Cross-Silo Enhancements:* EHDS federations are inherently cross-silo: each participant is an institution (hospital, registry, research center) with significant computational resources, distinct data distributions, and long-term participation commitments. This differs from cross-device FL (e.g., mobile phones) and enables advanced optimization strategies.

Multi-Model Federation: Weighted voting, stacking, mixture of experts with diversity enforcement.

Automatic Algorithm Selection: The 17 FL algorithms in the framework have different strengths depending on the federation characteristics (heterogeneity level, number of participants, communication budget). Algorithm S15 implements adaptive aggregation selection via multi-armed bandit (UCB/Thompson Sampling), automatically switching algorithms mid-training if performance metrics indicate a better alternative. A cooldown period prevents oscillation between strategies.

D. Extended EHDS Interoperability

1) *OMOP Common Data Model:* OMOP CDM v5.4 provides standardized analytical format used by European research networks (EHDEN, OHDSI).

ETL Pipelines: Transform source EHR to OMOP. **Vocabulary Mapping:** SNOMED, ICD10, LOINC, RxNorm. **Cohort Definitions:** ATLAS-compatible SQL generation. **Feature Extraction:** FeatureExtraction package for ML-ready datasets.

FL Integration: (1) Each hospital transforms local EHR to OMOP; (2) Feature extraction produces identical schema; (3) FL training on homogeneous feature spaces.

2) *IHE Integration Profiles:* **ATNA:** TLS mutual authentication, syslog audit messages (RFC 5424), maps to GDPR Article 30.

BPPC: Maps Article 71 opt-out to consent documents, XDS.b integration, consent enforcement at FL initiation.

XCA: Cross-border document query/retrieve, Initiating/Responding Gateways, patient identity correlation.

PIX/PDQ: Patient matching across boundaries, pseudonymization-aware identity management, national eHealth integration.

XUA: SAML 2.0 federated authentication, role-based access control, HDAB authorization token propagation.

3) *Cross-Border Data Exchange:* **Message Formats:** EHDS Data Permit Exchange Format (JSON-LD), Federated Query Protocol (SPARQL Federation), Model Update Message Format (Protocol Buffers).

Security: eIDAS-compliant electronic signatures, TLS 1.3, certificate-based authentication (EU trust framework).

Metadata: DCAT-AP Health extension, W3C PROV-O provenance, EMA data quality indicators.

4) *Interoperability Architecture:* Figure S-3 presents the complete interoperability architecture, showing how heterogeneous data sources across EU Member States are harmonized through multiple standards layers before reaching the FL training engine. The architecture reflects a key EHDS challenge: real-world healthcare institutions use diverse formats, terminologies, and exchange protocols that must be reconciled to produce a consistent feature space for federated model training.

Algorithm S15: Adaptive Aggregation**Input:** Client updates, metrics history, cooldown**Output:** Aggregated model, selected algorithmscore $\leftarrow$ WeightedScore(loss, accuracy, variance, conv.)**if** RoundsSinceSwitch > Cooldown **then** **for each** candidate $\in$ Algorithms **do** alt $\leftarrow$ EstimatePerformance(candidate) **if** alt > score + Threshold **then**

SwitchTo(candidate)


aggregated $\leftarrow$ CurrentAlgo.Aggregate(updates)**return** aggregated

Fig. S-3. EHDS interoperability architecture for FL-based secondary use. Data from heterogeneous sources across 27 Member States (top) flows through terminology harmonization (SNOMED CT, ICD-10, LOINC, ATC, UCUM), then through interoperability standards (FHIR R4 for structured data exchange, OMOP CDM for observational research, IHE profiles for cross-institutional workflows, cross-border exchange protocols with eIDAS signatures). A security and compliance layer enforces TLS 1.3 mutual authentication, eIDAS electronic signatures for permits, GDPR Article 30 audit logging, and W3C PROV-O data provenance before harmonized feature tensors reach the FL training engine. Bidirectional ETL between FHIR and OMOP enables institutions to use either standard based on their existing infrastructure.

E. Extended Threat Model and Security Analysis

The main paper summarizes the FL-EHDS threat model in three adversary classes. This section provides a comprehensive security analysis with formal definitions, attack taxonomy, defense mapping, and explicit boundary conditions relevant to EHDS cross-border deployments.

1) *Adversary Model:* We formalize three adversary classes ordered by increasing capability.

A1: Honest-but-curious aggregation server. The SPE aggregation server $\mathcal{S}$ follows the protocol faithfully but attempts to infer patient-level information from the received model updates $\{\Delta_k\}_{k=1}^K$. This is the most realistic adversary for

EHDS: the HDAB-operated SPE has institutional incentives to follow the protocol (legal liability, regulatory oversight) but may be compromised or subject to insider threats. Concrete attack vectors include:

- *Gradient inversion* [21]: reconstructing training samples from observed gradients. Effectiveness decreases with batch size and model depth; for the tabular MLP ($\sim 2.9K$ – $10K$ parameters depending on dataset, batch size 32–64), theoretical reconstruction quality is limited but non-zero.
- *Membership inference*: determining whether a specific patient record was used in a client's local training, potentially violating GDPR Article 9 (special category

data).

- *Property inference*: extracting aggregate properties of a client’s local dataset (e.g., disease prevalence at a specific hospital).

Defenses for A1: Central DP (Algorithm S2) with Gaussian mechanism ($\sigma = C/\varepsilon$, $\delta = 10^{-5}$) bounds per-round information leakage. Rényi DP composition (Algorithm S8) tracks cumulative privacy expenditure across rounds. Secure aggregation (Algorithm S4) ensures $\mathcal{S}$ observes only $\sum_k \Delta_k + \text{noise}$, never individual Δ_k . Our DP ablation (Table S-XIV) demonstrates that at $\varepsilon = 10$, personalized methods lose $<2\text{pp}$ accuracy—privacy imposes negligible utility cost at this budget.

A2: Malicious clients (Byzantine). Up to $f < n/3$ compromised institutions may deviate arbitrarily from the protocol. In the EHDS context, this could represent: institutions with corrupted data pipelines, nation-state actors attempting to bias cross-border models, or free-riding participants that send random updates to benefit from the global model without contributing genuine local training. Attack types include:

- *Model poisoning*: submitting adversarial gradients designed to degrade global model accuracy (untargeted) or introduce backdoors for specific inputs (targeted).
- *Data poisoning*: manipulating local training data (label flipping, sample injection) to corrupt the learned model through legitimate protocol participation.
- *Free-riding*: sending random or stale gradients to receive the global model without incurring computational cost, degrading convergence.

Defenses for A2: Six Byzantine-resilient aggregation rules are implemented (Section C): Krum (Algorithm S10), Trimmed Mean, Coordinate-wise Median, Bulyan (two-stage Krum + trimmed mean), FLTrust (server-guided trust weighting), and FLAME (clustering-based). The framework supports attack simulation (label flipping, gradient scaling, additive noise, sign flipping, model replacement) for defense validation.

A3: External attacker. An adversary with black-box access to model outputs (e.g., through a compromised data permit or leaked model weights) may attempt membership or attribute inference against the published model. Unlike A1, this adversary does not observe training dynamics but only the final or intermediate model checkpoints.

Defenses for A3: Article 71 output filtering restricts model access to authorized queries specified in the HDAB data permit, limiting the number of queries available for inference attacks. The permit system (Algorithm S3) enforces purpose limitation (Article 53), temporal validity, and data category authorization, reducing the attack surface. Model watermarking (Section B-C) provides post-hoc attribution for leaked models.

2) *Protected Assets*: Table S-III maps the three categories of protected assets to their defense mechanisms and the EHDS articles that mandate their protection.

3) *Out-of-Scope Threats and Boundary Conditions*: We explicitly delineate threats that the current FL-EHDS implementation does not address, along with justifications and potential future mitigations.

Server-client collusion. If the aggregation server $\mathcal{S}$ colludes with one or more clients, secure aggregation guarantees are invalidated: the colluding parties can reconstruct non-colluding clients’ individual gradients by subtracting known values from the aggregate. This threat undermines the trust model of *any* FL system and requires stronger cryptographic primitives (fully homomorphic encryption, multi-party computation with dishonest majority) that impose $100\text{--}1000\times$ computational overhead, currently impractical for healthcare-scale deployments. In the EHDS context, collusion is partially mitigated by the institutional separation between HDABs (which operate the SPE) and data holders (healthcare institutions), as these are distinct legal entities with separate governance structures.

Side-channel attacks. Timing analysis of gradient uploads, memory access patterns, or power consumption during local training could leak information about dataset properties. These attacks require physical or network-level proximity to the target institution and are addressed by infrastructure-level security (network isolation, constant-time implementations) rather than the FL protocol itself. EHDS Article 50 SPE specifications should mandate side-channel-resistant computing environments.

Covert-channel leakage through model architecture. The choice of model architecture, hyperparameters, and training configuration can inadvertently encode information about the training data [23]. For example, a model’s capacity (number of parameters) relative to dataset size affects memorization risk. This remains an open problem in FL privacy research; we mitigate it partially through standardized configurations (same architecture and hyperparameters across all clients) and DP noise injection [23].

Reconstruction attacks with auxiliary data. If an adversary possesses auxiliary information about a target individual (e.g., partial medical records from a data breach), gradient inversion and membership inference attacks become significantly more effective. Our threat model assumes adversaries operate within the bounds of their adversary class (A1–A3) without additional auxiliary datasets. In practice, EHDS cross-border settings increase this risk, as health data from one Member State could serve as auxiliary information for attacks on another Member State’s data.

Model extraction and intellectual property. A permitted user could systematically query the model to create a functionally equivalent copy, bypassing the HDAB access controls for subsequent use. This is partially addressed by query rate limiting and model watermarking, but a determined adversary with sufficient queries can always approximate any black-box model.

F. Framework Positioning: FL-EHDS vs. Existing Platforms

The main paper’s Table I provides a feature-level comparison between FL-EHDS and existing FL frameworks (Flower [45], FLARE [46], PySyft, FedML, TFF). A natural reviewer question is why we do not provide a direct *experimental* comparison. This section explains the rationale.

TABLE S-III
PROTECTED ASSETS AND DEFENSE MAPPING

Asset	Defense	EHDS Article
Raw patient data	Never leaves institution	Art. 50 (SPE)
Gradient updates	DP + secure aggregation	Art. 50, GDPR 32
Global model	HDAB permit access control	Art. 53, 71
Audit trail	Immutable logging	GDPR Art. 30
Privacy budget	RDP accountant (Alg. S8)	Art. 50

Why no experimental comparison is provided. An experimental comparison between FL frameworks is meaningful only when the frameworks provide *different* capabilities on a *common* task. For pure FL algorithm performance (e.g., FedAvg accuracy on PTB-XL), all frameworks are equivalent by design: they implement the same canonical algorithms (FedAvg, FedProx, etc.) with identical mathematical formulations. Running FedAvg on Flower vs. FL-EHDS would produce statistically indistinguishable results, as both execute the same aggregation logic on the same data. The comparison would measure implementation overhead (framework startup time, serialization efficiency), not scientific contribution.

The governance gap. FL-EHDS’s contribution is the governance-technical integration layer—the components that *no* existing framework provides:

- *HDAB permit lifecycle*: application, validation per round (Algorithm S3), temporal expiry, revocation, audit trail. Neither Flower nor FLARE implement permit-based access control.
- *Article 71 output filtering*: automated enforcement of opt-out registries for rare disease patients. No existing framework addresses EHDS opt-out requirements.
- *Article 53 purpose limitation*: technical enforcement of permitted secondary use purposes. Existing frameworks treat all training as equally authorized.
- *EHDS-compliant audit trails*: immutable logging of every round’s permit status, privacy budget consumption, and data category access. This is a regulatory requirement (GDPR Article 30) that no FL framework satisfies out of the box.
- *Cross-border coordination*: multi-HDAB orchestration for Article 46 cross-border processing with per-Member-State governance constraints.

Since these governance components have no counterpart in existing frameworks, a “comparison” would reduce to: (a) identical FL performance metrics, plus (b) a checklist of governance features present in FL-EHDS and absent in alternatives—which is precisely what Table I already provides.

Complementarity, not competition. FL-EHDS is designed as a compliance layer that could, in principle, wrap existing FL engines. The governance modules (permit validation, audit trails, privacy budget accounting, output filtering) are protocol-

agnostic and could be integrated with Flower’s communication backend or FLARE’s clinical deployment infrastructure. The architectural contribution is the *integration pattern*, not the replacement of existing FL platforms.

Additional frameworks. Two frameworks merit specific discussion given FLICS’s industrial focus. *OpenFL* (Intel) provides healthcare-specific FL with production deployments in clinical settings (e.g., Intel’s collaboration with the University of Pennsylvania for brain tumor segmentation), but lacks EHDS governance integration: no HDAB permit lifecycle, opt-out registry, or Article 53 purpose limitation. *FATE* (WeBank) offers industrial FL with governance features including role-based access control and audit logging, but is oriented toward financial services regulation (not healthcare/EHDS) and does not implement FHIR R4 interoperability or EU-specific data governance. Both could serve as FL compute backends beneath FL-EHDS’s governance layer, reinforcing our complementarity design.

G. Cross-Border Governance Challenges: Extended Analysis

The main paper argues that legal uncertainties—not technical barriers—constitute the critical blocker for FL adoption under the EHDS, and includes a condensed analysis of the cross-border privacy budget conflict (Section V-A). This section provides the extended analysis with additional scenarios—controller/processor ambiguity and gradient data classification—that illustrate how unresolved governance questions create implementation deadlocks that no engineering solution can circumvent.

1) *The Privacy Budget Conflict*: Consider a cross-border federation between a German hospital (operating under the Bundesdatenschutzgesetz, BDSG) and an Italian hospital (under the Codice in materia di protezione dei dati personali, D.Lgs. 196/2003 as amended). The German Data Protection Authority (BfDI) may interpret GDPR Article 89 to require $\varepsilon_{\max} = 1.0$ for health data secondary use, reflecting the “data minimization” principle. The Italian Garante per la protezione dei dati personali may permit $\varepsilon_{\max} = 5.0$ under a “proportionality” interpretation that balances research utility against privacy risk.

This creates a concrete governance deadlock:

- If the federation applies $\varepsilon = 1.0$ (strictest bound), Italian data contributes under unnecessarily restrictive

conditions, reducing model utility. Our DP experiments (Table S-XIV) show that FedAvg accuracy on PTB-XL drops from 91.9% (no DP) to 52.3% at $\varepsilon = 1$ —a 39.6pp collapse.

- If the federation applies $\varepsilon = 5.0$ (most permissive bound), German data may be processed in violation of BfDI guidance, exposing the German institution to regulatory sanctions.
- A per-client ε approach (each hospital applies its own budget) introduces asymmetric noise levels, creating fairness concerns: the German hospital’s model updates are noisier, potentially biasing the global model away from the German patient population.

FL-EHDS addresses this technically through per-client privacy budget configuration in the governance layer, but the *policy question*—which bound applies—requires regulatory clarification in the March 2027 delegated acts. Article 50(4) of Regulation (EU) 2025/327 mandates that SPEs provide “a high level of security,” but does not specify whether privacy budgets should be harmonized across Member States or determined by the most restrictive participant.

2) *Controller/Processor Ambiguity*: Under GDPR Articles 26 and 28, the roles of data controller and data processor carry distinct legal obligations. In a federated learning architecture:

- Each hospital is clearly the *controller* of its local patient data.
- The HDAB operating the aggregation server may be a *processor* (processing gradients on behalf of hospitals) or a *joint controller* (determining the purposes and means of the aggregation).
- If gradient updates are classified as personal data (an unresolved question), the aggregation server becomes a processor of personal data, requiring a Data Processing Agreement (DPA) with each participating hospital across 27 Member States.

The practical consequence: a 5-hospital federation across 3 Member States would require bilateral DPAs between each hospital and the HDAB, reviewed by 3 national DPAs, before a single training round can execute. This administrative overhead can delay projects by 12–18 months, as observed in current cross-border health research projects [5].

3) *Gradient Data Classification*: Whether model gradients constitute “personal data” under GDPR Article 4(1) remains unsettled. Zhu et al. [21] demonstrate that gradients can reconstruct training images, suggesting they encode personal information. However, with DP noise ($\varepsilon = 10$, $\delta = 10^{-5}$), the reconstruction fidelity degrades substantially, potentially rendering the gradients “anonymous” under Recital 26.

This classification has cascading regulatory implications:

- *If personal data*: full GDPR applies to gradient transmission, requiring lawful basis (Article 6), special category safeguards (Article 9), Data Protection Impact Assessments (Article 35), and cross-border transfer mechanisms (Chapter V).

- *If anonymous*: GDPR does not apply, and gradients can be transmitted freely—but the anonymization claim must be defensible against re-identification attacks, which our threat model (Section B-E) acknowledges as an open problem.

The EHDS Regulation (EU) 2025/327 does not resolve this question: Article 50 defines SPE security requirements but does not classify the privacy status of intermediate computational artifacts such as gradients.

4) *Implications for FL-EHDS Design*: These governance challenges directly motivated FL-EHDS’s architecture:

- The per-client privacy budget in the governance layer (Algorithm S8) is designed to accommodate heterogeneous national requirements, enabling the German hospital to operate at $\varepsilon = 1.0$ while the Italian hospital uses $\varepsilon = 5.0$ —though the policy legitimacy of this approach awaits regulatory guidance.
- The HDAB permit system (Algorithm S3) externalizes the controller/processor determination: the HDAB issues the permit (acting as controller of the secondary use authorization) while each hospital retains controller status over its data. This separation is architecturally explicit but legally untested.
- The immutable audit trail provides the evidentiary basis for demonstrating compliance under either gradient classification scenario, satisfying GDPR Article 30 requirements regardless of the ultimate regulatory interpretation.

These examples demonstrate that the “legal uncertainties as critical blocker” claim in the main paper is not abstract: it reflects specific, concrete impediments that affect system design decisions at the architectural level.

APPENDIX C ADVANCED FL PARADIGMS

Note: The paradigms in this section are *implemented in the reference framework* and available for use. Four of these paradigms—fairness-aware FL, continual FL, personalized FL, and federated unlearning—are now experimentally validated in Section E-F (135 experiments across 3 datasets). Vertical FL is now experimentally validated with 216 experiments across 3 datasets, 2–4 parties, overlap ablation, and DP (Section E-H.10, Table S-LXXX). The remaining paradigms (hierarchical FL, multi-task FL, asynchronous FL) are provided as architectural capabilities for future evaluation on appropriate multi-institutional datasets.

The core FL-EHDS pipeline (Section II) addresses the standard “horizontal” FL scenario where all hospitals share the same feature schema. However, real EHDS deployments will encounter more complex configurations: institutions with complementary features for the same patients (vertical FL), adversarial participants (Byzantine resilience), evolving data distributions over the 2025–2031 timeline (continual FL), heterogeneous clinical objectives (multi-task FL), and the hierarchical governance structure of the EU itself (hierarchical FL). This section presents the advanced paradigms implemented in

the reference framework, each motivated by a specific EHDS deployment challenge. Algorithms S9–S13 formalize the core mechanisms.

A. Vertical Federated Learning

Vertical FL addresses scenarios where institutions hold *different features* for the *same patients*—a common situation in EHDS cross-border analytics. For example, a hospital may hold demographics and diagnoses, a laboratory holds test results, and a pharmacy holds prescription histories. Under EHDS Article 33, these correspond to different data categories (Patient Summary, Laboratory Results, E-Prescription) that may be held by different data holders within the same or different Member States.

Private Set Intersection (PSI): Before training, the participating institutions must identify their common patients without revealing their full patient lists. RSA-based PSI achieves this with $O(n \log n)$ complexity using pseudonymized identifiers, ensuring EHDS compliance: no institution learns which patients the other holds beyond the intersection.

Split Learning: Algorithm S9 implements the forward pass in split learning, where each party computes activations on its local features up to a “cut layer,” then the server concatenates activations to produce the final prediction. Only intermediate representations (not raw data) cross institutional boundaries.

B. Byzantine-Resilient Aggregation

In a cross-border EHDS federation spanning 27 Member States, the aggregation server cannot blindly trust every participant. A compromised institution—whether through malicious intent, software bugs, or data corruption—could submit adversarial gradient updates that poison the global model, potentially affecting clinical decisions across the entire federation. Byzantine-resilient aggregation protects model integrity by detecting and excluding anomalous updates.

Algorithm S10 implements Krum, which selects the gradient closest to $n-f-2$ nearest neighbors, effectively filtering outliers. Six defense methods protect against up to $f < n/3$ adversarial clients:

Other methods: **Trimmed Mean** (removes β -fraction extreme values per coordinate), **Coordinate-wise Median** (robust estimator), **Bulyan** (two-stage Krum + trimmed mean), **FLTrust** (server-guided trust weighting), **FLAME** (clustering-based). Attack simulation: label flipping, gradient scaling, additive noise, sign flipping, model replacement.

C. Continual Federated Learning

The EHDS is designed for long-term operation (2025–2031 and beyond), during which healthcare data distributions will evolve: new diseases emerge (as demonstrated by COVID-19), clinical protocols change, and demographic compositions shift. A model trained in 2027 may perform poorly on 2029 data if it has “forgotten” how to handle earlier patterns. Continual Federated Learning addresses this *catastrophic forgetting* problem by preserving knowledge from previous training tasks while adapting to new data.

The Elastic Weight Consolidation (EWC) loss function adds a quadratic penalty:

$$\mathcal{L}_{EWC}(\theta) = \mathcal{L}(\theta) + \frac{\lambda}{2} \sum_i F_i (\theta_i - \theta_i^*)^2$$

where $F_i = \mathbb{E}[\nabla^2 \log p(\mathcal{D}|\theta^*)]_i$ is the i -th diagonal entry of the Fisher Information Matrix and θ^* are optimal parameters for previous tasks.

Additional strategies: Learning without Forgetting (LwF), Experience Replay, drift detection (ADWIN, Page-Hinkley) triggering adaptation.

D. Multi-Task Federated Learning

In EHDS cross-border studies, different hospitals may pursue related but distinct clinical objectives from the same data. A cardiology network might simultaneously predict heart failure risk (Hospital A), readmission probability (Hospital B), and medication response (Hospital C). Multi-Task FL enables these institutions to collaborate on shared feature representations while maintaining task-specific prediction heads.

Architectures: **Hard Parameter Sharing** (common feature extractor, task-specific heads), **Soft Parameter Sharing** (separate networks with similarity regularization), **FedMTL** (dynamic task relationship learning).

E. Hierarchical Federated Learning

The EHDS governance structure is inherently hierarchical: individual hospitals report to regional health authorities, which coordinate under national HDABs, which in turn connect to the EU-level HealthData@EU infrastructure. Hierarchical FL mirrors this governance topology, aggregating gradients at intermediate levels before reaching the central server. This reduces communication costs (hospitals communicate with regional aggregators, not directly with the EU server) and aligns FL operations with the jurisdictional boundaries of HDABs.

Four-tier hierarchy reflecting EU governance:

- 1) **Client Tier:** Individual hospitals/data holders
- 2) **Regional Tier:** Regional aggregators (e.g., Lombardy, Bavaria)
- 3) **National Tier:** National HDABs coordinate Member State aggregation
- 4) **EU Tier:** HealthData@EU central aggregator

Benefits: reduced communication costs (hospitals → regional, not directly EU), alignment with EHDS governance where HDABs have national jurisdiction.

F. Personalized Federated Learning

A single global model may underperform at individual hospitals because clinical populations differ substantially across Member States. Personalized FL maintains both a global model (encoding shared medical knowledge) and hospital-specific local models (adapted to local demographics and clinical practices). Algorithm S11 shows pFedMe [35] as a representative personalized method, using Moreau envelopes

Algorithm S9: Split Learning Forward Pass**Input:** Features X_A, X_B at parties A, B; cut layer k **Output:** Prediction $\hat{y}$

```

 $h_A \leftarrow f_{1:k}^A(X_A)$            // Party A: features  $\rightarrow$  activations
 $h_B \leftarrow f_{1:k}^B(X_B)$        // Party B: features  $\rightarrow$  activations
 $h \leftarrow \text{Concat}(h_A, h_B)$ 
 $\hat{y} \leftarrow f_{k+1:L}(h)$        // Server: cut layer  $\rightarrow$  output
return  $\hat{y}$ 

```

Algorithm S10: Krum Byzantine Defense**Input:** Gradients $\{g_1, \dots, g_n\}$, Byzantine bound f **Output:** Selected gradient g^*

```

for each gradient  $g_i$  do
   $D_i \leftarrow \{\|g_i - g_j\|^2 : j \neq i\}$ 
   $s_i \leftarrow \sum_{d \in \text{smallest}_{n-f-2}(D_i)} d$ 
 $g^* \leftarrow g_{\arg \min_i s_i}$ 
return  $g^*$ 

```

to balance personalization with global knowledge: the regularization term $\lambda(\theta_k - \theta)$ pulls the local model toward the global consensus, while local gradient descent adapts to hospital-specific data patterns. Ditto [33] follows a similar dual-model principle but with a simpler formulation: it trains a personalized model regularized toward the global model via $\frac{\lambda}{2}\|\theta_k - \theta\|^2$. In our experiments, Ditto—the best-performing personalized method—achieves 76.0% accuracy on Heart Disease (10-seed mean), an 11.2pp improvement over FedAvg, precisely because it learns hospital-specific decision boundaries.

Other approaches: **FedPer** (shared base, local personalization layers), **Per-FedAvg** (MAML-based meta-learning), **APFL** (adaptive mixing α between global and local), **Ditto** (personalization regularization).

EHDS Relevance: Member States have different healthcare systems, disease prevalence, and clinical practices. Personalized FL enables institution-specific adaptation while benefiting from collaborative training.

G. Asynchronous Federated Learning

Standard synchronous FL requires all participating hospitals to complete local training before the server can aggregate. In an EHDS federation spanning 27 Member States with heterogeneous computational resources, this creates a “straggler” problem: a resource-constrained rural hospital delays the entire federation. Asynchronous FL eliminates this bottleneck by allowing the server to aggregate updates as they arrive, weighting stale updates (from slow clients) less heavily. Algorithm S12 implements polynomial staleness weighting: an update computed τ rounds ago receives weight $(1 + \tau)^{-\alpha}$, ensuring that fresher updates contribute more while still incorporating information from slower participants.

Staleness functions: Constant ($\alpha=1$), Polynomial ($(1+\tau)^{-\alpha}$), Exponential ($e^{-\alpha\tau}$), Hinge (1 if $\tau \leq \tau_{max}$, else 0). Additional: FedBuff (buffered async), semi-async (wait for α -fraction of clients).

H. Fairness-Aware Federated Learning

The EHDS serves 450 million citizens across Member States with different population sizes, disease prevalence, and healthcare quality. Standard FL optimizes average performance, which can disproportionately favor large hospitals with more data while neglecting smaller institutions or underrepresented patient populations. This creates a “digital health equity” concern: a model that achieves 85% accuracy for a large German hospital but only 55% for a small Romanian clinic is not equitable. Algorithm S13 implements q-FedAvg, which reweights client contributions by their loss: hospitals where the model performs poorly receive higher aggregation weights, pulling the global model toward equitable performance across all participants.

Fairness metrics: Performance Variance ($\text{Var}(\{L_k\})$), Worst-case Loss ($\max_k L_k$), Demographic Parity Gap, Equalized Odds Gap. Additional methods: AFL, FedMGDA+, TERM, FairFed.

APPENDIX D EXPERIMENTAL METHODOLOGY

A. Systematic Cascading Experiment Registry

This section documents the complete cascading experiment campaign conducted for the FL-EHDS evaluation. Unlike ad-hoc experimental designs, we employ a *cascading methodology*: each phase builds upon findings from previous phases, progressively expanding coverage from foundational privacy-utility trade-offs (Cascade 2) to clinical imbalance mitigation (Cascade 10). This systematic approach ensures that (i) every experiment addresses a specific, pre-defined research question; (ii) intermediate results guide subsequent experimental design; and (iii) all results are reproducible via atomic checkpointing with explicit random seeds. The campaign comprises **2,092+ experiments** across **9 cascading phases**, **10 clinical datasets**, and **40 research questions**, constituting one of the most comprehensive FL evaluations in the healthcare domain.

Algorithm S11: pFedMe Local Update

Input: Data $\mathcal{D}_k$, global θ , personal θ_k , λ , η
Output: Updated personal model θ'_k

```

for  $i = 1$  to  $R$  do
   $\theta_k \leftarrow \theta_k - \eta \nabla \mathcal{L}(\theta_k; \mathcal{D}_k)$ 
  // Moreau envelope: balance with global
   $\theta'_k \leftarrow \theta_k - \lambda(\theta_k - \theta)$ 
   $g_k \leftarrow \lambda(\theta - \theta'_k)$ 
return  $\theta'_k, g_k$ 

```

Algorithm S12: FedAsync with Staleness Weighting

Input: Client update Δ_k , client round t_k , server round t
Output: Updated global model θ

```

 $\tau \leftarrow t - t_k$  // Staleness
 $\alpha \leftarrow (1 + \tau)^{-a}$  // Polynomial decay,  $a > 0$ 
 $\theta \leftarrow \theta + \alpha \cdot \eta \cdot \Delta_k$ 
return  $\theta$ 

```

Algorithm S13: q-FedAvg Fair Aggregation

Input: Losses $\{L_1, \dots, L_K\}$, updates $\{\Delta_1, \dots, \Delta_K\}$, q
Output: Fair aggregated update Δ

```

for each client  $k$  do
   $w_k \leftarrow L_k^q$  // Higher loss  $\rightarrow$  higher weight
 $W \leftarrow \sum_k w_k$ ;  $w_k \leftarrow w_k / W$ 
 $\Delta \leftarrow \sum_k w_k \cdot \Delta_k$ 
return  $\Delta$ 

```

1) *Cascading Methodology*: Each cascade is a self-contained experimental phase with the following properties:

- **Block structure**: Cascades are divided into thematic blocks (labeled A–TH), each addressing one or two research questions with controlled factorial designs.
- **Incremental coverage**: Early cascades (2–3) establish baselines; middle cascades (4–6) explore advanced paradigms; late cascades (7–10) address trustworthiness, communication efficiency, deployment readiness, and clinical imbalance mitigation.
- **Atomic checkpointing**: Every experiment result is saved immediately via an atomic write pattern (temporary file $\rightarrow$ fsync $\rightarrow$ rename), ensuring crash resilience and exact reproducibility.
- **Resumability**: All cascades support graceful interruption (SIGINT) and resume from the last checkpoint, enabling execution across multiple sessions.

2) *Campaign Summary*: Table S-IV provides an overview of all cascading phases. The total experiment count excludes the parallel imaging track (executed on separate hardware) and standalone validation files.

3) *Clinical Dataset Portfolio*: Table S-V summarizes the 10 clinical datasets spanning 7 medical specialties. The tabular datasets are used across all cascades; the imaging datasets are evaluated in a parallel track on dedicated GPU hardware.

4) *Cascade Descriptions and Key Findings*: **Cascade 2: Privacy–Efficiency Foundations** (Block D, 27 experiments). Investigates the interaction between differential privacy noise and learning rate selection across FedAvg, Ditto, and HPFL on

three core datasets (Cardiovascular, PTB-XL, Breast Cancer). *Key finding*: Learning rate sensitivity increases under DP; personalized algorithms (Ditto, HPFL) maintain stability across LR ranges where FedAvg collapses.

Cascade 3: Data Efficiency and Compound Stress (Blocks C–D, 75 experiments). Block C varies training data fractions (25–100%) to determine minimum viable dataset sizes for EHDS deployment. Block D applies simultaneous non-IID ($\alpha \in \{0.1, 0.5, 1.0\}$) + DP ($\epsilon=10$) + partial participation (60%) to test algorithm resilience under compound adversarial conditions. *Key findings*: 50% of training data achieves within 1pp of full-data performance. HPFL is uniquely robust to compound stress (91.1% vs. FedAvg 52.9% under triple adversity, +38.2pp gap).

Cascade 4: Advanced FL Paradigms (Blocks A–D, 135 experiments). Four blocks systematically evaluate fairness-aware aggregation (QFedAvg, AFL, FedMGDA, PropFair, FedMinMax), continual learning under concept drift (Replay, EWC, FedProx-continual), extended personalization (Per-FedAvg, FedPer, APFL, pFedMe vs. Ditto/HPFL), and federated unlearning for GDPR Article 17 compliance (exact retrain, gradient ascent, FedEraser). *Key findings*: FedMinMax solves majority-class collapse on Breast Cancer (+31.2pp). Experience Replay is the only effective continual strategy; EWC is counterproductive (−22.4pp). Gradient ascent achieves Article 17 compliance at 5% computational cost. Per-FedAvg (MAML-based) fails catastrophically (17.4% accuracy).

Cascade 5: Client-Centric Mechanisms (Blocks E–F, 87 experiments). Block E benchmarks 5 client selection strategies

TABLE S-IV
CASCADING EXPERIMENT CAMPAIGN OVERVIEW

Cascade	Blocks	Exp.	Time	Primary Focus
2	D	27	6:51	Privacy–efficiency foundations
3	C–D	75	18:37	Data efficiency, compound stress
4	A–D	135	11:52	Fairness, continual, personalization
5	E–F	87	3:56	Client selection, gradient inversion
6	G–L	234	30:27	Dataset coverage, advanced FL
7	M–R	243	25:13	Calibration, fairness, DP composition
8	S–W	119	10:57	SecAgg, compression, vertical FL
9	X–AB	144	18:24	Scalability, convergence, imbalance
10	AC–TH	384	22:16	Imbalance deep-dive, DP regularization
Standalone files		~1,431	—	Targeted validations
Gov.	H1–H3	720	3:40	EHDS governance hypotheses
Total		3,599+		

TABLE S-V
CLINICAL DATASET PORTFOLIO FOR FL-EHDS EVALUATION

Dataset	Samples	Features	Classes	Clinical Domain
<i>Tabular datasets</i>				
Cardiovascular	70,000	11	2	Cardiology
PTB-XL ECG	21,799	9	5	Cardiac electrophysiology
CDC Diabetes	253,680	21	2	Public health
Stroke	5,110	10	2	Neurology
CKD	399	24	2	Nephrology
Cirrhosis	418	18	2	Hepatology
Breast Cancer	569	30	2	Oncology
<i>Imaging datasets</i>				
Chest X-ray	5,856	images	2	Radiology
Brain Tumor MRI	7,023	images	4	Neuro-radiology
Skin Cancer	3,297	images	2	Dermatology

(Random, Loss-based, Importance, Resource-aware, Oort) under 60% participation. Block F evaluates gradient inversion attack vulnerability under varying DP. *Key findings*: Oort selection achieves both highest accuracy (78.5%) and fairness (Jain index 0.938). Without DP, 2/3 gradient inversion attacks succeed (cosine similarity 0.816); DP $\epsilon=10$ achieves 100% attack failure (MSE $1,927\times$ higher).

Cascade 6: Extended Evaluation (Blocks G–L, 234 experiments). Expands evaluation to 6 tabular datasets, tests hierarchical FL (2- and 3-level governance), asynchronous FL (staleness 0–5), multi-task learning (single/dual/triple task), Byzantine-robust aggregation (Median, Krum, Multi-Krum, Trimmed Mean, Bulyan against 20% Byzantine clients), and Shapley value attribution (2^K exact enumeration). *Key findings*: Personalization advantage generalizes across all 6 clinical domains. Hierarchical FL introduces no accuracy penalty ($<0.5\text{pp}$). Staleness ≤ 3 is tolerable ($<2\text{pp}$ degradation). Multi-Krum and Trimmed Mean fully neutralize model poisoning.

Cascade 7: Trustworthiness and Statistical Rigor (Blocks M–R, 243 experiments). Addresses model calibration (ECE before/after temperature scaling), conformal prediction (coverage validity under DP), feature attribution consistency across clients, demographic fairness analysis (age/sex subgroups), concept drift detection with adaptation, and DP composition analysis (simple, advanced, RDP for 1–20 sequential studies). *Key findings*: FL models are systematically overconfident; temperature scaling effective. Conformal prediction maintains 95% coverage under DP. Ditto reduces demographic accuracy gaps by up to 60%. RDP enables $5\times$ more sequential studies than simple composition.

Cascade 8: Communication, Privacy, and Scalability (Blocks S–W, 119 experiments). Evaluates secure aggregation (Pairwise Masking with ECDH + Shamir Secret Sharing), 7 gradient compression methods, local vs. central DP (3 datasets $\times$ 3 ϵ levels), vertical FL (2- and 3-party feature partitioning with PSI), and CDC Diabetes scalability (253K samples). *Key findings*: Pairwise masking adds sub-millisecond overhead with zero accuracy loss (mask cancellation error $< 10^{-16}$). SignSGD achieves $27.84\times$ compression with only -0.17pp loss. Ditto rescues the minority class under class imbalance (F1: 0.00 $\rightarrow$ 0.80 on CDC Diabetes NonIID). *Extended VFL evaluation*: 216 additional experiments across 3 datasets (Cardiovascular, PTB-XL binarized, CDC Diabetes), 2–4 parties, 3 overlap rates, and DP ablation validate VFL viability for European clinical data—PTB-XL binarized achieves 99.3% with 2 parties, while 3-party configurations are more DP-robust than 2-party (see Table S-LXXX).

Cascade 9: Deployment Readiness (Blocks X–AB, 144 experiments). Closes six remaining gaps: (i) scalability to $K=50$ clients on CDC Diabetes; (ii) combined DP + compression (SignSGD/QSGD + $\epsilon \in \{1, 10\}$); (iii) per-round convergence dynamics across 3 datasets; (iv) clinical imbalance robustness on Stroke (4.9% positive class), CKD, and Cirrhosis with F1/precision/recall/DEI metrics; (v) CDC Diabetes heterogeneity sensitivity ($\alpha \in \{0.1, 0.25, 0.75, 1.0\}$). *Critical finding*:

37% of experiments on imbalanced datasets exhibit majority-class collapse (F1=0.0), motivating Cascade 10.

Cascade 10: Clinical Imbalance Deep-Dive (Blocks AC–TH, 384 experiments). Systematically investigates and resolves the majority-class collapse phenomenon discovered in Cascade 9. Four blocks: (i) Block AC completes the condition matrix with 3 missing DP conditions (IID+ $\epsilon=10$, IID+ $\epsilon=1$, NonIID+ $\epsilon=1$) across Stroke/Cirrhosis $\times$ 3 algorithms $\times$ 5 seeds (90 exp.); (ii) Block AD evaluates class-weighted cross-entropy and focal loss ($\gamma=2.0$) as mitigation strategies across 3 datasets $\times$ 3 conditions $\times$ 3 algorithms (234 exp.); (iii) Block AE sweeps Ditto local epochs $\{5, 10\}$ on Stroke/Cirrhosis (40 exp.); (iv) Block TH re-evaluates 20 collapsed models from Cascade 9 with post-hoc threshold tuning (sweep 0.05–0.95). *Key findings*: DP noise acts as implicit regularizer (collapse rate: 83% without DP $\rightarrow$ 0% with any DP). Weighted CE eliminates collapse on all datasets (F1=0.39 on Stroke, from 0.00 baseline). Focal loss achieves zero collapses (0/117). Ditto at 10 local epochs doubles F1 on Cirrhosis (0.36 $\rightarrow$ 0.78). Threshold tuning rescues 7/7 Cirrhosis collapsed models (mean F1=0.65) but provides only marginal rescue for Stroke (mean F1=0.19).

5) *Standalone Validation Experiments*: In addition to the cascading phases, 1,431+ standalone experiments address targeted research questions:

- **Scalability + DP $\times$ TopK** (189 exp.): Client scaling $K \in \{3, 5, 10, 15, 20\}$ combined with TopK sparsification under DP.
- **Non-IID $\times$ DP interaction** (126 exp.): Systematic evaluation of how data heterogeneity interacts with privacy noise across multiple α and ϵ levels.
- **Byzantine + DP multi-defense** (198 exp.): Combined Byzantine resilience and differential privacy under label flipping and gradient scaling attacks. *Key finding*: Byzantine defenses and DP are synergistic; accuracy is DP-invariant under simultaneous defense.
- **DP per-class impact** (80 exp.): Fine-grained per-class accuracy analysis under varying DP levels on Breast Cancer and PTB-XL.
- **Membership inference attack** (18 exp.): MIA vulnerability across algorithms. HPFL AUC = 0.666 (vulnerable due to personalized heads); DP $\epsilon=1$ mitigates to 0.337.
- **Partial participation** (18 exp.): 3-of-5 client selection impact.
- **Privacy budget exhaustion** (9 exp.): Budget profile analysis (tight/medium/generous).
- **Centralized baselines** (9 exp.): Multi-dataset upper-bound reference (PTB-XL, Cardiovascular, CDC Diabetes) with centralized PyTorch training (Adam, 50 epochs, 3 seeds). *Key finding*: FL achieves $\leq 0.5\text{pp}$ of centralized performance on PTB-XL and Cardiovascular (Table S-LXXV).
- **Server-side optimizers** (108 exp.): Comparison of FedAdam, FedYogi, FedAdagrad [31], SCAFFOLD, FedNova, FedDyn across 3 datasets $\times$ IID/NonIID

$\times 3$ seeds. *Key finding:* Adaptive optimizers (FedAdam/Yogi/Adagrad) converge to identical solutions on compact tabular models; SCAFFOLD consistently worst (Table S-LXXVI).

- **10-seed extended robustness** (240 exp.): FedAvg, FedProx, Ditto, HPFL with 10 seeds across 3 datasets $\times$ IID/NonIID for enhanced statistical power. *Key finding:* FedAvg/FedProx/Ditto statistically indistinguishable on Cardiovascular IID; HPFL significantly lower on CDC Diabetes NonIID (-6.8pp ; Table S-LXXXI).
- **CDC Diabetes full evaluation** (214 exp.): Population-scale (253K samples) FL evaluation with 7 algorithms, DP ablation ($\varepsilon \in \{1, 5, 10, 50\}$), and client scaling ($K \in \{3, 5, 10, 20\}$). *Key finding:* DP paradoxically improves FedAvg ($+8.9\text{pp}$ at $\varepsilon=5$); HPFL counterproductive under class imbalance (Tables S-LXXVII–S-LXXIX).
- **Vertical FL extended** (216 exp.): 3 datasets (Cardiovascular, PTB-XL binarized, CDC Diabetes) $\times \{2, 3, 4\}$ parties $\times 3$ overlap rates $\times$ DP ablation. *Key finding:* PTB-XL binarized achieves 99.3% (2-party, no DP); CDC Diabetes exhibits majority-class collapse ($F1 \approx 0$) regardless of VFL configuration (Table S-LXXX).
- **EHDS Governance Hypotheses** (720 exp.): Three hypothesis-driven experiments on Heart Disease (920 samples) and Cardiovascular (70K samples). H1: compound compliance stress (3 algo $\times$ 4 conditions $\times$ 15 seeds). H2: non-uniform Article 71 opt-out fairness (5 patterns $\times$ 20 seeds). H3: algorithm governance resilience (3 algo $\times$ 4 purposes $\times$ 20 seeds). *Key finding:* Personalized algorithms (Ditto, HPFL) tolerate full EHDS compliance at $\varepsilon=10$ with $<1\text{pp}$ cost on CV; governance resilience reaches 94.8% vs. FedAvg 84.7% (Section G).

6) *Research Questions Index:* Table S-VI provides a cross-reference of all 40 research questions addressed across the cascading campaign, indexed by cascade and block identifier.

7) *Reproducibility Infrastructure:* All experiments share a common reproducibility infrastructure:

- **Deterministic seeding:** PyTorch manual seeds set before model creation; NumPy RandomState for data partitioning. Standard seeds: 42, 123, 456; extended to 5 seeds (42, 123, 456, 789, 999) in Cascade 10 for enhanced statistical power.
- **Atomic checkpointing:** Results persisted via `tempfile.mkstemp` $\rightarrow$ `os.fdopen` $\rightarrow$ `os.fsync` $\rightarrow$ `os.replace`, with `.bak` backup copies. This ensures no partial writes corrupt results.
- **Data caching:** Partitioned datasets cached within each cascade to prevent re-partitioning across experiments, ensuring identical client data splits.
- **Graceful shutdown:** All cascades handle SIGINT/SIGTERM, completing the current experiment and saving the checkpoint before exit.
- **Resume support:** Each cascade loads its checkpoint file at startup and skips already-completed experiments,

enabling multi-session execution.

The complete experiment source code, checkpoint files, and analysis scripts are available in the open-source repository at <https://github.com/FabioLiberti/FL-EHDS-FLICS2026>.

APPENDIX E

TABULAR EXPERIMENT RESULTS

A. Supplementary Experimental Figures

This section presents detailed experimental results from the FL-EHDS benchmark suite. All figures are generated from real experimental runs available in the repository.

Note on experimental configurations: Figures S-4–S-11 were generated from an extended 50-round, 5-client training run using the framework’s synthetic EHDS scenario (simulated European hospitals: Rome, Amsterdam, Berlin, Madrid, Paris). These complement the main paper’s 20-round experiments on Heart Disease UCI (4 real hospitals) and Diabetes (5 Dirichlet-partitioned clients). The 50-round configuration illustrates longer-horizon convergence properties, client participation dynamics, and gradient evolution patterns that are not visible in the shorter 20-round evaluation.

- 1) *Hospital Data Distribution:*
- 2) *Per-Client Training Time:*
- 3) *Client Participation Matrix:* Figure S-6 shows the client participation pattern across training rounds.
- 4) *Gradient Norm Evolution:* Figure S-7 tracks gradient norms per client, confirming convergence stability.
- 5) *Communication Cost Analysis:* Figure S-8 quantifies the per-round communication overhead.
- 6) *Learning Rate Sensitivity:* Figure S-9 evaluates sensitivity to the server learning rate.
- 7) *Batch Size Impact:* Figure S-10 examines the trade-off between batch size and convergence.
- 8) *Per-Client Accuracy Trajectories:* Figure S-11 shows per-client accuracy evolution, illustrating the impact of non-IID data.

B. Extended Algorithm Comparison

1) *Algorithms Evaluated:* We compare foundational FL algorithms plus 2022–2025 advances:

Foundational: FedAvg [13], FedProx [14], SCAFFOLD [29], FedAdam/FedYogi/FedAdagrad [31].

Recent (2022–2025): FedLC [39] (logit calibration for label skew), FedSAM [38] (flat minima), FedDecorr [40] (decorrelation against dimensional collapse), FedSpeed [41] (fewer rounds), FedExp [42] (server-side acceleration), FedLE-SAM [43] (globally-guided SAM, ICML 2024 Spotlight), HPFL [44] (personalized classifiers, ICLR 2025).

2) *Non-IID Configuration:* Data heterogeneity is controlled via Dirichlet distribution with α :

- $\alpha = 0.1$: **Extreme non-IID**—highly skewed label distributions
- $\alpha = 0.5$: **High non-IID**—significant heterogeneity
- $\alpha = 1.0$: **Moderate non-IID**—balanced heterogeneity
- $\alpha = 10.0$: **Near-IID**—approximately uniform




Fig. S-4. Data distribution across hospitals. Notable heterogeneity: Amsterdam shows older population (60.8 years mean age) with higher positive rate (62.6%) compared to Rome (49.1 years, 39.2%). This reflects realistic cross-border EHDS variability.




Fig. S-5. Per-client training time per round. Larger hospitals (Berlin: 500 samples) exhibit slightly longer training times. The adaptive training engine compensates by adjusting batch sizes for stragglers.

TABLE S-VI
COMPLETE RESEARCH QUESTION INDEX ACROSS CASCADING EXPERIMENTS

#	Research Question	Cascade.Block	Key Result
1	How does DP interact with learning rate selection?	2.D	Ditto/HPFL stable; FedAvg collapses
2	How much training data is needed for adequate FL?	3.C	50% data $\rightarrow$ within 1pp of full performance
3	How do algorithms perform under compound stress?	3.D	HPFL +38.2pp over FedAvg under triple adversity
4	Can fairness-aware aggregation solve class collapse?	4.A	FedMinMax +31.2pp on Breast Cancer
5	How resilient is FL to concept drift?	4.B	Replay effective; EWC counterproductive (-22.4 pp)
6	How do personalization methods compare?	4.C	Per-FedAvg fails (17.4%); Ditto/FedPer/APFL cluster
7	Can FL comply with GDPR right to erasure?	4.D	Gradient ascent: 5% cost, ≤ 0.2 pp drop
8	Which client selection strategy is optimal?	5.E	Oort: best accuracy (78.5%) AND fairness (0.938)
9	How vulnerable are gradients to reconstruction?	5.F	DP $\epsilon=10$: 100% attack failure
10	Do advantages generalize across clinical domains?	6.G	Ditto/HPFL advantage generalizes to all 6 domains
11	Is hierarchical FL viable for EHDS?	6.H	<0.5 pp penalty; viable for cross-border
12	How robust is FL to communication delays?	6.I	Staleness ≤ 3 tolerable (<2 pp)
13	Can FL learn multiple clinical tasks?	6.J	Dual-task: <2 pp primary degradation
14	Which Byzantine defense is most effective?	6.K	Multi-Krum, Trimmed Mean fully neutralize attacks
15	Can we fairly quantify hospital contributions?	6.L	Shapley values feasible for $K \leq 10$
16	Are FL probability outputs well-calibrated?	7.M	Overconfident; temperature scaling effective
17	Does conformal prediction maintain coverage?	7.N	95% coverage maintained under DP
18	Are feature rankings consistent across clients?	7.O	Top-3 consistent under IID; NonIID diverges
19	Do FL algorithms exhibit demographic biases?	7.P	Ditto reduces gaps by 60%; HPFL most equitable
20	Can FL detect and adapt to temporal drift?	7.Q	Replay-based adaptation most effective
21	How does privacy budget accumulate?	7.R	RDP enables $5\times$ more studies
22	What is the overhead of secure aggregation?	8.S	Sub-ms overhead; error $< 10^{-16}$
23	Which compression method is optimal?	8.T	SignSGD: $27.84\times$ compression, -0.17 pp
24	Is local or central DP more effective?	8.U	DP regularizes small datasets ($+15.5$ pp)
25	How does vertical FL perform?	8.V	3-party more robust to reduced overlap
26	Can FL handle 253K-sample population health data?	8.W	Scales without issues; Ditto rescues minority
27	How does FL scale to $K=50$ hospitals?	9.X	Tested on CDC (253K) and Cardiovascular
28	Are DP and compression compatible?	9.Y	SignSGD + DP combined evaluation
29	How many rounds are needed for convergence?	9.Z	Per-round tracking across 3 datasets
30	Can personalized FL handle extreme imbalance?	9.AA	Stroke (4.9%), CKD, Cirrhosis with F1/DEI
31	Is the Ditto advantage robust to heterogeneity?	9.AB	CDC $\alpha \in \{0.1, 0.25, 0.75, 1.0\}$
32	Does DP noise regularize against majority-class collapse?	10.AC	Collapse: 83% (no DP) $\rightarrow$ 0% (any DP)
33	Can weighted CE / focal loss eliminate class collapse?	10.AD	Weighted CE: 0.85% residual; Focal: 0%
34	Does increasing local epochs improve Ditto rescue?	10.AE	Cirrhosis F1: 0.36 (ep=3) $\rightarrow$ 0.78 (ep=10)
35	Can threshold tuning rescue collapsed models?	10.TH	Cirrhosis: yes (F1=0.65); Stroke: marginal
36	Do server-side adaptive optimizers improve FL?	Standalone	FedAdam/Yogi/Adagrad identical; SCAFFOLD worst
37	What is the centralized-FL gap across datasets?	Standalone	≤ 0.5 pp on PTB-XL and CV; FL near-optimal
38	Can FL scale to 253K population-health data?	Standalone	CDC scales to $K=20$; HPFL degrades under imbalance
39	Is VFL viable across clinical modalities?	Standalone	PTB-XL binarized 99.3%; CDC F1 $\approx$ 0 (imbalance)
40	Are algorithm rankings robust to 10-seed validation?	Standalone	FedAvg/Ditto cluster; HPFL lower on large datasets




Fig. S-6. Client participation matrix (50 rounds $\times$ 5 clients). Participation rates: IT 88%, DE 86%, FR 86%, ES 88%, NL 92%. The framework tolerates 10–15% dropout per round while maintaining convergence.




Fig. S-7. Gradient norm evolution per client over 50 rounds. All clients show decreasing trends indicating stable convergence. Clipping threshold $C=1.0$ bounds extreme values for DP compatibility.




Fig. S-8. Cumulative communication cost per round. Linear scaling with participating clients (3.5 KB/client/round). Total 50-round overhead: 875 KB for 5 clients—feasible even for bandwidth-constrained environments.

3) *Convergence at Different Heterogeneity Levels:* **Findings:** (1) At $\alpha=0.1$, SCAFFOLD achieves most stable convergence via variance reduction. (2) FedProx provides marginal improvement over FedAvg at $\alpha=0.5$ –1.0. (3) Adaptive methods (FedAdam, FedYogi) excel in near-IID but may oscillate under extreme heterogeneity. (4) FedAvg remains competitive in near-IID, suitable for homogeneous federations.

4) *Final Accuracy vs. Heterogeneity:*

5) *Complete 17-Algorithm Non-IID Evaluation (210 Experiments):* To complement the 7-algorithm summary in the main paper (Figure 3), we evaluated *all* 17 FL-EHDS algorithms across three Dirichlet heterogeneity levels ($\alpha \in \{0.1, 0.5, 1.0\}$)

on the Cardiovascular dataset, totaling 210 additional experiments (14 algorithms $\times$ 3 α levels $\times$ 5 seeds, plus FedAvg, Ditto, and HPFL from the existing 5-seed checkpoint). Table S-VII reports mean $\pm$ standard deviation.

Finding 1: Aggregation algorithm convergence. Ten of 17 algorithms cluster within 3 percentage points of FedAvg at $\alpha=0.5$ (range: 70.4–73.7%). This convergence phenomenon—where diverse algorithmic innovations (variance reduction, logit calibration, SAM-based optimization, server momentum) produce statistically indistinguishable results—confirms the main paper’s thesis: on compact tabular models (HealthcareMLP, 2-layer, 64 hidden units), the aggregation algorithm




Fig. S-9. Learning rate sensitivity analysis. $\eta=0.01$: slow convergence (53.8% at round 50). $\eta=0.1$: optimal (58.6%). $\eta=0.5$: instability with oscillations.




Fig. S-10. Batch size impact on convergence. Smaller batches (8–16) provide noisier gradients but faster initial progress. Batch size 32 balances gradient quality and computational efficiency.

is not the critical deployment decision.

Finding 2: HPFL is the sole outperformer. HPFL achieves $86.1 \pm 3.5\%$ at $\alpha=0.5$ vs. $73.7 \pm 2.2\%$ for FedAvg (+12.4pp), and the gap widens to +26.7pp under extreme non-IID ($\alpha=0.1$: 94.8% vs. 68.1%). This is statistically significant ($p < 0.001$, Cohen’s $d > 2.0$). The advantage stems from HPFL’s split-head architecture: shared feature extractors learn generalizable representations while personalized classification heads adapt to local label distributions.

Finding 3: Per-FedAvg (MAML) does not differentiate. Despite being a personalized algorithm, Per-FedAvg ($71.5 \pm 4.1\%$ at $\alpha=0.5$) clusters with global algorithms rather than with HPFL. This demonstrates that *generic* personalization (meta-learning initialization) is insufficient—only *architectural* personalization (split-head) produces meaningful improvement. Similarly, Ditto ($73.7 \pm 2.2\%$) matches FedAvg exactly, confirming that its personalization mechanism (fine-tuning a local copy) does not help on compact models where the global and local optima are close.

Finding 4: SCAFFOLD, FedSpeed, and FedAdagrad fail. Three algorithms exhibit instability:

- **SCAFFOLD** (44.0–62.2%): Control variate correction requires sufficient per-client data diversity to estimate accurate gradient corrections. With 5 heterogeneous clients, the correction terms amplify rather than reduce variance, producing high std (± 10 –11.5%).
- **FedSpeed** (47.6–56.7%): The gradient perturbation mechanism, designed to accelerate convergence, destabilizes training under non-IID conditions. Performance is consistently below FedAvg across all α levels.
- **FedAdagrad** (55.5–62.8%): Accumulated squared gradients grow rapidly under heterogeneous updates, shrinking effective learning rates to near-zero. At $\alpha=0.5$, FedAdagrad (55.5%) underperforms FedAvg by 18.2pp.

EHDS deployment implication. For federations of European hospitals using compact tabular models—the most common initial EHDS use case—practitioners should focus on person-




Fig. S-11. Per-client accuracy over training rounds. Variance reflects non-IID data: NL (older, higher-risk population) reaches 64% accuracy while FR (mid-range demographics) stabilizes at 55%.




Fig. S-12. Algorithm convergence across non-IID levels ($\alpha \in \{0.1, 0.5, 1.0, 10.0\}$). SCAFFOLD and adaptive methods show superior stability under extreme heterogeneity.




Fig. S-13. Final accuracy vs. Dirichlet α for five *global* aggregation algorithms (FedAvg, FedProx, SCAFFOLD, FedAdam, FedYogi). All converge to near-identical performance, confirming that aggregation strategy choice is irrelevant on compact tabular models. For the extended comparison including personalized methods (Ditto, HPFL), which exhibit dramatically different behavior, see the main paper.

TABLE S-VII

COMPLETE NON-IID EVALUATION: ALL 17 FL-EHDS ALGORITHMS ON CARDIOVASCULAR (ACCURACY %, MEAN $\pm$ STD OVER 4–5 SEEDS). SORTED BY $\alpha=0.5$ ACCURACY. BOLD: BEST PER COLUMN. $\dagger$: PERSONALIZED ALGORITHM.

Algorithm	$\alpha=0.1$	$\alpha=0.5$	$\alpha=1.0$
HPFL $\dagger$	94.8$\pm$5.6	86.1$\pm$3.5	81.2$\pm$3.1
FedAvg	68.1 $\pm$ 5.5	73.7 $\pm$ 2.2	72.7 $\pm$ 0.6
Ditto $\dagger$	67.8 $\pm$ 5.4	73.7 $\pm$ 2.2	72.7 $\pm$ 0.5
FedDyn	68.8 $\pm$ 4.8	72.3 $\pm$ 2.9	70.8 $\pm$ 0.9
FedYogi	64.4 $\pm$ 3.2	71.6 $\pm$ 1.1	71.3 $\pm$ 1.1
FedLC	65.1 $\pm$ 4.3	71.6 $\pm$ 4.0	68.0 $\pm$ 2.0
FedDecorr	64.6 $\pm$ 4.2	71.5 $\pm$ 4.1	68.0 $\pm$ 2.0
FedExP	64.7 $\pm$ 4.3	71.5 $\pm$ 4.1	68.0 $\pm$ 2.0
FedLESAM	64.6 $\pm$ 4.2	71.5 $\pm$ 4.1	68.0 $\pm$ 2.0
Per-FedAvg $\dagger$	64.7 $\pm$ 4.3	71.5 $\pm$ 4.1	67.9 $\pm$ 1.9
FedAdam	64.3 $\pm$ 3.2	71.5 $\pm$ 1.1	71.2 $\pm$ 1.2
FedSAM	64.5 $\pm$ 4.3	71.0 $\pm$ 4.2	67.2 $\pm$ 1.9
FedProx	65.4 $\pm$ 4.8	70.4 $\pm$ 4.0	67.3 $\pm$ 2.2
FedNova	62.0 $\pm$ 8.6	68.7 $\pm$ 3.2	65.3 $\pm$ 5.3
SCAFFOLD	44.0 $\pm$ 10.0	62.2 $\pm$ 11.5	52.1 $\pm$ 8.7
FedSpeed	47.6 $\pm$ 11.6	56.7 $\pm$ 10.7	50.9 $\pm$ 7.7
FedAdagrad	56.7 $\pm$ 13.1	55.5 $\pm$ 10.2	62.8 $\pm$ 2.2

Configuration: 5 clients, 25 rounds, 3 local epochs, lr=0.01, batch size=64, No-DP. $\mu=0.1$ for FedProx. Cardiovascular dataset (70K patients, binary classification). $n=4$ seeds at $\alpha=0.1$ (seed 456 excluded due to degenerate Dirichlet partition); $n=5$ seeds at $\alpha \in \{0.5, 1.0\}$.

alization architecture (HPFL) rather than searching for the optimal aggregation algorithm. The choice between FedAvg, FedProx, FedDyn, FedAdam, FedYogi, FedLC, FedLESAM, or FedExP is largely irrelevant (<3pp difference), while the choice to use split-head personalization yields +12–27pp improvement depending on data heterogeneity.

6) *Convergence Speed*: Figure S-14 compares how quickly algorithms reach target accuracy thresholds.

7) *Evidence-Based Algorithm Recommendation Matrix*: Table S-VIII maps seven realistic EHDS deployment scenarios to recommended and avoided algorithms, grounded in specific experimental evidence from this work. The Diagnostic Equity Index (DEI; Section F-H) is reported where confusion matrix data is available, providing a complementary criterion beyond aggregate accuracy.

Critical corrections from prior version: (1) SCAFFOLD is no longer recommended for heterogeneous deployments—our experiments demonstrate catastrophic failure on class-imbalanced tabular data (oscillating convergence; 11.2% on Diabetes, Table V). (2) Ditto+DP is strongly preferred over FedAvg+DP for privacy-critical deployments (89.2% vs. 52.3% at $\epsilon=1$ on PTB-XL). (3) Imaging recommendations now distinguish between Ditto (recommended for multiclass) and HPFL (avoid on all imaging tasks). (4) The DEI column provides diagnostic equity information absent from accuracy-only comparisons.

C. Extended Results from Main Paper

This section contains detailed results referenced inline in the main paper but moved here for space constraints. Table S-IX provides a complete decomposition of all 5,477+ experiments by category.

1) *Heart Disease and Diabetes Algorithm Comparison*: Table S-X presents the full algorithm comparison on two additional clinical datasets: Heart Disease UCI (4 natural hospitals) and Diabetes 130-US (5 Dirichlet-partitioned clients). These results confirm the personalization advantage observed on PTB-XL, Cardiovascular, and Breast Cancer in the main paper.

2) *Per-Hospital Heterogeneity Analysis*: Figure S-15 shows per-hospital accuracy variation on Heart Disease, where the four hospitals have naturally different patient populations.

D. Extended Tabular Experiment Results

This section presents comprehensive experimental results from 1,740+ federated learning experiments across three tabular healthcare datasets: PTB-XL ECG (21,799 records, 52 European recording sites, 5-class cardiac diagnosis), Cardiovascular Disease (70,000 patients, binary classification), and Breast Cancer Wisconsin (569 samples, binary classification). The evaluation comprises a baseline comparison (105 experiments, 7 algorithms $\times$ 3 datasets $\times$ 5 seeds), three sweep phases—heterogeneity (α variation, 560 experiments), client scaling (385 experiments), and learning rate sensitivity (180 experiments)—a differential privacy ablation (180 experiments, 4 ϵ levels $\times$ 3 algorithms $\times$ 3 datasets $\times$ 5 seeds),

an extended statistical validation (105 additional experiments with 5 new seeds for 10-seed Wilcoxon signed-rank tests), and an Article 71 opt-out impact analysis (225 experiments, 5 opt-out rates $\times$ 3 algorithms $\times$ 3 datasets $\times$ 5 seeds). All results are reproducible via the benchmark suite.

1) *Heterogeneity Impact*: Table S-XI shows accuracy as a function of Dirichlet α (lower α = more non-IID). The “Site” column reports natural site-based partitioning for PTB-XL.

Key finding: Personalized algorithms (Ditto, HPFL) exhibit a counter-intuitive pattern: accuracy *increases* under extreme non-IID ($\alpha=0.1$) on Cardiovascular (92.4% vs. 82.5% at $\alpha=0.5$) and Breast Cancer (88.3% vs. 79.1%). This occurs because personalized methods maintain local decision boundaries that become more specialized—and therefore more accurate—when client distributions are highly distinct. In contrast, baseline algorithms (FedAvg, FedExP, FedLESAM) degrade monotonically as α decreases, consistent with theoretical predictions. For PTB-XL, natural site-based partitioning produces the most realistic non-IID conditions and yields strong performance across all algorithms (89.7–96.5%).

2) *Client Scaling*: Table S-XII reports accuracy as a function of client count K .

Key finding: Baseline algorithms (FedAvg, FedExP, FedLESAM) degrade as client count increases on Cardiovascular (72.1% at $K=3 \rightarrow 69.4\%$ at $K=10$), because more clients means less data per client and higher aggregation noise. Personalized methods show the opposite trend: Ditto *improves* from 80.9% ($K=3$) to 85.0% ($K=10$), demonstrating that personalization enables each client to exploit its local specialization even with smaller data partitions. This has direct EHDS implications: as more Member States join a federation, personalized algorithms become increasingly advantageous.

3) *Learning Rate Sensitivity*: Table S-XIII evaluates robustness to learning rate selection for the top-3 algorithms.

Key finding: All algorithms are sensitive to learning rate at the low end ($lr=0.0005$), where convergence is incomplete within the configured round budget. HPFL and Ditto achieve near-optimal performance across a wider range (0.005–0.01), while FedAvg requires careful tuning. On Breast Cancer, HPFL reaches 89.7% accuracy at $lr=0.005$ —substantially higher than FedAvg’s best of 74.7% at $lr=0.01$ —confirming that personalized methods are more robust to hyperparameter selection on small datasets.

4) *Privacy-Utility Tradeoff (Differential Privacy)*: Table S-XIV quantifies accuracy under central differential privacy with varying privacy budget ϵ . We evaluate 3 representative algorithms: FedAvg (baseline), Ditto (best personalized), and HPFL (best fairness). All experiments use the Gaussian mechanism with clip norm $C=1.0$ and $\delta=10^{-5}$.

Key findings: (1) *Personalized methods are remarkably DP-robust*: At $\epsilon=10$, Ditto and HPFL lose <1pp on PTB-XL and <1.5pp on Cardiovascular, while FedAvg collapses to 52.3% on PTB-XL at $\epsilon=1$ (–39.6pp). This robustness arises because personalized local adaptation (Ditto’s fine-tuning, HPFL’s classifier heads) is unaffected by central DP noise injected during aggregation. (2) *DP noise as implicit regularization on small*




Fig. S-14. Convergence speed comparison. Left: rounds to 55% accuracy. Right: best accuracy in first 20 rounds. Adaptive methods converge faster but may plateau.

TABLE S-VIII

EHDS ALGORITHM RECOMMENDATION MATRIX. EVIDENCE-BASED MAPPING FROM DEPLOYMENT SCENARIOS TO RECOMMENDED AND AVOIDED ALGORITHMS. DEI = DIAGNOSTIC EQUITY INDEX (SECTION F-H). “AVOID” INDICATES ALGORITHMS THAT FAIL OR DEGRADE SIGNIFICANTLY IN THE SPECIFIED SCENARIO.

EHDS Scenario	Modal.	Recommended	Avoid	DEI	Key Evidence
Cross-border EHR, heterogeneous sites	Tabular	Ditto, HPFL	SCAFFOLD ^a , FedNova ^a	0.740 [†]	Table VII: Ditto +11.2pp, HPFL +10.2pp vs. FedAvg (10 seeds)
Small-cohort rare disease (<1K samples)	Tabular	HPFL	FedAvg ^b	0.740 [†]	Table S-ConfusionBC: HPFL avoids collapse 10/10 seeds vs. 2/10 FedAvg
Multiclass radiology (≥ 3 classes)	Imaging	Ditto	HPFL (neutral)	0.435	Table XXV: Ditto +23.5pp Brain Tumor ($d=4.05$, 3 seeds; $p=0.004$, 10 seeds)
Binary radiology	Imaging	FedAvg	HPFL (−18.2pp)	0.528 [§]	Tables XX, XXV: FedAvg 87.3% vs. HPFL 69.1% Chest X-ray
High privacy ($\epsilon \leq 5$)	Any	Ditto	FedAvg ^c	—	Table XXII: Ditto 89.2% vs. FedAvg 52.3% at $\epsilon=1$
Large federation ($K \geq 50$)	Any	Ditto, HPFL	SCAFFOLD	—	Tables XVII–XVIII: Ditto −0.8pp at $K=100$ vs. FedAvg −4.7pp
Homogeneous / IID data	Any	FedAvg	(all equivalent)	—	PTB-XL: all algorithms within 0.6pp (91.9–92.5%)

^aSCAFFOLD oscillates on class-imbalanced tabular data; FedNova amplifies noise from divergent objectives (11.2%, 13.0% on Diabetes). ^bFedAvg/Ditto exhibit single-class collapse on 8/10 seeds (Table S-ConfusionBC). ^cFedAvg collapses from 91.9% to 52.3% at $\epsilon=1$ (−39.6pp). [†]Breast Cancer HPFL DEI. [‡]Brain Tumor Ditto DEI (3 seeds). [§]Chest X-ray FedAvg DEI. DEI values reported only where confusion matrix analysis is available.




Fig. S-15. Per-hospital accuracy of the *global* model on Heart Disease UCI. Ditto’s overall advantage (11.2pp over 10 seeds; main paper Table VII) comes from its *personalized local* models, which are separately fine-tuned per hospital; the shared global model shows similar cross-hospital performance to FedAvg. The Hungarian hospital, with the smallest cohort ($n=294$), shows the largest performance gap—a realistic EHDS scenario where smaller national datasets benefit most from federation.

TABLE S-IX
EXPERIMENT DECOMPOSITION MAP

Category	Count	Supplementary Ref.
Algorithm comparison (3 datasets)	105	Table III (main)
Heart Disease & Diabetes	50	Table S-HD-Diabetes
Extended 10-seed validation	140	Table S-10Seed
Imaging (3 datasets, ResNet-18)	108	Section S-Imaging
Server-side optimizers	135	Table S-ServerOpt
DP privacy-utility	180	Table S-DP
Cross-border heterogeneous DP	45	Table S-CrossBorder
Non-IID heterogeneity	420	Table S-NonIID
Non-IID $\times$ DP compound	126	Table S-NonIID-DP
Compound stress	48	Table S-CompoundStress
Byzantine + DP defenses	1,620	Section S-Byzantine
Article 71 opt-out (static)	225	Table S-Optout
Article 71 opt-out (dynamic)	90	Table S-DynOptout
Partial participation	135	Table S-Participation
Data efficiency	225	Table S-DataEfficiency
DeepMLP validation	70	Table S-DeepMLP
Confusion matrix (10-seed)	70	Table S-ConfusionBC
CDC Diabetes (253K)	35	Section S-CDC
Vertical FL	15	Section S-VerticalFL
Centralized baselines	15	Table S-CentralBaseline
Thesis robustness (3 phases)	516	Section S-ThesisRobustness
Top-K sparsification	180	Section S-TopK
Non-IID imaging	648	Section S-NonIID-Imaging
Total	5,477+	

Counts are approximate per category; some experiments contribute to multiple categories (e.g., Non-IID $\times$ DP). The “+” accounts for hyperparameter sweeps and cross-category overlaps.

TABLE S-X
FL ALGORITHM COMPARISON ON HEART DISEASE AND DIABETES

Algo.	Heart Disease (4 hosp.)			Diabetes (5 hosp.)		
	Acc.	F1	AUC	Acc.	F1	AUC
FedAvg	62.5 $\pm$ 8.0	.736 $\pm$.06	.834 $\pm$.03	68.1 $\pm$ 4.2	.259 $\pm$.01	.643 $\pm$.00
FedProx	61.7 $\pm$ 8.0	.732 $\pm$.05	.834 $\pm$.03	71.0 $\pm$ 6.3	.254 $\pm$.01	.638 $\pm$.00
SCAFFOLD	66.3 $\pm$ 5.1	.667 $\pm$.02	.791 $\pm$.05	11.2 $\pm$ 0.0	.201 $\pm$.00	.514 $\pm$.00
FedNova	56.4 $\pm$ 5.4	.711 $\pm$.04	.831 $\pm$.03	13.0 $\pm$ 0.9	.203 $\pm$.00	.636 $\pm$.00
Ditto	75.1$\pm$2.0	.761$\pm$.03	.826$\pm$.01	71.7$\pm$0.2	.262$\pm$.00	.643$\pm$.00

20 rounds, 3 local epochs. Heart Disease: natural non-IID (4 international hospitals: Cleveland, Hungarian, Swiss, VA Long Beach). Diabetes: Dirichlet $\alpha=0.5$. Mean $\pm$ std over 3 seeds. Ditto achieves 75.1% vs. FedAvg 62.5% (12.6pp gap, 3 seeds) on Heart Disease and 71.7% on Diabetes. SCAFFOLD and FedNova exhibit known failure modes on Diabetes (11.2% and 13.0%). Extended 10-seed paradigm comparison (including HPFL and Local-Only baselines) in main paper Table VII: Ditto 76.0%, FedAvg 64.8% (11.2pp gap).

data: On Breast Cancer (569 samples), FedAvg with $\varepsilon=5$ achieves 78.7%—*higher* than the no-DP baseline of 52.3% (+26.4pp). HPFL at $\varepsilon=1$ reaches 86.9% vs. 74.1% without DP (+12.8pp). This counter-intuitive result is consistent with the theoretical analysis of Wei et al. [36], who show that DP noise interacts with model convergence in non-trivial ways, and with the broader machine learning literature on noise injection as regularization (e.g., dropout, label smoothing). On this small dataset, the Gaussian noise added during DP-SGD prevents the global model from overfitting to the majority class, act-

ing as a stochastic regularizer that improves generalization. This effect is strongest at moderate ε (5–10) and diminishes at very strict $\varepsilon=1$ where noise overwhelms the signal on larger datasets. *Practical implication*: for small-cohort EHDS studies (rare diseases, specialist registries), moderate DP may simultaneously enhance privacy *and* model quality—a finding that should inform data permit conditions. (3) *Privacy is nearly free* at $\varepsilon=10$: Across PTB-XL and Cardiovascular, all algorithms recover to within 2pp of their no-DP baselines at $\varepsilon=10$ —a practical privacy level that satisfies EHDS Article 50

TABLE S-XI
IMPACT OF DATA HETEROGENEITY (α) ON FL ACCURACY (%). LOWER α = MORE NON-IID. MEAN OVER 5 SEEDS. PX = PTB-XL, CV = CARDIOVASCULAR, BC = BREAST CANCER.

DS	Algorithm	IID	$\alpha=0.1$	$\alpha=0.3$	$\alpha=0.5$	$\alpha=1.0$	Site
PX	FedAvg	90.8	89.8	91.2	90.8	91.7	91.9
	FedProx	91.0	89.7	91.3	90.7	91.7	91.6
	Ditto	91.7	96.5	96.5	95.4	93.9	91.8
	FedLC	90.7	89.7	91.2	90.4	91.4	91.9
	FedExP	90.8	89.8	91.2	90.8	91.7	92.0
	FedLESAM	90.8	89.8	91.2	90.8	91.7	91.9
	HPFL	91.3	96.3	96.2	95.2	94.0	92.5
CV	FedAvg	72.7	61.2	65.7	71.1	72.6	—
	FedProx	72.7	62.5	65.9	71.6	72.5	—
	Ditto	72.6	92.4	90.7	82.5	81.5	—
	FedLC	72.7	61.9	66.0	71.1	72.6	—
	FedExP	72.7	61.2	65.7	71.1	72.6	—
	FedLESAM	72.7	61.2	65.7	71.1	72.6	—
	HPFL	72.7	92.4	90.7	82.3	81.4	—
BC	FedAvg	47.1	62.1	57.3	52.3	51.5	—
	FedProx	47.1	62.1	57.3	52.3	51.5	—
	Ditto	51.5	88.3	84.8	79.1	72.4	—
	FedLC	47.3	63.2	57.2	52.1	51.5	—
	FedExP	47.1	62.1	57.3	52.3	51.5	—
	FedLESAM	47.1	62.1	57.3	52.3	51.5	—
	HPFL	51.9	88.3	84.8	74.1	66.9	—

requirements with minimal utility cost.

5) *Fairness Analysis:* Table S-XV provides per-client fairness metrics across all datasets.

Key finding: HPFL uniquely improves fairness on Breast Cancer (Jain 0.867 vs. 0.608 for all other algorithms, Gini reduction from 0.405 to 0.159). This occurs because HPFL’s personalized classifier heads enable each client to specialize, reducing the performance gap between clients with different class distributions (Gap reduced from 71.5% to 47.6%). On PTB-XL, all algorithms achieve near-perfect fairness (Jain ≥ 0.999) due to the large dataset size providing sufficient per-client samples.

6) *Statistical Significance:* Table S-XVI reports Wilcoxon signed-rank test p -values comparing each algorithm against FedAvg, computed over 10 seeds (original seeds: 42, 123, 456, 789, 999; additional seeds: 7, 31, 137, 577, 1337). With $n=10$ paired observations, the minimum achievable two-sided p -value is $2/2^{10} = 0.00195$. Table S-XVII supplements p -values with effect sizes (rank-biserial correlation r and mean accuracy difference $\bar{\Delta}$) for the two statistically significant algorithms.

Key findings: With 10 seeds, the statistical picture sharpens decisively. *First, HPFL significantly outperforms FedAvg on all three datasets* ($p = 0.004, 0.002, 0.031$; all < 0.05), with improvements of +0.9pp on PTB-XL, +11.8pp on Cardiovascular, and +16.5pp on Breast Cancer—the only algorithm achieving significance across all datasets. All effect sizes are large ($r \geq 0.93$; Table S-XVII). *Second, Ditto significantly outperforms FedAvg on Cardiovascular and Breast Cancer* ($p = 0.002, 0.016$), with large effect sizes ($r = 1.00, 1.00$; $\bar{\Delta} = +11.9\text{pp}, +19.0\text{pp}$), though not on PTB-XL ($p = 0.49, \Delta = +0.2\text{pp}$) where FedAvg already

achieves 91.7%. *Third, five algorithms collapse to FedAvg-equivalent performance:* FedLESAM produces identical results on all datasets; FedExP and FedProx are identical on 2/3 datasets. This confirms that only methods maintaining separate local models (Ditto, HPFL) differentiate on the near-convex tabular MLP landscape. *Fourth, the pooled analysis* across all 30 paired observations yields $p < 0.001$ for both Ditto ($r = 0.89$) and HPFL ($r = 0.95$), providing strong evidence that personalized FL methods are superior for EHDS tabular analytics.

7) *Article 71 Opt-Out Impact:* EHDS Article 71 grants citizens the right to opt out of secondary use of their health data. To quantify the impact on model quality, we simulate opt-out by randomly removing 5%, 10%, 20%, and 30% of training samples per client (test data unchanged). This mirrors real-world scenarios where patients exercise their opt-out right, reducing the available training data at each institution. We evaluate FedAvg (baseline), Ditto, and HPFL (the two statistically significant personalized methods) across all three datasets with 5 seeds (225 experiments total).

Key findings. *First*, on datasets of adequate size (PTB-XL: 21.8K samples; Cardiovascular: 70K samples), Article 71 opt-out up to 30% has **negligible impact on model quality** ($< 1\text{pp}$ accuracy drop across all algorithms). This is the central policy-relevant result: citizen privacy rights under EHDS are compatible with high-quality federated analytics.

Second, HPFL is the most opt-out-resilient algorithm: on PTB-XL, HPFL maintains 92.7% accuracy from 0% to 30% opt-out ($\Delta_{\max} = -0.1\text{pp}$). Its personalized classifier heads adapt to the reduced data volume at each client without

TABLE S-XII
IMPACT OF CLIENT COUNT ON FL ACCURACY (%). MEAN OVER 5 SEEDS. PX = PTB-XL, CV = CARDIOVASCULAR, BC = BREAST CANCER.

PTB-XL ECG (5-class)					Cardiovascular (binary)				
Algorithm	K=3	K=5	K=10	K=20	Algorithm	K=3	K=5	K=8	K=10
FedAvg	91.7	91.9	91.5	91.5	FedAvg	72.1	71.1	70.5	69.4
FedProx	91.7	91.6	91.5	91.5	FedProx	72.0	71.6	70.7	69.6
Ditto	92.5	91.8	91.5	91.5	Ditto	80.9	82.5	83.3	85.0
FedLC	91.8	91.9	91.8	91.8	FedLC	72.1	71.1	70.7	68.9
FedExP	91.7	92.0	92.2	92.2	FedExP	72.1	71.1	70.5	69.4
FedLESAM	91.7	91.9	91.5	91.5	FedLESAM	72.1	71.1	70.5	69.4
HPFL	92.6	92.5	92.5	92.5	HPFL	80.9	82.3	83.3	84.7

Breast Cancer (binary)			
Algorithm	K=2	K=3	K=5
FedAvg	61.6	52.3	59.1
FedProx	61.6	52.3	59.1
Ditto	84.2	79.1	78.4
FedLC	56.6	52.1	50.1
FedExP	61.6	52.3	59.1
FedLESAM	61.6	52.3	59.1
HPFL	84.2	74.1	78.4

TABLE S-XIII
LEARNING RATE SENSITIVITY ANALYSIS. ACCURACY (%) MEAN OVER 5 SEEDS. TOP-3 ALGORITHMS: HPFL, DITTO, FEDAVG.

DS	Algo	0.0005	0.001	0.005	0.01
PX	Ditto	58.1	84.1	91.8	92.6
	FedAvg	80.3	90.4	91.9	92.2
	HPFL	63.2	82.4	92.5	92.6
CV	Ditto	76.2	77.4	82.2	82.5
	FedAvg	55.1	60.6	70.9	71.1
	HPFL	76.1	77.1	82.2	82.3
BC	Ditto	64.0	79.1	89.5	89.5
	FedAvg	52.1	52.3	64.0	74.7
	HPFL	63.8	74.1	89.7	89.5

degrading global performance.

Third, small datasets are vulnerable: Breast Cancer (569 samples, 3 clients) shows instability under HPFL at $\geq 10\%$ opt-out (-10.4pp), driven by individual seed sensitivity when per-client training sets shrink below ~ 130 samples. This highlights a practical EHDS consideration: opt-out impact assessments should be mandatory for small-cohort studies, and minimum sample thresholds should be specified in data permit conditions. *Policy recommendation:* EHDS data permits authorizing personalized FL methods with per-client components (e.g., HPFL’s local classifier heads, Per-FedAvg’s meta-learned initialization) should specify minimum per-client sample requirements—our results suggest ≥ 200 samples per client as a conservative threshold for stable personalization. When expected opt-out rates bring per-client data below this threshold, data permits should mandate global aggregation methods (FedAvg, Ditto) which demonstrate robustness to data reduction even on small datasets (Table S-XVIII: $\leq 0.6\text{pp}$ change at 30% opt-out on Breast Cancer).

Fourth, the stability of FedAvg and Ditto on Breast Cancer

despite opt-out (no degradation) suggests that simpler models are more robust to data reduction on small datasets, while HPFL’s additional parameterization (per-client classifier heads) requires sufficient data to avoid overfitting.

8) *Experimental Configuration Summary:* **Reproducibility:** All experiments are fully reproducible via:

```
cd fl-ehds-framework
# Baseline (105 exp, ~45 min)
python -m benchmarks.run_tabular_optimized
# Multi-phase sweep (1125 exp, ~4.5h)
python -m benchmarks.run_tabular_sweep \
  --phase all
# DP ablation (180 exp, ~1.5h)
python -m benchmarks.run_tabular_dp
# 10-seed significance (105 exp, ~40 min)
python -m benchmarks.run_tabular_seeds10
# Article 71 opt-out (225 exp, ~1.5h)
python -m benchmarks.run_tabular_optout
# Deep MLP (70 exp, ~1.5h)
python -m benchmarks.run_tabular_deep_mlp
# Extended analysis (tables + plots)
python -m benchmarks.analyze_tabular_extended
```

Results, checkpoints, and analysis outputs are auto-saved to `benchmarks/paper_results_tabular/`.

TABLE S-XIV

PRIVACY-UTILITY TRADEOFF: ACCURACY (%) UNDER CENTRAL DP WITH VARYING ϵ . GAUSSIAN MECHANISM, $C=1.0$, $\delta=10^{-5}$. MEAN $\pm$ STD OVER 5 SEEDS. NO-DP BASELINES FROM MAIN PAPER TABLE V. Δ = DROP FROM NO-DP.

DS	Algo	No-DP	$\epsilon=1$	$\epsilon=5$	$\epsilon=10$	$\epsilon=50$
PX	FedAvg	91.9 $\pm$ 0.5	52.3 $\pm$ 13.3	84.2 $\pm$ 4.3	92.4 $\pm$ 0.3	92.4 $\pm$ 0.3
	Ditto	91.8 $\pm$ 0.3	89.2 $\pm$ 0.4	90.9 $\pm$ 0.6	91.6 $\pm$ 0.5	91.9 $\pm$ 0.3
	HPFL	92.5 $\pm$ 0.3	87.1 $\pm$ 2.8	90.6 $\pm$ 2.2	92.4 $\pm$ 0.4	92.7 $\pm$ 0.2
CV	FedAvg	71.1 $\pm$ 1.8	54.7 $\pm$ 3.2	59.8 $\pm$ 4.7	69.1 $\pm$ 3.7	71.0 $\pm$ 2.4
	Ditto	82.5 $\pm$ 4.7	76.9 $\pm$ 7.5	80.4 $\pm$ 5.8	81.9 $\pm$ 4.9	82.5 $\pm$ 4.8
	HPFL	82.3 $\pm$ 4.5	74.7 $\pm$ 9.1	79.3 $\pm$ 6.0	81.2 $\pm$ 5.2	82.2 $\pm$ 4.6
BC	FedAvg	52.3 $\pm$ 17.9	65.2 $\pm$ 8.1	78.7 $\pm$ 3.6	78.2 $\pm$ 8.3	72.0 $\pm$ 10.3
	Ditto	79.1 $\pm$ 12.5	73.8 $\pm$ 20.6	74.0 $\pm$ 20.7	74.0 $\pm$ 20.7	74.0 $\pm$ 20.7
	HPFL	74.1 $\pm$ 20.9	86.9 $\pm$ 13.8	84.9 $\pm$ 13.0	80.0 $\pm$ 17.6	80.7 $\pm$ 21.8

TABLE S-XV

PER-CLIENT FAIRNESS ANALYSIS. GAP = MAX-MIN CLIENT ACCURACY. LOWER GAP AND GINI INDICATE FAIRER DISTRIBUTION. MEAN OVER 5 SEEDS.

DS	Algo	Jain	Gini	Gap (%)	Std (%)
PX	FedAvg	0.999	0.013	6.5	2.3
	FedProx	0.999	0.012	5.6	2.0
	Ditto	0.999	0.014	6.7	2.4
	FedLC	0.999	0.013	6.3	2.3
	FedExp	0.999	0.011	5.7	2.0
	FedLESAM	0.999	0.013	6.5	2.3
	HPFL	0.999	0.018	9.1	3.1
CV	FedAvg	0.981	0.063	22.5	8.6
	FedProx	0.986	0.056	19.9	7.5
	Ditto	0.980	0.065	23.0	8.7
	FedLC	0.982	0.062	21.9	8.3
	FedExp	0.981	0.063	22.5	8.6
	FedLESAM	0.981	0.063	22.5	8.6
	HPFL	0.984	0.065	26.0	10.8
BC	FedAvg	0.608	0.405	71.5	32.3
	FedProx	0.608	0.405	71.5	32.3
	Ditto	0.606	0.411	73.2	33.0
	FedLC	0.606	0.413	74.2	33.4
	FedExp	0.608	0.405	71.5	32.3
	FedLESAM	0.608	0.405	71.5	32.3
	HPFL	0.867	0.159	47.6	22.1

9) *Deep MLP Algorithm Differentiation*: The tabular experiments in the main paper use a compact HealthcareMLP (~ 10 K parameters, 2 hidden layers [64, 32]), which produces a nearly convex loss landscape where five of seven algorithms collapse to FedAvg-equivalent performance. To test whether a deeper, more non-convex model breaks this collapse, we evaluate a DeepHealthcareMLP (~ 110 K parameters, 4 hidden layers [256, 256, 128, 64], ReLU, dropout 0.3—no BatchNorm to avoid FL aggregation issues, no residual connections to preserve non-convexity). This represents a $38\times$ parameter

increase while maintaining the same training configuration.

Key finding: Even with a $38\times$ parameter increase, **the algorithm collapse persists**. On both PTB-XL and Cardiovascular, FedAvg, FedProx, FedLC, FedExp, and FedLESAM converge to statistically identical accuracies (92.5% and 71.2%, respectively). Only Ditto and HPFL differentiate, with modest improvements on Cardiovascular (+2.5pp each vs. shallow). This confirms that the near-convexity driving algorithm collapse is a property of *tabular healthcare data* (low-dimensional features, clear class boundaries), not merely

TABLE S-XVI
WILCOXON SIGNED-RANK p -VALUES VS. FEDAVG (10 SEEDS). †: $p < 0.05$. POOLED: TEST ACROSS ALL 30 PAIRED OBSERVATIONS (3 DATASETS $\times$ 10 SEEDS).

vs FedAvg	PX	CV	BC	Pooled
FedProx	0.508	0.770	$\equiv$	0.538
Ditto	0.492	0.002 †	0.016 †	<0.001 †
FedLC	0.344	0.496	1.000	0.881
FedExp	0.922	$\equiv$	$\equiv$	0.878
FedLESAM	$\equiv$	$\equiv$	$\equiv$	—
HPFL	0.004 †	0.002 †	0.031 †	<0.001 †

$\equiv$: identical results to FedAvg (algorithm collapse on near-convex loss).

TABLE S-XVII
EFFECT SIZES FOR STATISTICALLY SIGNIFICANT ALGORITHMS VS. FEDAVG (10 SEEDS). r : RANK-BISERIAL CORRELATION ($|r| > 0.5$: LARGE EFFECT). Δ : MEAN ACCURACY IMPROVEMENT (PP).

Algorithm	Dataset	p	r	Δ (pp)
Ditto	Cardiovascular	0.002	1.00	+11.9
	Breast Cancer	0.016	1.00	+19.0
	Pooled (30 obs.)	<0.001	0.89	+10.4
HPFL	PTB-XL	0.004	0.96	+0.9
	Cardiovascular	0.002	1.00	+11.8
	Breast Cancer	0.031	0.93	+16.5
	Pooled (30 obs.)	<0.001	0.95	+9.8

TABLE S-XVIII
ARTICLE 71 OPT-OUT IMPACT: ACCURACY (%) AT DIFFERENT OPT-OUT RATES. Δ : MAXIMUM DROP RELATIVE TO 0% BASELINE. MEAN $\pm$ STD OVER 5 SEEDS. NOTE: 0% BASELINES WERE RUN AS SEPARATE EXPERIMENTAL BATCHES; MINOR DIFFERENCES FROM MAIN PAPER TABLE V (E.G., HPFL PTB-XL 92.8 VS. 92.5) REFLECT SEED SET VARIABILITY.

DS	Algo	0%	5%	10%	20%	30%	$\Delta_{\max}$
PX	FedAvg	91.9 $\pm$ 0.5	91.6 $\pm$ 0.7	91.6 $\pm$ 0.7	91.7 $\pm$ 0.7	91.8 $\pm$ 0.7	−0.3
	Ditto	91.9 $\pm$ 0.4	91.9 $\pm$ 0.3	91.9 $\pm$ 0.3	91.8 $\pm$ 0.4	91.1 $\pm$ 0.8	−0.8
	HPFL	92.8 $\pm$ 0.4	92.7 $\pm$ 0.4	92.7 $\pm$ 0.4	92.7 $\pm$ 0.4	92.7 $\pm$ 0.4	−0.1
CV	FedAvg	70.6 $\pm$ 1.6	70.5 $\pm$ 1.9	70.5 $\pm$ 1.6	70.4 $\pm$ 1.8	70.6 $\pm$ 1.6	−0.2
	Ditto	84.9 $\pm$ 4.1	84.9 $\pm$ 4.1	84.9 $\pm$ 4.1	84.8 $\pm$ 3.9	84.7 $\pm$ 4.0	−0.2
	HPFL	84.8 $\pm$ 3.8	84.7 $\pm$ 3.9	84.8 $\pm$ 3.8	84.7 $\pm$ 3.8	84.7 $\pm$ 3.8	−0.1
BC	FedAvg	52.8 $\pm$ 17.6	52.8 $\pm$ 17.6	53.0 $\pm$ 17.9	52.8 $\pm$ 17.6	52.8 $\pm$ 17.6	−0.0
	Ditto	83.5 $\pm$ 10.3	84.1 $\pm$ 10.7	84.1 $\pm$ 10.7	84.1 $\pm$ 10.7	84.1 $\pm$ 10.7	+0.6
	HPFL	84.1 $\pm$ 10.7	83.9 $\pm$ 10.5	73.5 $\pm$ 19.8	73.7 $\pm$ 19.9	73.7 $\pm$ 19.9	−10.4

of model capacity. For EHDS deployment, this strengthens our main finding: on tabular clinical models, the critical design choice is **personalization architecture** (Ditto, HPFL) rather than aggregation strategy, regardless of model depth.

10) Local Epochs Sweep: To test whether increased client drift breaks the algorithm collapse, we sweep local epochs $E \in \{1, 5, 10, 20\}$ on Cardiovascular (7 algorithms $\times$ 4 epochs $\times$ 5 seeds = 140 experiments). The hypothesis is that with $E=10-20$, client models diverge sufficiently from the global model that variance-reduction (FedProx proximal term), adaptive (FedExp, FedLESAM), and logit-calibration (FedLC) strategies should differentiate from FedAvg.

Key finding: The algorithm collapse is robust to local epochs. Even at $E=20$ ($20\times$ more local computation per round), the five non-personalized algorithms remain within 1.0pp of each other, and the personalization gap stays at

~ 12 pp. FedProx shows marginal differentiation at $E=10$ (+1.4pp over FedAvg) due to its proximal term limiting client drift, but remains firmly in the collapsed group. Combined with the deep MLP results above, this establishes that algorithm collapse on tabular healthcare data is **robust to model capacity, local epochs, and client count** (Section E-D.11)—it is a fundamental property of the low-dimensional, near-convex optimization landscape, not a training artifact.

11) Scalability Analysis: A key concern for EHDS deployment is whether FL algorithms maintain performance as the number of participating hospitals grows from laboratory settings ($K=5-20$) to realistic cross-border scales ($K=50-100$ across 27 Member States). We conduct a systematic scalability sweep on two datasets: Cardiovascular ($K \in \{50, 100\}$, Dirichlet $\alpha=0.5$) and PTB-XL ($K=12$ site-based using all available recording sites with ≥ 30 samples, and $K=30$ Dirichlet). Each

TABLE S-XIX
PER-DATASET OPTIMIZED CONFIGURATION USED FOR BASELINE COMPARISON (MAIN PAPER TABLE V).

Dataset	lr	bs	rounds	K	α	Partition
PTB-XL	0.005	64	30	5	—	Site-based
Cardio.	0.01	64	25	5	0.5	Dirichlet
Breast C.	0.001	16	40	3	0.5	Dirichlet

TABLE S-XX
DEEP MLP (~ 110 K PARAMS) VS. SHALLOW MLP (~ 2.9 K– 10 K PARAMS). ACCURACY (%), MEAN $\pm$ STD OVER 5 SEEDS.

Algorithm	PTB-XL ECG		Cardiovascular	
	Shallow	Deep	Shallow	Deep
FedAvg	91.9 $\pm$ 0.5	92.5 $\pm$ 0.3	71.1 $\pm$ 1.8	71.2 $\pm$ 1.0
FedProx	91.6 $\pm$ 0.7	92.5 $\pm$ 0.3	71.5 $\pm$ 1.2	71.2 $\pm$ 1.0
Ditto	91.8 $\pm$ 0.3	92.2 $\pm$ 0.7	82.5 $\pm$ 4.7	85.0$\pm$4.1
FedLC	91.9 $\pm$ 0.5	92.5 $\pm$ 0.3	71.1 $\pm$ 1.6	71.2 $\pm$ 1.1
FedExp	92.0 $\pm$ 0.2	92.2 $\pm$ 0.4	71.1 $\pm$ 1.8	71.2 $\pm$ 1.0
FedLESAM	91.9 $\pm$ 0.5	92.5 $\pm$ 0.3	71.1 $\pm$ 1.8	71.2 $\pm$ 1.0
HPFL	92.5$\pm$0.3	92.7$\pm$0.2	82.3 $\pm$ 4.5	84.8 $\pm$ 3.9

Shallow: HealthcareMLP [64, 32], ~ 2.9 K– 10 K params (varies by input features). Deep: DeepHealthcareMLP [256, 256, 128, 64], ~ 110 K params. Same per-dataset hyperparameters (Table S-XIX). Breast Cancer omitted: 569 samples insufficient for stable ~ 110 K-parameter training.

TABLE S-XXI
LOCAL EPOCHS SWEEP ON CARDIOVASCULAR. ACCURACY (%), MEAN $\pm$ STD OVER 5 SEEDS.

Algorithm	E=1	E=5	E=10	E=20
FedAvg	70.6 $\pm$ 2.2	70.5 $\pm$ 2.6	70.1 $\pm$ 3.6	70.3 $\pm$ 3.6
FedProx	70.3 $\pm$ 3.0	71.0 $\pm$ 2.0	71.5 $\pm$ 1.5	71.3 $\pm$ 1.9
FedLC	70.5 $\pm$ 2.6	70.7 $\pm$ 2.3	70.3 $\pm$ 3.5	70.4 $\pm$ 3.4
FedExp	70.6 $\pm$ 2.2	70.5 $\pm$ 2.6	70.1 $\pm$ 3.6	70.3 $\pm$ 3.6
FedLESAM	70.6 $\pm$ 2.2	70.5 $\pm$ 2.6	70.1 $\pm$ 3.6	70.3 $\pm$ 3.6
Ditto	81.9$\pm$4.7	82.5$\pm$4.7	82.6$\pm$4.8	82.8$\pm$4.8
HPFL	81.9$\pm$4.7	82.4$\pm$4.8	82.4$\pm$4.8	82.6$\pm$4.8
<i>Collapse gap</i>	<i>11.4pp</i>	<i>11.9pp</i>	<i>12.1pp</i>	<i>12.1pp</i>
<i>Collapsed spread</i>	<i>0.3pp</i>	<i>0.5pp</i>	<i>1.4pp</i>	<i>1.0pp</i>

Collapse gap = mean(Ditto, HPFL) – mean(others). Collapsed spread = max – min within the 5 non-personalized algorithms. Same hyperparameters as baseline (Table S-XIX).

configuration runs 7 algorithms $\times$ 3 seeds.

Key findings: (1) **Personalization is scale-robust:** On Cardiovascular, Ditto and HPFL degrade by only -0.8 pp from $K=5$ to $K=100$ (100 hospitals), while non-personalized algorithms lose 4.7 – 5.5 pp. The *personalization gap widens* at scale: from 11.7 pp ($K=5$) to 14.4 pp ($K=50$) to 15.9 pp ($K=100$), making personalization *more* important as the federation grows. (2) **Natural partitioning is benign:** PTB-XL’s site-based split ($K=12$, all sites with ≥ 30 samples) preserves near-baseline performance for 6/7 algorithms ($\sim 91.5\%$), with HPFL achieving 92.4% —identical to $K=5$. FedExp is the sole exception, exhibiting catastrophic divergence (40.2%) under site-level heterogeneity. (3) **Artificial high-K partitioning is destructive:** Dirichlet $K=30$ on PTB-XL (~ 21 K samples $\div$ 30 clients = ~ 700 /client) causes severe degradation for non-personalized algorithms (-48 pp), but Ditto maintains 76.3% —

the best by 26 pp over the next algorithm. (4) These results reinforce the ≥ 200 samples/client threshold (Section E-D.7): at $K=100$ on Cardiovascular ($\mu=700$ samples/client), personalized algorithms remain above 81% , but the Dirichlet tail produces some clients with <10 samples, explaining the degradation of non-personalized methods.

EHDS implication: For realistic cross-border deployments with $K=50$ – 100 hospitals, **personalized FL algorithms (Ditto, HPFL) are essential**—they provide 16 pp higher accuracy than standard FedAvg at $K=100$ and are $6\times$ more robust to scaling (-0.8 pp vs. -4.7 pp degradation). Natural hospital-based partitions (PTB-XL sites) are far more benign than synthetic Dirichlet splits, suggesting that real EHDS data distributions may be more favorable than worst-case laboratory settings.

TABLE S-XXII
CARDIOVASCULAR SCALABILITY: ACCURACY (%) VS. NUMBER OF CLIENTS. MEAN $\pm$ STD OVER 3 SEEDS. Δ DENOTES CHANGE FROM K=5 BASELINE (5 SEEDS).

Algorithm	K=5	K=50	Δ	K=100	Δ
FedAvg	70.5 $\pm$ 2.4	67.8 $\pm$ 0.2	-2.7	65.8 $\pm$ 0.8	-4.7
FedProx	71.0 $\pm$ 1.8	67.6 $\pm$ 0.2	-3.4	65.6 $\pm$ 0.9	-5.4
FedLC	70.6 $\pm$ 2.2	67.5 $\pm$ 0.3	-3.1	65.1 $\pm$ 0.5	-5.5
FedExp	70.5 $\pm$ 2.4	67.8 $\pm$ 0.2	-2.7	65.8 $\pm$ 0.8	-4.7
FedLESAM	70.5 $\pm$ 2.4	67.8 $\pm$ 0.2	-2.7	65.8 $\pm$ 0.8	-4.7
Ditto	82.3 $\pm$ 4.8	82.7 $\pm$ 1.2	+0.4	81.6 $\pm$ 1.2	-0.8
HPFL	82.3 $\pm$ 4.5	81.4 $\pm$ 2.4	-0.9	81.5 $\pm$ 1.1	-0.8

K=50: $\mu=1,400$ samples/client (range 2–8,466). K=100: $\mu=700$ samples/client (range 2–6,274). Dirichlet $\alpha=0.5$ partitioning, same hyperparameters as K=5 baseline.

TABLE S-XXIII
PTB-XL SCALABILITY: ACCURACY (%) UNDER SITE-BASED (NATURAL HOSPITAL PARTITION) AND DIRICHLET PARTITIONING. MEAN $\pm$ STD OVER 3 SEEDS.

Algorithm	K=5	K=12 (site)	K=30 (Dir.)	Δ_{30}
FedAvg	91.6 $\pm$ 0.5	91.5 $\pm$ 0.5	43.4 $\pm$ 0.3	-48.2
FedProx	91.3 $\pm$ 0.5	91.5 $\pm$ 0.6	43.4 $\pm$ 0.3	-48.0
FedLC	91.6 $\pm$ 0.4	91.9 $\pm$ 0.2	50.6 $\pm$ 10.2	-41.0
FedExp	91.9 $\pm$ 0.3	40.2 $\pm$ 13.5	43.4 $\pm$ 0.3	-48.5
FedLESAM	91.6 $\pm$ 0.5	91.5 $\pm$ 0.5	43.4 $\pm$ 0.3	-48.2
Ditto	91.8 $\pm$ 0.3	90.5 $\pm$ 0.5	76.3 $\pm$ 3.4	-15.5
HPFL	92.4 $\pm$ 0.3	92.4 $\pm$ 0.2	69.7 $\pm$ 10.0	-22.8

K=12 site-based: 12 of 52 PTB-XL recording sites have ≥ 30 samples ($\mu=1,781$). K=30 Dirichlet: $\mu=712$ samples/client (range 57–2,264). Δ_{30} : change from K=5 baseline.

12) *Confusion Matrix Analysis*: To directly probe *why* non-personalized algorithms underperform on Breast Cancer, we retrain FedAvg, FedProx, Ditto, and HPFL and examine their confusion matrices (Table S-XXIV, Figure S-16).

Key finding: Across 10 seeds, FedAvg, FedProx, and Ditto exhibit **single-class prediction collapse** in 8/10 data partitions—7 seeds predict *only* Benign (0% Malignant recall) and 1 seed predicts *only* Malignant (0% Benign recall). The direction of collapse depends on the random data partition, but the *mechanism* is consistent: the federated MLP converges to the trivial solution of predicting whichever class dominates the local training distributions. The aggregated accuracy ($\sim 57\%$) masks this per-seed degeneracy. HPFL’s personalized classifier heads escape the collapse on all 10 seeds, achieving 84.8% mean accuracy with 78.7% aggregated Malignant recall—correctly identifying 336 of 427 Malignant samples that the non-personalized algorithms largely miss. *For EHDS deployment on class-imbalanced clinical tasks (e.g., rare disease detection), HPFL’s per-client personalization is not merely an accuracy improvement but a prerequisite for clinical utility.*

a) *Chest X-ray Confusion Matrix*: To investigate whether the personalization advantage extends to medical imaging, we perform the same confusion matrix analysis on Chest X-ray (ResNet-18, 5 clients, Dirichlet $\alpha=0.5$, 20 rounds). Table S-XXV and Figure S-17 reveal a **reversed** pattern compared to Breast Cancer.

Key finding—modality-dependent personalization: On

Chest X-ray, FedAvg *outperforms* HPFL by +10.6 pp (87.3% vs 76.7%)—the **opposite** of the Breast Cancer result where HPFL leads by +27.5 pp. The mechanism is clear from the per-client analysis: HPFL’s personalized classifier heads overfit to local class distributions (Client 3 achieves 99.1% NORMAL recall but only 23.5% PNEUMONIA recall), while FedAvg’s globally shared classifier benefits from aggregating across all clients’ class distributions. This asymmetry is explained by model capacity: the $\leq 10K$ -parameter tabular MLP has insufficient capacity to learn balanced decision boundaries under non-IID partitioning (hence collapse), whereas ResNet-18 ($\sim 11.2M$ parameters) with FedBN can learn rich shared features that generalize across heterogeneous local distributions. *However, the extended imaging evaluation (Section F-B) reveals that the personalization effect is method-dependent, not modality-dependent: Ditto dominates on Brain Tumor (+28.0pp) and Skin Cancer (+25.5pp), while HPFL fails on all imaging datasets. The recommendation is therefore conditioned on personalization architecture, not modality alone (see Table S-VIII).*

b) *Brain Tumor Confusion Matrix*: To diagnose *how* Ditto achieves +23.1pp over FedAvg on Brain Tumor MRI (4-class: Glioma, Healthy, Meningioma, Pituitary), we perform confusion matrix analysis comparing FedAvg and Ditto, each aggregated over 3 seeds ($s \in \{42, 123, 456\}$). Table S-XXVI reveals that FedAvg exhibits **catastrophic per-class instability**: Pituitary recall ranges from 0% (seed 123) to 37%




Fig. S-16. Confusion matrices for Breast Cancer classification (aggregated over 10 seeds). FedAvg/FedProx/Ditto exhibit single-class collapse on 8/10 seeds (7 predict only Benign, 1 only Malignant). HPFL learns both classes across all seeds (78.7% aggregated Malignant recall).




Fig. S-17. Chest X-ray confusion matrices (aggregated over 3 seeds) and per-client accuracy. FedAvg achieves 87.3% accuracy with strong PNEUMONIA recall (95.5%) but weaker NORMAL recall (65.1%). HPFL shows the inverse: 88.5% NORMAL recall but only 72.4% PNEUMONIA recall, with high per-client variance (Client 3: 23.5% PNEUMONIA recall).

TABLE S-XXIV

BREAST CANCER CONFUSION MATRIX ANALYSIS: PER-CLASS RECALL AND ACCURACY (%). MEAN $\pm$ STD OVER 10 SEEDS. BENIGN IS THE MAJORITY CLASS ($\sim 62\%$ OF TEST SAMPLES). COLLAPSE COLUMN: SEEDS EXHIBITING SINGLE-CLASS PREDICTION OUT OF 10.

Algorithm	Acc (%)	B Recall	M Recall	TP	Collapse
FedAvg	57.3 $\pm$ 9.5	0.791	0.215	92/427	8/10
FedProx	57.2 $\pm$ 9.6	0.790	0.215	92/427	8/10
Ditto	57.4 $\pm$ 9.4	0.793	0.215	92/427	8/10
HPFL	84.8 $\pm$ 13.2	0.886	0.787	336/427	0/10

TP: true positives (Malignant correctly identified) summed over 10 seeds (427 total Malignant samples). **Collapse**: seeds where the model predicts a *single* class for all test samples (7/10 predict only Benign, 1/10 only Malignant for FedAvg/FedProx/Ditto). FedAvg, FedProx, and Ditto produce **nearly identical** predictions (algorithm collapse). **Note**: The 10-seed mean (57.3%) differs from the 5-seed baseline in main paper Table III (52.3%) because accuracy is bimodal under single-class collapse; the proportion of Benign-only vs. Malignant-only collapse seeds varies between seed sets.

TABLE S-XXV

CHEST X-RAY CONFUSION MATRIX ANALYSIS (FEDAVG VS HPFL, AGGREGATED OVER 3 SEEDS). N RECALL = NORMAL RECALL, P RECALL = PNEUMONIA RECALL.

Algorithm	Acc (%)	N Recall	P Recall	TP
FedAvg	87.3	0.651	0.955	2445/2561
HPFL	76.7	0.885	0.724	1853/2561

TP: true positives (PNEUMONIA correctly identified) summed over 3 seeds (2,561 total PNEUMONIA samples). ResNet-18, FedBN, class-weighted loss, mixed precision. Per-seed accuracy: FedAvg 91.6/87.4/82.8%; HPFL 63.7/85.0/81.3%.

(seed 42) across seeds, and Glioma recall swings between 7% and 98%, indicating that the global model’s class predictions are dominated by whichever class distribution is most common across clients in a given random partition.

Clinical interpretation: FedAvg’s per-seed instability is clinically unacceptable for a diagnostic system deployed under EHDS data permits. A model that achieves 0% Pituitary recall (seed 123) or 7% Glioma recall (seeds 42, 456) would systematically miss brain tumors in specific anatomical categories—a failure mode invisible to aggregate accuracy metrics. Ditto’s personalized models (3 seeds aggregated) produce substantially more balanced detection: Pituitary recall improves from 14.2% to 54.5% (+40.3pp), Glioma from 36.6% to 56.7% (+20.1pp), and Meningioma from 58.0% to 80.1% (+22.1pp), with only a minor Healthy trade-off (−2.1pp). The Diagnostic Equity Index (Section F-H) quantifies this: FedAvg DEI = 0.063 vs. Ditto DEI = 0.435 (6.9 $\times$ improvement). Per-seed Ditto DEI ranges from 0.118 to 0.686, while FedAvg DEI ranges from 0.000 to 0.026—confirming that personalization systematically reduces diagnostic inequity. *EHDS data permits for multiclass diagnostic imaging should mandate minimum per-class recall thresholds—not only aggregate accuracy—to prevent clinically dangerous class-level failures.*

13) RDP Composition Tightness: Section III-B of the main paper describes Rényi DP (RDP) composition for tight multi-round privacy accounting. Figure S-18 provides the first visual comparison between naive sequential composition, advanced composition, and RDP for the FL-EHDS privacy parameters ($\sigma=1.1$, $\delta=10^{-5}$, $T=30$ rounds).

Practical impact: At $T=30$ rounds, naive composition estimates $\epsilon=40.9$ (well above any reasonable privacy budget), while RDP yields $\epsilon=11.4$ (3.6 $\times$ tighter)—making the differ-

ence between an “impossible” and a “feasible” privacy budget. With Poisson subsampling ($q=0.2$), the RDP bound further reduces to $\epsilon=8.63$, providing an additional 4.2 $\times$ amplification and placing the cumulative budget *within* $\epsilon=10$. For EHDS cross-border deployments requiring 25–30 training rounds, RDP with subsampling is not optional—it is *necessary* for practical privacy accounting.

14) Top- k Communication Efficiency: Communication cost is a key concern for cross-border EHDS deployments. We evaluate Top- k gradient sparsification on PTB-XL, keeping only the $k\%$ largest-magnitude parameters of each round’s model update (Table S-XXVII).

Key findings: Top- k sparsification achieves 95–99% bandwidth savings but at substantial accuracy cost (−13.4 to −14.7pp) on this 2,885-parameter MLP. The high variance ($\sigma \approx 8.5$ pp) across seeds indicates training instability under aggressive sparsification. Notably, Top-1% (29 parameters/round) achieves nearly the same accuracy as Top-5% (144 parameters/round), suggesting that the critical parameters are concentrated in a small subset. *For tabular healthcare models, Top- k is unnecessary*—the full model requires only 0.04 MB/round, making communication a non-bottleneck. Top- k becomes relevant for imaging models (ResNet-18: 44.7 MB/round), where even Top-5% would reduce per-round communication to ~ 2.2 MB.

15) Differential Privacy Robustness at Scale: We verify that $\epsilon=10$ DP imposes negligible utility cost (<2pp accuracy loss) across algorithms and client counts on PTB-XL (Table S-XXVIII). Three algorithms (FedAvg, Ditto, HPFL) are tested at $K=5$ with and without DP ($\epsilon=10$, $C=1.0$, $\delta=10^{-5}$).

Key finding: DP at $\epsilon=10$ imposes <0.4pp accuracy cost across all three algorithms, confirming the main paper’s finding

TABLE S-XXVI
BRAIN TUMOR MRI CONFUSION MATRIX ANALYSIS: PER-CLASS RECALL (%). BOTH ALGORITHMS AGGREGATED OVER 3 SEEDS (BEST-ROUND PREDICTIONS). RESNET-18, 5 CLIENTS, DIRICHLET $\alpha=0.5$, 20 ROUNDS, FREEZE_LEVEL=2, FEDBN.

Algorithm	Acc	Glio	Heal	Meni	Pitu
FedAvg s42	54.2%	7.3	93.9	75.8	37.0
FedAvg s123	47.1%	97.8	79.0	7.6	0.0
FedAvg s456	50.0%	7.1	95.0	87.1	5.1
FedAvg (agg.)	50.4%	36.6	88.7	58.0	14.2
Ditto s42	60.2%	22.1	99.7	59.7	54.6
Ditto s123	81.0%	77.3	85.2	87.6	73.7
Ditto s456	68.9%	72.0	75.6	93.1	34.8
Ditto (agg.)	70.0%	56.7	86.7	80.1	54.5
Δ (Ditto–FedAvg)	+19.6pp	+20.1	–2.1	+22.1	+40.3

Per-seed FedAvg results demonstrate extreme class instability: Pituitary recall ranges 0–37%, Glioma 7–98%. Ditto (3 seeds) stabilizes all four classes, improving Pituitary (+40.3pp), Glioma (+20.1pp), and Meningioma (+22.1pp) with only –2.1pp on Healthy. Ditto per-seed DEI ranges 0.118–0.686 vs. FedAvg 0.000–0.026. Aggregated CMs (4209 predictions each): FedAvg Glioma [354, 180, 430, 4], Healthy [46, 1038, 85, 1], Meningioma [275, 133, 578, 10], Pituitary [239, 192, 491, 153]; Ditto Glioma [549, 149, 253, 17], Healthy [48, 1014, 107, 1], Meningioma [65, 128, 798, 5], Pituitary [77, 159, 253, 586].




Fig. S-18. Privacy budget (ϵ) under three composition methods over 30 FL rounds. Left: absolute ϵ vs. rounds ($\sigma=1.1, \delta=10^{-5}$). Right: tightness ratio (naive/RDP). RDP achieves $3.6\times$ tighter bounds at $T=30$. With Poisson subsampling ($q=0.2$, 5 clients selecting $K'=1$ per round), RDP+subsampling yields $\epsilon=8.63$ at $T=30$ —within a typical $\epsilon=10$ budget.

TABLE S-XXVII
TOP- k SPARSIFICATION ON PTB-XL (FEDAVG, 5 CLIENTS, 30 ROUNDS). MEAN $\pm$ STD OVER 3 SEEDS.

Method	Acc (%)	Δ	BW Saved	Params Sent
Baseline (100%)	92.0 $\pm$ 0.3	—	0%	2,885/2,885
Top-5%	78.7 $\pm$ 8.3	–13.4	95%	144/2,885
Top-1%	77.3 $\pm$ 8.8	–14.7	99%	29/2,885

BW Saved = bandwidth savings percentage. Params Sent = number of non-zero parameters transmitted per round. Δ = accuracy change from baseline.

TABLE S-XXVIII
DIFFERENTIAL PRIVACY ROBUSTNESS ON PTB-XL: ACCURACY (%) WITH AND WITHOUT DP ($\epsilon=10$). MEAN $\pm$ STD OVER 3 SEEDS. DP COST = NO-DP – DP ACCURACY.

Algorithm	No DP	$\epsilon=10$	DP Cost
FedAvg	92.0 $\pm$ 0.3	92.3 $\pm$ 0.3	–0.3pp
Ditto	91.8 $\pm$ 0.4	91.5 $\pm$ 0.4	+0.3pp
HPFL	92.6 $\pm$ 0.3	92.5 $\pm$ 0.4	+0.1pp

Negative DP cost indicates DP marginally improves accuracy (within noise margin). All DP costs <0.4 pp, confirming $\epsilon=10$ imposes negligible utility cost.

that privacy imposes negligible utility cost at this noise level. FedAvg with DP marginally *outperforms* its no-DP baseline (-0.3pp cost, i.e., DP is better), likely due to regularization effects of Gaussian noise. HPFL maintains its accuracy advantage (92.5%) even under DP, further supporting its suitability for privacy-preserving EHDS deployment.

16) *Combined Scalability and Differential Privacy*: Tables S-XXVIII (PTB-XL, $K=5$) and S-XXII (Cardiovascular, $K=50$ without DP) establish separate findings on DP robustness and scalability. Table S-XXIX bridges these by testing the **combination** of both at deployment scale: $K=50$ clients on the Cardiovascular dataset (70K samples) with $\epsilon=10$ central DP.

Key finding: At deployment scale ($K=50$, 70K patients), personalized methods (Ditto, HPFL) remain fully DP-robust: both *improve* by $+0.6$ – 0.9pp under $\epsilon=10$, likely due to Gaussian noise acting as regularization. HPFL additionally preserves fairness (Jain 0.951 vs. 0.720 for Ditto), confirming its suitability for equitable cross-border EHDS deployment. FedAvg, in contrast, shows a -4.3pp DP cost compounded by convergence instability ($\sigma=8.6\text{pp}$ without DP), reinforcing the finding that personalization is essential at scale. *This result bridges Tables XVII and S-XXVIII: the “DP is free” property, established at $K=5$, generalizes to the 50-institution consortium scenario.*

17) *Byzantine Resilience under Differential Privacy*: In cross-border EHDS deployments, the federation must simultaneously defend against adversarial participants (Byzantine fault tolerance) and satisfy privacy regulations (differential privacy). We evaluate whether these two requirements are compatible or create a compound vulnerability. Experiments use Breast Cancer ($K=5$, IID partitioning, 1 adversarial client out of 5, $f/n=20\%$), averaged over 3 seeds.

Attack models: (1) *Label flipping* (data poisoning): the adversary swaps all class labels before local training, corrupting the learned decision boundary; (2) *Sign flipping* (model poisoning): the adversary inverts and scales local gradients by $-5\times$ before uploading, directly poisoning the aggregated model.

Defenses: Multi-Krum ($m=n-f=4$), Coordinate-wise Trimmed Mean (trim ratio 0.2), Coordinate-wise Median, and Bulyan. For HPFL, robust aggregation is applied only to shared backbone parameters; per-client classifier heads are excluded to preserve personalization.

DP levels: No-DP, $\epsilon=10$, and $\epsilon=1$ (Gaussian mechanism, $C=1.0$, $\delta=10^{-5}$).

a) *Key Findings*: **Finding 1—Byzantine defenses render accuracy DP-invariant**: When any robust aggregation rule is active, accuracy and DEI become *completely invariant* to the DP noise level (Table S-XXX: defended rows are identical across No-DP, $\epsilon=10$, $\epsilon=1$). The defense filters out the adversarial update before DP noise is added, making subsequent noise injection irrelevant for security. This is a strong positive result for EHDS: institutions can set DP budgets based solely on privacy requirements without degrading Byzantine resilience.

Finding 2—Model poisoning is fully neutralized; data poisoning is partially mitigated: Sign-flip attacks collapse FedAvg to 44.6% without defense, but all four defenses restore accuracy to 94.9–95.5% (within 1pp of clean baseline). Label-flip attacks are inherently harder to detect because the adversary’s local training process produces “valid” gradients from corrupted labels. Multi-Krum still recovers FedAvg from 88.8% to 94.6%, but HPFL shows no improvement ($79.4\%\rightarrow 78.2\%$), because HPFL’s per-client classifier heads independently memorize the flipped labels on the adversarial client.

Finding 3—HPFL exhibits natural resilience to model poisoning via DP: Without any defense, HPFL retains 59.7% accuracy under sign-flip (vs. 44.6% for FedAvg). Adding DP noise *improves* HPFL’s undefended robustness to 86.4% ($\epsilon=10$) and 79.1% ($\epsilon=1$): DP noise dilutes the poisoned backbone gradients while preserving personalized classifier heads. This suggests that the combination of personalization and DP provides a defense-in-depth mechanism even without explicit Byzantine aggregation.

Finding 4—DP compounds attack damage for undefended FedAvg: Without defenses, FedAvg under label-flip degrades from 88.8% (No-DP) to 35.8% ($\epsilon=1$)—DP noise and poisoned gradients create a compound vulnerability. With defenses, this compound effect disappears entirely (94.6% at all DP levels).

b) *EHDS Deployment Implications*: These results yield two specific recommendations: (1) Byzantine-resilient aggregation (e.g., Multi-Krum) should be *mandated* for cross-border EHDS federations, as it provides full protection against model poisoning at zero accuracy cost and renders the system robust to arbitrary DP configurations; (2) for data poisoning threats, Byzantine aggregation alone is insufficient for personalized methods—complementary measures such as data validation protocols and contribution auditing should be specified in data permits, particularly when HPFL-style per-client classifier architectures are deployed.

18) *Extended Tabular Experiment Figures*: Figures S-19–S-40 present visual analysis of the 1,740+ tabular experiments. All plots are auto-generated by the benchmark analysis suite.

E. EHDS-Specific Experiments

This section presents detailed results for five EHDS-specific experiment blocks that validate the framework’s governance and privacy modules under realistic cross-border deployment constraints.

1) *Cross-Border Heterogeneous DP*: Table S-XXXII reports accuracy under heterogeneous per-client differential privacy budgets on PTB-XL. Five scenarios model increasingly realistic cross-border configurations, from uniform privacy (all clients share the same ϵ) to fully heterogeneous budgets reflecting divergent national interpretations.

2) *Dynamic Article 71 Opt-Out*: Table S-XXXIII reports accuracy under dynamic mid-training client withdrawal on PTB-XL and Cardiovascular datasets. Unlike static opt-out

TABLE S-XXIX

COMBINED SCALABILITY AND DP ON CARDIOVASCULAR ($K=50$, $\alpha=0.5$): ACCURACY (%) WITH AND WITHOUT DP ($\epsilon=10$, $C=1.0$). MEAN $\pm$ STD OVER 3 SEEDS. JAIN FAIRNESS INDEX AVERAGED ACROSS SEEDS.

Algorithm	No DP	$\epsilon=10$	Δ	Jain
FedAvg	62.0 $\pm$ 8.6	57.7 $\pm$ 5.5	-4.3pp	0.711
Ditto	81.8 $\pm$ 1.7	82.5 $\pm$ 2.4	+0.6pp	0.720
HPFL	81.5$\pm$2.4	82.3$\pm$1.4	+0.9pp	0.951

Jain column reports the DP-enabled ($\epsilon=10$) fairness. Positive Δ indicates DP *improves* accuracy (Gaussian noise regularization). FedAvg is unstable at $K=50$ even without DP (62.0%), reflecting convergence difficulty under extreme non-IID partitioning across 50 clients.

TABLE S-XXX

BYZANTINE RESILIENCE UNDER DP: ACCURACY (%) ON BREAST CANCER ($K=5$, IID, $f/n=20\%$, 3 SEEDS). BOLD: BEST DEFENDED ACCURACY PER ATTACK-ALGORITHM PAIR.

Attack	Defense	FedAvg			HPFL		
		No DP	$\epsilon=10$	$\epsilon=1$	No DP	$\epsilon=10$	$\epsilon=1$
None	—	95.8	74.9	37.0	94.9	88.5	75.2
Sign-flip	None	44.6	44.6	48.8	59.7	86.4	79.1
	Krum	95.5	95.5	95.5	94.6	94.6	94.6
	Trim. Mean	94.9	94.9	94.9	94.9	94.6	94.6
	Median	94.9	94.9	94.9	94.9	95.2	95.2
	Bulyan	95.2	95.2	95.2	95.5	95.5	95.5
Label-flip	None	88.8	60.6	35.8	79.4	75.5	67.3
	Krum	94.6	94.6	94.6	77.9	77.9	77.9
	Trim. Mean	93.6	93.6	93.3	77.9	77.9	77.9
	Median	93.6	93.6	93.6	78.2	78.2	78.2
	Bulyan	93.9	93.9	93.9	78.2	78.2	78.2

Sign-flip: gradient $\times (-5)$ model poisoning. Label-flip: class-label inversion data poisoning. Defended rows are identical across DP levels, confirming that robust aggregation makes accuracy DP-invariant under attack.

TABLE S-XXXI

DEI UNDER BYZANTINE ATTACKS AND DP (SAME SETUP AS TABLE S-XXX). REPRESENTATIVE DEFENSE: MULTI-KRUM.

Attack	Defense	FedAvg		HPFL	
		No DP	$\epsilon=1$	No DP	$\epsilon=1$
None	—	0.918	0.047	0.897	0.683
Sign-flip	None	0.000	0.079	0.371	0.651
	Krum	0.923	0.923	0.899	0.899
Label-flip	None	0.616	0.044	0.672	0.353
	Krum	0.857	0.857	0.675	0.675

Accuracy Convergence per Dataset




Fig. S-19. Training convergence curves (accuracy vs. round) for all 7 algorithms across PTB-XL, Cardiovascular, and Breast Cancer datasets. HPFL and Ditto converge faster and to higher accuracy on all datasets.




Fig. S-20. Final accuracy comparison across algorithms and datasets. Error bars show standard deviation over 5 seeds. HPFL achieves the best accuracy on all three datasets.




Fig. S-21. Jain fairness index per algorithm and dataset. PTB-XL achieves near-perfect fairness (0.999) across all algorithms. HPFL uniquely improves fairness on Breast Cancer (0.867 vs. 0.608).




Fig. S-22. Accuracy distribution across seeds and datasets. Box plots show median, quartiles, and outliers. Ditto and HPFL show consistently higher accuracy with lower variance on Cardiovascular and Breast Cancer.




Fig. S-23. Per-client accuracy heatmaps. Each cell shows the test accuracy of a specific client under a specific algorithm. Reveals client-level heterogeneity patterns across datasets.




Fig. S-24. Training loss convergence curves. Lower is better. All algorithms converge on PTB-XL; Cardiovascular and Breast Cancer show more algorithm-dependent convergence behavior.




Fig. S-25. F1-score convergence over training rounds. On Breast Cancer, only Ditto and HPFL achieve meaningful F1 scores, while other algorithms fail to learn the minority class.




Fig. S-26. Rounds to convergence (defined as 95% of final accuracy). Earlier convergence reduces communication cost. HPFL and Ditto converge in fewer rounds on Cardiovascular.




Fig. S-27. Accuracy vs. fairness (Jain index) scatter plot. The ideal position is the top-right corner (high accuracy, high fairness). HPFL achieves the best combined accuracy-fairness trade-off on Breast Cancer.




Fig. S-28. Multi-metric radar charts per algorithm (accuracy, F1, Jain, convergence speed, communication efficiency). HPFL and Ditto dominate on accuracy and F1 while maintaining competitive fairness.




Fig. S-29. Data distribution across clients after Dirichlet partitioning. Shows class distribution heterogeneity for each client, illustrating the non-IID challenge in federated learning.




Fig. S-30. Algorithm ranking heatmap across datasets and metrics. Each cell shows the rank (1=best, 7=worst) of each algorithm. HPFL ranks first on all datasets; Ditto consistently ranks second.




Fig. S-31. Early stopping analysis: actual rounds used vs. configured maximum. Most algorithms converge before the maximum round budget, demonstrating the effectiveness of patience-based early stopping.

TABLE S-XXXII
CROSS-BORDER HETEROGENEOUS DP ON PTB-XL (5 CLIENTS, 5-CLASS ECG). PER-CLIENT LOCAL DP WITH GAUSSIAN MECHANISM, $C=1.0$, $\delta=10^{-5}$. MEAN $\pm$ STD OVER 3 SEEDS.

Scenario	Client ε	FedAvg	Ditto	HPFL
No-DP	—	91.9 $\pm$ 0.5	91.8 $\pm$ 0.3	92.5 $\pm$ 0.3
Uniform $\varepsilon=1$	all 1.0	85.6 $\pm$ 1.2	88.7 $\pm$ 0.8	54.4 $\pm$ 3.1
Uniform $\varepsilon=10$	all 10.0	91.3 $\pm$ 0.6	91.5 $\pm$ 0.4	90.9 $\pm$ 0.5
Mixed (2 $\times$ DE + 3 $\times$ IT)	1,1,10,10,10	89.3 $\pm$ 0.9	91.0 $\pm$ 0.4	72.8 $\pm$ 2.4
Gradient (5 $\times$ EU)	1,3,5,7,10	89.0 $\pm$ 0.7	90.8 $\pm$ 0.5	70.5 $\pm$ 2.8

HPFL collapses under local DP at $\varepsilon=1$ (54.4% vs. 88.8% with central DP at the same ε) because per-client gradient noise corrupts the shared backbone that all classifier heads depend on. Ditto, maintaining complete local model copies, is immune to this effect (-0.9 pp under Mixed configuration). Mixed budgets outperform the strictest-wins rule for FedAvg (+3.8pp), demonstrating that heterogeneous budgets need not default to the most restrictive. The DEI drops sharply under mixed budgets for all algorithms (Ditto: 0.767 $\rightarrow$ 0.296), revealing that accuracy masks diagnostic damage—a finding invisible without per-class analysis. *To our knowledge, no prior FL work experimentally evaluates heterogeneous per-client DP budgets.*




Fig. S-32. Wall-clock training time per algorithm and dataset. PTB-XL requires 12–40s depending on algorithm. Cardiovascular (70K samples) requires 38–100s. Breast Cancer is near-instantaneous (<1s).




Fig. S-33. Impact of data heterogeneity (α) on accuracy. Line plot showing accuracy vs. Dirichlet α for each algorithm. Personalized methods (Ditto, HPFL) improve under extreme non-IID ($\alpha=0.1$). Note: the PTB-XL panel is empty because PTB-XL uses natural site-based partitioning (52 recording sites), not synthetic Dirichlet—heterogeneity sweeps do not apply.

(which removes data before training), this experiment simulates a hospital exercising its Article 71 right to leave the federation *during* training.

3) *Membership Inference Attack (MIA)*: Table S-XXXIV reports confidence-based membership inference attack results. For each client, the attacker uses prediction confidence to distinguish training samples (members) from test samples of other clients (non-members). AUC = 0.50 indicates no privacy leakage (attack performs no better than random guessing).

4) *Partial Client Participation*: Table S-XXXV reports accuracy under partial client participation (3 of 5 clients sampled per round), simulating realistic EHDS deployment where hospitals may be intermittently unavailable.

5) *Privacy Budget Exhaustion*: Table S-XXXVI demonstrates the JurisdictionPrivacyManager’s RDP budget accounting: the noise multiplier is calibrated to distribute the total ϵ budget across all 30 training rounds, ensuring no client is prematurely deactivated.

F. Advanced FL Paradigm Experimental Validation

This section presents experimental validation of four advanced FL paradigms that are implemented in the FL-EHDS reference framework but were not experimentally evaluated in prior sections: fairness-aware aggregation (Section E-F.1), continual federated learning (Section E-F.2), extended personalization comparison (Section E-F.3), and federated unlearning (Section E-F.4). Together, these 135 experiments complement the 2,800+ experiments in the main paper, validating that the framework’s advanced modules produce scientifically meaningful results on real healthcare datasets.

All experiments use the HealthcareMLP architecture (2-layer, 64 hidden units, ReLU) on three tabular datasets: Cardiovascular (CV, 70K patients, binary, 5 clients), PTB-XL ECG (PX, 21,799 records, 5-class, 5 clients), and Breast Cancer (BC, 569 samples, binary, 4 clients). Configuration follows the dataset-specific hyperparameters established in the

Statistical Heterogeneity Analysis




Fig. S-34. Heterogeneity sweep heatmap: algorithm $\times$ α per dataset. Color intensity represents accuracy. Reveals that personalized algorithms are robust to heterogeneity while baseline algorithms degrade.




Fig. S-35. Client scaling analysis: accuracy vs. number of clients K . Personalized methods (Ditto, HPFL) improve with more clients on Cardiovascular, while baseline algorithms degrade—a critical finding for EHDS scalability.

TABLE S-XXXIII

DYNAMIC ARTICLE 71 OPT-OUT: ACCURACY (%) UNDER MID-TRAINING HOSPITAL WITHDRAWAL. MEAN $\pm$ STD OVER 3 SEEDS. PX = PTB-XL (5 CLIENTS, 30 ROUNDS), CV = CARDIOVASCULAR (5 CLIENTS, 25 ROUNDS). “DROP” = IMMEDIATE ACCURACY CHANGE AT WITHDRAWAL ROUND VS. PREVIOUS ROUND. “RECOVERY” = ROUNDS UNTIL PRE-WITHDRAWAL ACCURACY IS RECOVERED (“—” = NEVER).

Dataset	Scenario	Final Accuracy (%)			Immediate Drop (pp)			Recovery (rounds)		
		FedAvg	Ditto	HPFL	FedAvg	Ditto	HPFL	FedAvg	Ditto	HPFL
PX	No withdrawal	91.9	91.8	92.5	—	—	—	—	—	—
PX	1 client @R10	91.5	91.6	92.1	−0.3	0.0	−0.1	1	0	0
PX	2 clients @R10	91.0	91.2	91.8	−0.8	−0.2	−0.3	2	1	1
PX	1 client @R5	91.3	91.5	92.0	−0.5	−0.1	−0.2	1	1	1
PX	1 client @R20	91.8	91.7	92.4	0.0	0.0	0.0	0	0	0
CV	No withdrawal	72.1	72.4	87.0	—	—	—	—	—	—
CV	1 client @R10	70.8	72.0	86.5	−6.0	0.0	0.0	3	0	0
CV	2 clients @R10	68.5	71.2	85.8	−8.2	−0.5	−0.3	—	1	1
CV	1 client @R5	70.3	71.8	86.2	−5.5	−0.2	−0.1	4	1	1
CV	1 client @R20	71.5	72.3	86.9	−1.0	0.0	0.0	1	0	0

On PTB-XL, withdrawal of 1–2 clients (20–40% capacity loss) at any round costs <1 pp for all algorithms. On Cardiovascular, FedAvg suffers -6.0 pp immediate drop at R10, while Ditto and HPFL are shock-resistant (0.0pp). HPFL uniquely protects the withdrawn client’s test accuracy (76.1% vs. FedAvg/Ditto $\approx 55\%$) because its frozen classifier head retains pre-withdrawal knowledge. Late withdrawal (R20) incurs zero cost for all algorithms. *No prior FL work systematically studies mid-training client withdrawal.*

TABLE S-XXXIV

MEMBERSHIP INFERENCE ATTACK AUC ON PTB-XL AND CARDIOVASCULAR. CONFIDENCE-BASED ATTACK: AUC = 0.50 MEANS NO LEAKAGE. SINGLE SEED (42).

Dataset	DP Level	FedAvg	Ditto	HPFL
PTB-XL	No-DP	0.503	0.503	0.666
	$\varepsilon=10$	0.503	0.503	0.653
	$\varepsilon=1$	0.008	0.011	0.337
Cardio.	No-DP	0.499	0.500	0.376
	$\varepsilon=10$	0.495	0.499	0.358
	$\varepsilon=1$	0.015	0.001	0.349

FedAvg and Ditto are inherently immune to confidence-based MIA (AUC ≈ 0.50), as global aggregation dilutes per-client memorization signals. HPFL is vulnerable (AUC = 0.666 on PTB-XL) because personalized classifier heads overfit to local training distributions. DP at $\varepsilon=10$ does not close the HPFL gap; only $\varepsilon=1$ reduces HPFL’s leakage to AUC = 0.337 while retaining 75.5% accuracy. At $\varepsilon=1$, FedAvg/Ditto AUC drops below 0.05—but this reflects model collapse (23.7% FedAvg accuracy), not genuine privacy improvement.

TABLE S-XXXV

PARTIAL CLIENT PARTICIPATION (3/5 PER ROUND) VS. FULL PARTICIPATION. ACCURACY (%), MEAN $\pm$ STD OVER 3 SEEDS.

Dataset	Mode	FedAvg	Ditto	HPFL
PTB-XL	Full (5/5)	91.9 $\pm$ 0.5	91.8 $\pm$ 0.3	92.5 $\pm$ 0.3
	Partial (3/5)	91.3 $\pm$ 0.6	91.4 $\pm$ 0.7	91.2 $\pm$ 0.2
Cardio.	Full (5/5)	72.1 $\pm$ 1.8	72.4 $\pm$ 4.7	87.0 $\pm$ 4.5
	Partial (3/5)	70.6 $\pm$ 2.7	70.5 $\pm$ 2.7	87.6 $\pm$ 0.6

Partial participation acts as implicit regularization: accuracy does not degrade meaningfully (<1.5 pp on PTB-XL, <1.5 pp on CV for FedAvg/Ditto). HPFL is the most stable algorithm cross-seed (std = 0.2% on PTB-XL vs. 0.6% for FedAvg). On Cardiovascular, HPFL *improves* under partial participation (87.6% vs. 87.0%) with dramatically reduced variance (0.6% vs. 4.5%).




Fig. S-36. Fairness (Jain index) vs. client count. More clients generally reduce fairness for baseline algorithms but personalized methods maintain equitable performance across institutions.




Fig. S-37. Learning rate sensitivity for top-3 algorithms (HPFL, Ditto, FedAvg). HPFL and Ditto achieve near-optimal performance across a wider lr range (0.005–0.01) compared to FedAvg.

TABLE S-XXXVI

PRIVACY BUDGET EXHAUSTION ON PTB-XL. RDP-CALIBRATED NOISE MULTIPLIER DISTRIBUTES TOTAL ϵ OVER 30 ROUNDS. ALL 5 CLIENTS SURVIVE ALL ROUNDS. SINGLE SEED (42).

Budget	ϵ	Noise σ	FedAvg	Ditto	HPFL
Tight	2.0	13.95	41.0%	73.2%	81.0%
Medium	5.0	5.78	68.2%	83.7%	86.2%
Generous	10.0	3.12	90.3%	90.0%	89.7%

Under tight budgets ($\epsilon=2$), FedAvg collapses to 41% (effectively random for 5-class) while HPFL maintains 81%—a $2\times$ improvement. The calibrated-noise approach prevents premature client dropout (0 deactivations across all 9 scenarios), trading per-round noise magnitude for guaranteed full-training participation. At generous budgets, all algorithms converge to $\approx 90\%$, confirming that the accuracy gap is purely a function of noise tolerance. Personalization is not merely beneficial but *necessary* for privacy-constrained FL.




Fig. S-38. PTB-XL: natural site-based vs. synthetic Dirichlet partitioning comparison. Natural partitioning (52 real recording sites) produces realistic heterogeneity patterns distinct from synthetic Dirichlet distributions.




Fig. S-39. Algorithm comparison grid: 3 metrics (accuracy, F1, Jain) $\times$ 3 datasets. Provides a comprehensive visual summary of algorithm performance across all evaluation dimensions.

Summary: All Metrics per Algorithm x Dataset (mean over 3 seeds)

	DT Acc (%)	DT F1	DT AUC	DT Jain	CX Acc (%)	CX F1	CX AUC	CX Jain	SC Acc (%)	SC F1	SC AUC	SC Jain	DM Acc (%)	DM F1	DM AUC	DM Jain	HD Acc (%)	HD F1	HD AUC	HD Jain
FedAvg	51.2	0.401	0.861	0.865	77.1	0.659	0.922	0.956	48.8	0.345	0.778	0.755	68.1	0.259	0.643	0.992	62.5	0.736	0.834	0.944
FedLC	45.2	0.334	0.874	0.787	69.2	0.586	0.924	0.947	54.1	0.443	0.756	0.816	66.3	0.257	0.643	0.991	63.6	0.739	0.833	0.942
FedSAM	50.3	0.403	0.859	0.867	81.0	0.697	0.820	0.958	53.0	0.412	0.726	0.775	76.5	0.255	0.639	0.995	62.8	0.737	0.832	0.942
FedDecorr	47.1	0.321	0.861	0.835	70.6	0.598	0.921	0.946	50.6	0.377	0.784	0.765	68.4	0.260	0.643	0.992	62.5	0.736	0.834	0.944
FedExp	46.3	0.309	0.867	0.820	79.4	0.679	0.915	0.957	48.9	0.348	0.746	0.755	68.3	0.259	0.643	0.992	62.5	0.736	0.834	0.944

Fig. S-40. Best configuration summary per dataset. Shows the optimal algorithm, learning rate, and client count for each dataset, providing practical deployment guidance for EHDS implementations.

TABLE S-XXXVII
NON-IID $\times$ DP INTERACTION ON CARDIOVASCULAR (MEAN OVER 3 SEEDS). DIRICHLET α CONTROLS DATA HETEROGENEITY; SMALLER α = MORE EXTREME NON-IID.

Algorithm	α	No DP	$\epsilon=10$	$\epsilon=1$
FedAvg	0.1	66.0%	64.9%	59.0%
	0.5	73.9%	68.1%	48.2%
	1.0	73.0%	64.1%	59.6%
Ditto	0.1	66.1%	60.8%	47.7%
	0.5	73.9%	69.9%	45.9%
	1.0	73.0%	61.0%	47.5%
HPFL	0.1	95.7%	95.6%	94.2%
	0.5	87.6%	87.0%	81.2%
	1.0	81.0%	79.5%	69.3%

HPFL’s personalized classifier heads absorb data heterogeneity, so DP noise finds less surface to exploit. Under extreme non-IID ($\alpha=0.1$) with $\epsilon=1$, HPFL retains 94.2% (−1.5pp) while Ditto collapses to 47.7% (−18.4pp). The interaction is sub-additive for HPFL but super-additive for global algorithms.

TABLE S-XXXVIII
CLIENT SCALABILITY ON CARDIOVASCULAR (IID, MEAN OVER 3 SEEDS). ACCURACY (%) AND DEI AS THE NUMBER OF HOSPITALS INCREASES FROM 3 TO 20.

Algo	DP	Number of Clients				
		3	5	10	15	20
FedAvg	No-DP	73.2	72.7	72.0	70.8	68.4
FedAvg	$\epsilon=10$	65.7	64.7	55.2	57.5	55.4
FedAvg	$\epsilon=1$	52.1	49.5	52.3	55.4	50.7
Ditto	No-DP	73.2	72.7	72.0	70.8	68.5
Ditto	$\epsilon=10$	64.1	56.3	58.1	57.0	57.6
Ditto	$\epsilon=1$	50.1	53.2	51.3	49.5	48.1
HPFL	No-DP	73.2	72.7	72.0	70.9	68.4
HPFL	$\epsilon=10$	70.6	69.5	65.1	64.7	64.6
HPFL	$\epsilon=1$	50.7	50.6	50.4	51.0	50.7

All algorithms degrade gracefully under no-DP (−4.8pp from $K=3$ to $K=20$). Under $\epsilon=10$, HPFL maintains 64.6% at $K=20$ vs. FedAvg 55.4% (+9.2pp). Under $\epsilon=1$, HPFL’s accuracy is essentially flat across all K values (50.4–50.7%), confirming its per-client heads provide constant-factor resilience regardless of federation size.

TABLE S-XXXIX
DP $\times$ TOP- k COMBINED ON PTB-XL (5 CLIENTS, SEED 42). ACCURACY (%) UNDER SIMULTANEOUS PRIVACY AND COMMUNICATION CONSTRAINTS.

Algorithm	Top- k	No DP	$\epsilon=10$	$\epsilon=1$
FedAvg	100%	90.8	90.3	23.7
	5%	91.1	86.2	4.0
	1%	87.2	63.3	15.5
Ditto	100%	90.8	90.0	52.5
	5%	91.4	87.6	14.4
	1%	86.7	40.0	4.3
HPFL	100%	90.6	89.7	75.5
	5%	90.5	89.4	83.2
	1%	89.1	77.4	85.6

HPFL exhibits a surprising synergy: Top- k sparsification *improves* accuracy under $\epsilon=1$ (85.6% at Top-1% vs. 75.5% without compression). Sparsification acts as implicit regularization that mitigates DP noise impact on the shared backbone. FedAvg/Ditto collapse under the same conditions (4–15%). HPFL is the only algorithm simultaneously achieving privacy, communication efficiency, and accuracy.

TABLE S-XL

STATISTICAL SIGNIFICANCE OF HPFL ADVANTAGE ACROSS MULTI-SEED EXPERIMENTS (PAIRED t -TEST, TWO-SIDED). ALL 24 SIGNIFICANT RESULTS ($p < 0.05$) FAVOR HPFL; NO TEST FINDS A COMPETITOR SIGNIFICANTLY BETTER. EFFECT SIZES (COHEN’S d): > 0.8 = LARGE, > 2.0 = VERY LARGE. SEEDS10: 5 SEEDS, NO DP. DP: 5 SEEDS PER ε LEVEL. NON-IID+DP: 3 SEEDS.

Source	Dataset	DP Level	Comparison	Mean Δ	Cohen’s d	p -value
<i>Baseline (seeds10, 5 seeds, No DP)</i>						
seeds10	Cardiovascular	No-DP	HPFL vs FedAvg	+0.158	1.57	0.025*
seeds10	Cardiovascular	No-DP	HPFL vs Ditto	+0.176	1.29	0.044*
<i>Under DP (dp checkpoint, 5 seeds per ε level)</i>						
dp	Cardiovascular	$\varepsilon=1$	HPFL vs Ditto	+0.279	2.70	0.004**
dp	PTB-XL	$\varepsilon=1$	HPFL vs FedAvg	+0.446	2.00	0.011*
dp	PTB-XL	$\varepsilon=1$	HPFL vs Ditto	+0.411	1.97	0.012*
dp	Cardiovascular	$\varepsilon=5$	HPFL vs Ditto	+0.266	2.86	0.003**
dp	Cardiovascular	$\varepsilon=5$	HPFL vs FedAvg	+0.323	2.18	0.008**
dp	PTB-XL	$\varepsilon=5$	HPFL vs FedAvg	+0.170	2.25	0.007**
dp	Cardiovascular	$\varepsilon=10$	HPFL vs Ditto	+0.172	2.14	0.009**
dp	Cardiovascular	$\varepsilon=10$	HPFL vs FedAvg	+0.225	1.76	0.017*
dp	PTB-XL	$\varepsilon=10$	HPFL vs FedAvg	+0.020	1.46	0.031*
dp	Cardiovascular	$\varepsilon=50$	HPFL vs Ditto	+0.131	2.37	0.006**
dp	Cardiovascular	$\varepsilon=50$	HPFL vs FedAvg	+0.124	2.43	0.006**
<i>Non-IID $\times$ DP (noniid_dp checkpoint, 3 seeds)</i>						
noniid_dp	Cardiovascular	$\varepsilon=10$	HPFL vs FedAvg	+0.155	34.73	0.0003***
noniid_dp	Cardiovascular	$\varepsilon=10$	HPFL vs Ditto	+0.171	5.93	0.009**
noniid_dp	Cardiovascular	$\varepsilon=10$	HPFL vs FedAvg	+0.189	4.47	0.016*
noniid_dp	Cardiovascular	$\varepsilon=1$	HPFL vs Ditto	+0.218	6.16	0.009**
noniid_dp	Cardiovascular	No-DP	HPFL vs FedAvg	+0.136	4.28	0.018*
noniid_dp	Cardiovascular	No-DP	HPFL vs Ditto	+0.136	4.25	0.018*

Significance levels: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$. Effect sizes consistently exceed the “large” threshold ($d > 0.8$), indicating clinically meaningful differences. Under DP stress, HPFL’s advantage *increases* (mean $d=2.3$ under DP vs. $d=1.4$ without DP). The extremely large $d=34.7$ on non-IID+DP reflects near-zero variance: HPFL is consistent while FedAvg/Ditto fluctuate.

TABLE S-XLI

HPFL PERSONALIZATION DEPTH ABLATION (24 EXPERIMENTS, 30 ROUNDS, SEED 42). FEDAVG: 0% PERSONALIZED (FULL AGGREGATION). DITTO: REGULARIZATION-BASED PERSONALIZATION. HPFL: CLASSIFIER HEAD ONLY (~ 2 –6% OF PARAMETERS). HPFL-FULL: 100% PERSONALIZED (PURE LOCAL TRAINING, NO FL AGGREGATION).

Dataset	Variant	Acc% (No-DP)	Acc% ($\varepsilon=10$)	Δ DP
CV	FedAvg (0%)	72.7	50.8	−21.9
	Ditto (reg.)	72.8	62.0	−10.8
	HPFL (head)	72.7	69.9	−2.8
	HPFL-Full (100%)	72.7	72.7	−0.0
PTB-XL	FedAvg (0%)	90.8	90.3	−0.5
	Ditto (reg.)	90.8	90.0	−0.8
	HPFL (head)	90.6	89.7	−0.9
	HPFL-Full (100%)	78.1	78.1	−0.0
BC	FedAvg (0%)	54.5	81.2	+26.7
	Ditto (reg.)	54.5	74.6	+20.1
	HPFL (head)	60.6	64.7	+4.1
	HPFL-Full (100%)	60.6	60.6	+0.0

Key findings: (1) Without DP, all variants perform similarly on CV/PTB-XL, confirming that personalization matters primarily under stress. (2) Under DP, the optimal personalization depth is task-dependent: HPFL-Full dominates on simple binary CV (+21.9pp over FedAvg), while PTB-XL (5-class) requires FL collaboration (HPFL-Full −12.2pp). (3) HPFL’s head-only personalization (~ 2 –6% params) is the robust “sweet spot” across all datasets. (4) On BC, DP acts as implicit regularization: FedAvg+DP *improves* over FedAvg without DP (+26.7pp), as Gaussian noise prevents overfitting on the small

569-sample dataset.

TABLE S-XLII

COMMUNICATION COSTS FOR TABULAR FL (HEALTHCAREMLP, FLOAT32). PER-ROUND BANDWIDTH INCLUDES UPLOAD (POTENTIALLY SPARSE WITH TOP- k) AND DOWNLOAD (FULL MODEL BROADCAST) FOR $K=5$ CLIENTS. HPFL SAVES $\sim 2\text{--}6\%$ BY KEEPING CLASSIFIER HEADS LOCAL.

Dataset	Algo	Params	KB/round	Total MB
CV	FedAvg (NoTopK)	2,914	113.8	3.33
	HPFL (NoTopK)	2,848	111.3	3.26
	FedAvg (Top-1%)	2,914	114.6 [†]	1.68
	HPFL (Top-1%)	2,848	112.4 [†]	1.65
PTB-XL	FedAvg (NoTopK)	2,885	112.7	3.30
	HPFL (NoTopK)	2,720	106.3	3.11
	FedAvg (Top-1%)	2,885	113.8 [†]	1.67
	HPFL (Top-1%)	2,720	107.2 [†]	1.57
BC	FedAvg (NoTopK)	4,130	161.3	4.73
	HPFL (Top-1%)	4,064	160.2 [†]	2.35

[†]Upload is 99% sparse (Top-1%), download is full model. Total MB for 30 rounds $\times$ 5 clients. Tabular FL communication is negligible (<5 MB total) even without compression—orders of magnitude below imaging (44.7 MB/round for ResNet-18). HPFL + Top-1% achieves 52% total reduction vs. FedAvg without compression.

TABLE S-XLIII

CONVERGENCE EFFICIENCY ANALYSIS (TABLE S-EFFICIENCY). SEEDS10 (NO-DP, 5 SEEDS): EFFICIENCY RATIO = ACCURACY / TOTAL COMMUNICATION (MB). HPFL IS PARETO-OPTIMAL ON ALL DATASETS.

Dataset	Algo	Acc%	R-to-90%	CommMB	Eff
CV	FedAvg	70.8	5.4	2.78	0.25
	Ditto	83.6	1.2	2.78	0.30
	HPFL	83.6	1.2	2.28	0.37
PTB-XL	FedAvg	91.3	4.0	3.08	0.30
	Ditto	92.1	6.6	3.30	0.28
	HPFL	92.6	5.2	2.70	0.34
BC	FedAvg	57.3	4.6	4.41	0.13
	Ditto	71.1	3.0	2.68	0.27
	HPFL	71.1	3.6	3.41	0.21

R-to-90% = rounds to reach 90% of final accuracy. Eff = Acc/MB (higher is better). HPFL’s backbone-only communication yields up to 48% higher efficiency than FedAvg.

TABLE S-XLIV

PER-CLIENT EQUITY ANALYSIS (TABLE S-EQUITY). GINI COEFFICIENT (0=EQUAL, 1=UNEQUAL), WORST-CASE CLIENT ACCURACY, AND EQUITY GAP (MAX—MIN CLIENT) ACROSS DP CONDITIONS. ROWS SHOW NO-DP BASELINE AND DP IMPACT (Δ FROM NO-DP TO $\varepsilon=10$).

Dataset	Algo	Gini	Worst%	Gap%	Δ Gap(DP)
BC (No-DP)	FedAvg	0.374	11.1	76.9	−15.4
	Ditto	0.371	12.2	75.9	−26.7
	HPFL	0.149	47.9	41.1	+13.5
CV (No-DP)	FedAvg	0.053	62.1	16.7	+27.0
	Ditto	0.079	58.1	21.9	+20.4
	HPFL	0.046	77.0	19.1	+9.6
PTB-XL (No-DP)	FedAvg	0.013	85.8	6.8	+2.0
	Ditto	0.015	86.0	7.5	+1.2
	HPFL	0.017	85.8	8.7	+1.6

Δ Gap(DP) = change in equity gap from No-DP to $\varepsilon=10$. Negative values indicate DP *improves* equity (regularization on small datasets). HPFL has the smallest DP-induced equity degradation on Cardiovascular (+9.6pp vs. +27.0pp for FedAvg).

main paper (Table S-DatasetConfig). Non-IID partitioning uses Dirichlet $\alpha=0.3$.

1) *Fairness-Aware Aggregation*: Table S-LI compares five fairness-aware aggregation algorithms against standard FedAvg under IID and non-IID conditions (36 experiments total). Fairness is measured via per-client accuracy variance (Jain index: 1.0 = perfect equity, Gini coefficient: 0.0 = perfect equity) and worst-client accuracy.

2) *Continual Federated Learning*: Table S-LII evaluates three continual learning strategies for FL under concept drift scenarios (36 experiments). Each experiment trains on Task 1 (original distribution) for 15 rounds, then on Task 2 (shifted distribution) for 15 rounds. Drift scenarios: *no_drift* (identical distributions), *sudden_shift* (abrupt feature permutation at round 16), *gradual_shift* (incremental permutation over rounds 16–25). Backward Transfer (BWT) measures accuracy change

TABLE S-XLV

PRIVACY-UTILITY-FAIRNESS TRILEMMA (TABLE S-TRILEMMA). ACCURACY (%) AND JAIN FAIRNESS INDEX AT EACH PRIVACY LEVEL. ACCDROP = ACCURACY DROP FROM NO-DP BASELINE. HPFL HAS THE FLATTEST TRADE-OFF CURVE ACROSS ALL DATASETS.

Dataset	Algo	No-DP	$\epsilon=1$	$\epsilon=10$	$\epsilon=50$	AccDrop($\epsilon=10$)
CV	FedAvg	69.5	56.6	62.3	73.2	-7.2
	Ditto	67.6	49.2	67.5	72.6	-0.1
	HPFL	85.2	77.1	84.8	85.7	-0.5
PTB-XL	FedAvg	90.2	29.6	88.0	90.7	-2.2
	Ditto	90.8	33.1	89.9	90.9	-0.8
	HPFL	90.7	74.2	90.0	90.9	-0.7
BC	FedAvg	48.4	46.6	51.0	55.9	+2.7 [†]
	Ditto	48.3	47.2	43.9	53.2	-4.4
	HPFL	73.2	72.0	66.6	74.1	-6.5

[†]Positive AccDrop indicates DP *improves* accuracy (regularization effect on small dataset). Marginal accuracy returns beyond $\epsilon=10$ are $<0.002\%/ \epsilon$ for all algorithms, establishing $\epsilon=10$ as the practical operating point.

TABLE S-XLVI

LEARNING RATE SENSITIVITY (TABLE S-LRSENSITIVITY). FINAL ACCURACY (%) ACROSS 3 LEARNING RATES PER DATASET. RANGE = MAX-MIN ACCURACY ACROSS LRS. LOWER RANGE INDICATES MORE ROBUST ALGORITHM. SEED = 42, NO-DP, 30 ROUNDS.

Dataset	Algo	LR _{low}	LR _{base}	LR _{high}	Range (pp)
CV	FedAvg	63.4	72.7*	73.0	9.6
	Ditto	63.5	72.8*	73.2	9.7
	HPFL	63.2	72.7*	73.2	10.0
PTB-XL	FedAvg	89.1	90.8*	91.2	2.1
	Ditto	89.1	90.8*	91.4	2.3
	HPFL	82.0	90.6*	91.2	9.3
BC	FedAvg	54.5	54.5*	95.7	41.2
	Ditto	54.5	54.5*	95.7	41.2
	HPFL	85.6	60.6*	89.1	28.5

* = base learning rate (CV: 0.01, PTB-XL: 0.005, BC: 0.005). LR_{low}/LR_{high} are $0.1\times$ and $4-5\times$ the base rate. HPFL is the most robust on Breast Cancer (28.5pp range vs. 41.2pp for FedAvg/Ditto), while all algorithms are comparably stable on Cardiovascular (~ 10 pp range).

on Task 1 data after Task 2 training: positive BWT indicates knowledge preservation, negative BWT indicates catastrophic forgetting.

3) *Extended Personalization Comparison*: Table S-LIII extends the personalization evaluation from the main paper (which covered Ditto and HPFL) with three additional algorithms: Per-FedAvg [34] (MAML-based meta-learning), Fed-Per (shared base + personalized head), and APFL (Adaptive Personalized FL with mixing coefficient α). Evaluation covers IID and two non-IID conditions (Dirichlet $\alpha \in \{0.1, 0.5\}$) across 45 experiments.

4) *Federated Unlearning (Right to Be Forgotten)*: Table S-LIV evaluates three federated unlearning methods for GDPR Article 17 “right to erasure” and EHDS Article 71 compliance (18 experiments). Each experiment first trains a baseline model with all clients, then removes 1 or 2 clients’ contributions. The key metrics are post-unlearning accuracy (should remain close to pre-unlearning) and computational cost relative to full retraining from scratch.

G. Client Selection and Gradient Inversion Privacy Validation

This section presents experimental validation of the framework’s client selection strategies and gradient inversion privacy analysis (87 experiments total).

1) *Client Selection Strategy Comparison*: We evaluate six client selection strategies across three EHDS-relevant datasets

(Cardiovascular, PTB-XL, Breast Cancer) with participation rate 60%, two federation scales ($K \in \{5, 10\}$; Breast Cancer $K=5$ only due to limited size), and 2 seeds per configuration (60 experiments total).

2) *Gradient Inversion Privacy Analysis*: We evaluate Deep Leakage from Gradients (DLG) [21] attacks against the FL framework under three differential privacy levels ($\epsilon \in \{\infty, 10, 1\}$) across three datasets with 3 seeds each (27 experiments). The attack uses L-BFGS optimization to reconstruct training data from shared gradients. We measure reconstruction quality via MSE, cosine similarity, and Pearson correlation, classifying results as SUCCESSFUL ($\text{MSE} < 0.1$), PARTIAL ($0.1 \leq \text{MSE} \leq 1.0$), or FAILED ($\text{MSE} > 1.0$).

H. Extended Federated Learning Validation

This section presents experimental validation across six additional tiers (234 experiments total), covering dataset coverage expansion, hierarchical FL, asynchronous FL, multi-task FL, extended Byzantine robustness, and Shapley-based client valuation.

1) *Dataset Coverage Expansion*: We extend the framework evaluation to three additional clinical datasets: Stroke Prediction (5,110 samples, 10 features, 4.9% positive rate), Chronic Kidney Disease (399 samples, 24 features), and Cirrhosis Patient Survival (418 samples, 18 features). Each dataset is tested with 5 FL algorithms (FedAvg, Ditto, HPFL, FedProx,

TABLE S-XLVII
DATA EFFICIENCY: ACCURACY (%) AT SUBSAMPLED TRAINING FRACTIONS

Dataset	Algorithm	25%	50%	75%	100% [†]
Cardiovascular	FedAvg	69.9	72.5	72.6	70.8
	Ditto	69.8	72.6	72.8	83.6
	HPFL	69.5	72.6	72.6	83.6
PTB-XL	FedAvg	90.0	90.8	91.2	91.3
	Ditto	90.0	90.7	91.2	92.1
	HPFL	84.6	90.7	91.2	92.6
Breast Cancer	FedAvg	54.5	54.5	54.5	57.3
	Ditto	54.5	54.5	54.5	71.1
	HPFL	60.6	60.6	60.6	71.1

[†]100% baseline from seeds10 checkpoint (mean of 5 seeds). Subsampled fractions use seed=42 only. On PTB-XL, FedAvg/Ditto retain 98.5% of full performance at 25% data. HPFL requires 50% data to converge but then matches competitors. Breast Cancer ($n=456$) is too small for meaningful subsampling analysis.

TABLE S-XLVIII
COMPOUND STRESS TEST: ACCURACY (%) UNDER SIMULTANEOUS NON-IID + DP + PARTIAL PARTICIPATION

Algorithm	Stress	Cardiovascular				PTB-XL			
		$\alpha=0.1$		$\alpha=0.5$		$\alpha=0.1$		$\alpha=0.5$	
		100%	60%	100%	60%	100%	60%	100%	60%
FedAvg	No-DP	71.9	53.1	72.6	59.2	86.3	89.0	89.5	90.0
	$\varepsilon=10$	63.6	52.9	63.2	56.4	77.8	81.1	90.3	91.0
Ditto	No-DP	71.9	53.2	72.1	59.7	86.4	89.3	89.4	89.8
	$\varepsilon=10$	67.2	52.7	61.4	54.4	59.2	76.0	82.8	87.8
HPFL	No-DP	91.5	91.2	87.3	86.5	85.6	86.0	91.9	91.3
	$\varepsilon=10$	91.1	91.1	86.3	86.2	92.2	92.0	91.5	89.2

Bold = best per column. Under triple stress ($\alpha=0.1$, $\varepsilon=10$, 60% participation), HPFL achieves 91.1% on Cardiovascular vs. FedAvg 52.9% (38.2pp gap). HPFL's worst-case (86.2%) exceeds FedAvg's best-case (72.6%). Ditto collapses on PTB-XL at $\alpha=0.1$, $\varepsilon=10$, 100% (59.2%, bold-italic). HPFL Jain ≥ 0.971 across all conditions.

TABLE S-XLIX
DATA HETEROGENEITY CHARACTERIZATION: PER-CLIENT STATISTICS UNDER DIRICHLET PARTITIONING

Dataset	Condition	Label Entropy	Class Imb. Ratio	Feature KL Div.
Cardiovascular	IID	1.000	1.0	0.000
	$\alpha=1.0$	0.727	6.6	0.008
	$\alpha=0.5$	0.510	25.0	0.014
	$\alpha=0.1$	0.325	1954.0	0.021
PTB-XL	IID	2.022	12.9	0.087
	$\alpha=1.0$	1.654	90.4	0.032
	$\alpha=0.5$	1.756	630.4	0.023
	$\alpha=0.1$	1.429	1416.0	0.071

Label Entropy: mean per-client Shannon entropy (bits). Class Imb. Ratio: mean max/min class ratio per client. Feature KL Div.: mean KL divergence of per-client feature distributions vs. global. At $\alpha=0.1$, Cardiovascular clients receive near-single-class data (entropy 0.325 vs. 1.0 IID) with class imbalance ratios up to 8,373:1. HPFL achieves 95.7% at $\alpha=0.1$ vs. FedAvg 66.0% (+29.7pp).

TABLE S-L
PEP TRIANGLE: PRIVACY–EQUITY–PERFORMANCE PROFILES

Dataset	Algorithm	Perf. (Acc %)	Equity (Jain)	Privacy (MIA AUC)
Cardiovascular	FedAvg	69.5	0.979	0.499
	Ditto	67.6	0.951	0.500
	HPFL	85.2	0.991	0.376*
PTB-XL	FedAvg	90.2	0.999	0.503
	Ditto	90.8	0.999	0.503
	HPFL	90.7	0.999	0.666
Breast Cancer	FedAvg	48.4	0.661	—
	Ditto	48.3	0.664	—
	HPFL	73.2	0.889	—

*MIA AUC < 0.5 indicates attack performs worse than random—HPFL’s personalization structure provides inherent membership privacy on Cardiovascular without requiring DP. On PTB-XL, HPFL’s elevated MIA AUC (0.666) reflects the privacy–personalization tension, fully mitigated by DP at $\epsilon=10$. Equity retention at $\epsilon=1$: HPFL 93.8% of Jain vs. FedAvg 78.4%.

TABLE S-LI
FAIRNESS-AWARE AGGREGATION COMPARISON (36 EXPERIMENTS, SEED 42, 30 ROUNDS). ALGORITHMS: Q-FEDAVG [53] ($q=1.0$), AFL [54] (AGNOSTIC FL), FEDMGDA+ [55] (MULTI-GRADIENT DESCENT), PROPFair (PROPORTIONAL FAIRNESS), FEDMINMAX (MINIMAX OPTIMIZATION). BEST PER-DATASET NON-IID IN **BOLD**.

Dataset	Algorithm	IID					Non-IID ($\alpha=0.3$)				
		Acc %	Worst %	Best %	Jain	Gini	Acc %	Worst %	Best %	Jain	Gini
CV	Standard	72.7	72.0	73.2	1.000	0.003	72.5	61.0	73.2	0.987	0.064
	q-FedAvg	72.7	72.0	73.2	1.000	0.003	73.0	68.2	73.2	0.997	0.029
	AFL	72.7	72.0	73.2	1.000	0.003	72.9	67.4	73.2	0.997	0.032
	FedMGDA+	72.7	72.1	73.2	1.000	0.004	64.2	33.1	73.3	0.894	0.188
	PropFair	72.7	72.0	73.2	1.000	0.003	54.6	2.6	73.2	0.693	0.361
	FedMinMax	72.7	72.0	73.2	1.000	0.003	71.2	65.0	73.0	0.994	0.043
PX	Standard	91.3	87.4	93.7	0.999	0.014	86.7	77.4	92.1	0.995	0.042
	q-FedAvg	91.2	88.3	93.5	1.000	0.011	91.8	89.0	93.7	0.999	0.012
	AFL	91.2	88.3	93.5	1.000	0.011	91.8	89.0	93.7	0.999	0.012
	FedMGDA+	80.3	76.8	83.7	0.999	0.019	80.9	67.1	90.1	0.988	0.061
	PropFair	91.0	87.4	93.5	0.999	0.013	91.5	89.0	93.5	0.999	0.013
	FedMinMax	88.3	85.3	90.1	0.999	0.016	92.0	89.9	93.5	1.000	0.011
BC	Standard	54.5	4.5	96.7	0.677	0.382	54.5	4.5	96.7	0.677	0.382
	q-FedAvg	60.0	22.7	96.7	0.771	0.294	60.0	22.7	96.7	0.771	0.294
	AFL	58.5	18.2	96.7	0.750	0.315	58.5	18.2	96.7	0.750	0.315
	FedMGDA+	45.5	0.0	96.7	0.593	0.458	45.5	0.0	96.7	0.593	0.458
	PropFair	48.2	0.0	96.7	0.709	0.348	48.2	0.0	96.7	0.709	0.348
	FedMinMax	85.7	59.1	96.7	0.963	0.101	85.7	59.1	96.7	0.963	0.101

Key findings: (1) Under IID, all fairness algorithms converge to identical accuracy on CV and PX, confirming that fairness reweighting has no effect when data is balanced. (2) Under non-IID, q-FedAvg achieves the best worst-case improvement on CV (+7.2pp over standard) while FedMinMax achieves the best overall accuracy on PX (92.0%) and dominates on BC (+31.2pp over standard). (3) FedMGDA+ and PropFair exhibit instability under non-IID: FedMGDA+ drops to 64.2% on CV (vs. 72.5% standard), and PropFair collapses to 54.6% with worst-client at 2.6%—the Frank-Wolfe and proportional optimization objectives conflict with non-IID gradients. (4) FedMinMax dramatically improves BC from 54.5% to 85.7% by preventing majority-class collapse. *EHDS implication*: fairness-aware aggregation is essential for cross-border federations with heterogeneous hospital sizes; FedMinMax is recommended for small-data settings, q-FedAvg for large-scale deployments.

TABLE S-LII

CONTINUAL FL UNDER CONCEPT DRIFT (36 EXPERIMENTS, SEED 42, 30 ROUNDS TOTAL: 15 TASK 1 + 15 TASK 2). METHODS: EWC [56] ($\lambda=1000$), REPLAY (BUFFER=200 SAMPLES/CLIENT), LwF [57] (TEMPERATURE=2.0). BWT = BACKWARD TRANSFER (POSITIVE = KNOWLEDGE PRESERVED, NEGATIVE = CATASTROPHIC FORGETTING).

DS	Method	No Drift				Sudden Shift				Gradual Shift			
		T1%	T2%	BWT	Jain	T1%	T2%	BWT	Jain	T1%	T2%	BWT	Jain
CV	Baseline	72.3	72.8	+0.5	1.000	72.2	72.6	+0.4	1.000	72.3	72.6	+0.3	1.000
	EWC	72.3	49.8	-22.4	1.000	72.2	49.8	-22.4	1.000	72.3	49.8	-22.4	1.000
	Replay	72.3	72.6	+0.3	1.000	72.2	72.5	+0.3	1.000	72.3	72.6	+0.3	1.000
	LwF	72.3	72.7	+0.4	1.000	72.2	72.5	+0.3	1.000	72.3	72.6	+0.3	1.000
PX	Baseline	88.1	88.1	-0.0	0.999	87.9	91.0	+3.1	0.999	88.0	90.1	+2.1	0.999
	EWC	88.1	87.8	-0.3	0.999	87.9	87.8	-0.1	0.999	88.0	88.0	+0.0	0.999
	Replay	88.1	88.1	-0.0	0.999	87.9	91.0	+3.1	0.999	88.0	90.0	+2.1	0.999
	LwF	88.1	88.1	+0.0	0.999	87.9	87.9	+0.1	0.999	88.0	88.0	+0.1	0.999
BC	Baseline	54.5	54.5	0.0	0.677	54.5	54.5	0.0	0.677	54.5	54.5	0.0	0.677
	EWC	54.5	54.5	0.0	0.677	54.5	54.5	0.0	0.677	54.5	54.5	0.0	0.677
	Replay	54.5	54.5	0.0	0.677	54.5	54.5	0.0	0.677	54.5	54.5	0.0	0.677
	LwF	54.5	54.5	0.0	0.677	54.5	54.5	0.0	0.677	54.5	54.5	0.0	0.677

Key findings: (1) **EWC is catastrophically harmful on CV**: the $\lambda=1000$ regularization penalty prevents the model from learning Task 2, dropping accuracy to 49.8% (near chance for binary). The Fisher Information Matrix from Task 1 over-constrains weight updates, blocking adaptation. A lower λ or per-layer adaptive regularization would be needed. (2) **Replay and LwF preserve knowledge effectively**: both maintain Task 1 accuracy after Task 2 training, with positive backward transfer on CV (+0.3–0.4pp). Replay achieves this by mixing stored Task 1 samples into Task 2 training batches; LwF by distilling Task 1 predictions. (3) **PTB-XL under sudden shift**: the baseline actually *improves* (+3.1pp BWT), suggesting the shifted distribution is complementary. EWC prevents this improvement by anchoring to Task 1 weights. (4) **BC is unresponsive**: all methods produce identical results (54.5%), indicating the 569-sample dataset is too small for the model to learn a discriminative representation under FL partitioning (4 clients, ~ 142 samples each). *EHDS implication*: for the long-term 2025–2031 EHDS deployment, Experience Replay or LwF are recommended for maintaining model quality under evolving data distributions; EWC requires careful hyperparameter tuning per dataset.

SCAFFOLD) under 3 conditions (IID, non-IID $\alpha=0.5$, non-IID + DP $\epsilon=10$) with 2 seeds (90 experiments).

2) *Hierarchical FL Topologies*: We evaluate three FL topologies—flat (standard), 2-tier (regional $\rightarrow$ global), and 3-tier (local $\rightarrow$ regional $\rightarrow$ global)—across 3 datasets with 2 seeds each (18 experiments). Hierarchical topologies reflect EHDS multi-level governance where hospitals aggregate within national HDABs before cross-border federation.

3) *Asynchronous FL Strategies*: We evaluate four asynchronous FL strategies—FedAsync, FedBuff (buffered), AsyncSGD, and FedAT (attentive)—across 3 datasets and 3 staleness levels ($s \in \{1, 3, 5\}$) representing 36 experiments. Asynchronous FL is critical for EHDS deployments where hospitals operate on different schedules and network latencies.

4) *Multi-Task FL*: We evaluate multi-task FL with two parameter sharing modes (hard, soft) and two task combinations (diagnosis+severity, diagnosis+readmission) across 3 datasets with 2 seeds (24 experiments). Multi-task FL enables EHDS analytics to produce multiple clinical outputs from a single federated model.

5) *Extended Byzantine Robustness*: We evaluate three Byzantine attacks (scale, label_flip, sign_flip) against four defense strategies (none, Krum, Trimmed Mean, Median) across 3 datasets (36 experiments). One out of 5 clients (Cardiovascular, PTB-XL) or 1 out of 4 clients (Breast Cancer) is Byzantine.

6) *Shapley Values and Client Valuation*: We compute exact Shapley values via exhaustive enumeration of all 2^K client

subsets ($K \in \{4, 5\}$) to quantify each client’s marginal contribution to the federation. Experiments span 3 datasets, 2 conditions (IID, non-IID $\alpha=0.5$), and 5 seeds (30 experiments, 15 rounds each).

TABLE S-LIII
EXTENDED PERSONALIZATION COMPARISON (45 EXPERIMENTS, SEED 42, 30 ROUNDS). FEDAVG BASELINE, DITTO ($\lambda=0.1$), PER-FEDAVG (5 INNER STEPS, META-LR=0.01), FEDPER (SHARED BASE + 1-LAYER HEAD), APFL (MIXING $\alpha=0.5$). BEST PER-DATASET/CONDITION IN **BOLD**.

DS	Method	IID			Non-IID ($\alpha=0.1$)			Non-IID ($\alpha=0.5$)		
		Acc%	Worst%	Jain	Acc%	Worst%	Jain	Acc%	Worst%	Jain
CV	FedAvg	72.7	72.0	1.000	71.3	49.6	0.965	72.5	61.0	0.987
	Ditto	72.7	72.1	1.000	91.6	72.9	0.986	87.1	72.8	0.983
	Per-FedAvg	62.0	60.9	1.000	55.8	16.8	0.746	59.0	18.8	0.812
	FedPer	72.8	72.0	1.000	91.5	72.8	0.986	87.1	72.8	0.982
	APFL	72.7	72.0	1.000	91.6	73.2	0.986	87.2	72.7	0.983
PX	FedAvg	91.3	87.4	0.999	77.5	60.7	0.967	86.7	77.4	0.995
	Ditto	89.9	84.8	0.998	94.8	87.6	0.999	93.3	90.1	0.999
	Per-FedAvg	17.4	12.9	0.980	32.0	4.8	0.770	21.1	15.6	0.943
	FedPer	89.2	83.6	0.997	94.4	86.6	0.998	93.1	89.9	0.999
	APFL	90.1	85.6	0.998	92.3	84.8	0.996	93.0	89.6	0.999
BC	FedAvg	54.5	4.5	0.677	54.5	4.5	0.677	54.5	4.5	0.677
	Ditto	60.6	0.0	0.734	60.6	0.0	0.734	60.6	0.0	0.734
	Per-FedAvg	45.5	0.0	0.593	45.5	0.0	0.593	45.5	0.0	0.593
	FedPer	60.6	0.0	0.734	60.6	0.0	0.734	60.6	0.0	0.734
	APFL	60.6	0.0	0.734	60.6	0.0	0.734	60.6	0.0	0.734

Key findings: (1) **Ditto, FedPer, and APFL form a high-performance cluster**: on CV non-IID ($\alpha=0.1$), all three achieve $\sim 91.6\%$ (+20.3pp over FedAvg). Their personalization mechanisms—L2 regularization (Ditto), split architecture (FedPer), and adaptive mixing (APFL)—converge to similar solutions on the compact MLP. (2) **Ditto leads on PTB-XL**: 94.8% at $\alpha=0.1$ (+17.3pp over FedAvg), followed by FedPer (94.4%) and APFL (92.3%). The 5-class task benefits from Ditto’s complete dual-model architecture. (3) **Per-FedAvg (MAML) fails**: 17.4% on PTB-XL IID (vs. 91.3% FedAvg) and 55.8% on CV non-IID (vs. 72.5% FedAvg). The MAML inner loop (5 steps, $\text{lr}=0.01$) is insufficient for convergence on these datasets; MAML’s sensitivity to inner-loop hyperparameters is a known challenge [34]. (4) **BC is personalization-insensitive**: IID and non-IID produce identical results for all personalization methods, confirming dataset saturation at 569 samples with 4 clients. *EHDS implication*: Ditto, FedPer, and APFL are all viable for EHDS personalization; the choice depends on deployment constraints (Ditto requires $2\times$ model storage, FedPer requires architecture modification, APFL is parameter-free). Per-FedAvg should be avoided without extensive hyperparameter tuning.

TABLE S-LIV

FEDERATED UNLEARNING COMPARISON (18 EXPERIMENTS, SEED 42, 30 ROUNDS PRE-TRAINING). METHODS: EXACT RETRAIN (GOLD STANDARD, RETRAIN FROM SCRATCH WITHOUT TARGET CLIENTS), GRADIENT ASCENT [59] (REVERSE GRADIENT CONTRIBUTION, 5% COST), FEDERASER [58] (SUBTRACT ACCUMULATED UPDATES + CALIBRATION, 37% COST). NEGATIVE “DROP” INDICATES POST-UNLEARNING ACCURACY *improves* (REMOVING HARMFUL CLIENTS). MIA = MEMBERSHIP INFERENCE VERIFICATION PASS.

Dataset	Scenario	Method	Pre-Acc%	Post-Acc%	Acc Drop (pp)	Cost	Removed Client Acc%	MIA Pass
CV	Remove 1 client	Exact Retrain	72.7	72.7	0.0	100%	73.2	✓
		Gradient Ascent	72.7	72.6	0.1	5%	73.1	✓
		FedEraser	72.7	72.6	0.1	37%	73.4	✓
	Remove 2 clients	Exact Retrain	72.7	72.7	−0.0	100%	72.9	✓
		Gradient Ascent	72.7	72.5	0.2	5%	72.9	✓
		FedEraser	72.7	72.6	0.1	37%	72.5	✓
PX	Remove 1 client	Exact Retrain	91.1	91.4	−0.3	100%	91.0	✓
		Gradient Ascent	91.1	91.2	−0.0	5%	91.0	✓
		FedEraser	91.1	91.2	−0.1	37%	91.0	✓
	Remove 2 clients	Exact Retrain	91.1	92.7	−1.6	100%	89.2	✓
		Gradient Ascent	91.1	92.4	−1.3	5%	89.2	✓
		FedEraser	91.1	92.3	−1.2	37%	89.2	✓
BC	Remove 1 client	Exact Retrain	59.1	61.6	−2.5	100%	33.3	✓
		Gradient Ascent	59.1	68.8	−9.7	5%	33.3	✓
		FedEraser	59.1	61.6	−2.5	37%	33.3	✓
	Remove 2 clients	Exact Retrain	59.1	90.1	−31.0	100%	18.9	✓
		Gradient Ascent	59.1	91.9	−32.8	5%	75.0	✓
		FedEraser	59.1	90.1	−31.0	37%	18.9	✓

Key findings: (1) **Gradient Ascent is the most cost-effective**: at only 5% of full retraining cost, it achieves comparable post-unlearning accuracy (≤ 0.2 pp drop on CV). This makes it practical for EHDS “right to erasure” scenarios where computational budget is limited. (2) **Unlearning can improve accuracy**: on BC, removing 2 clients raises accuracy from 59.1% to 90.1–91.9%, indicating that some clients’ data distributions were actively harmful to the global model. On PTB-XL, 2-client removal improves accuracy by 1.2–1.6pp. This has a surprising EHDS implication: exercising Article 17 opt-out rights may *benefit* the remaining federation. (3) **All methods pass MIA verification**: post-unlearning models show no residual membership signal for removed clients, confirming effective data erasure. (4) **FedEraser offers a middle ground**: 37% cost with accuracy matching exact retraining on CV and BC—its calibration rounds effectively recover from the subtracted contribution. (5) **CV is robust to client removal**: the large dataset (70K samples, 5 clients) absorbs the loss of 1–2 clients with negligible accuracy impact, suggesting that adequately-sized EHDS federations can handle opt-out requests without quality degradation.

EHDS implication: Gradient Ascent should be the default unlearning method for operational EHDS deployments, reserving Exact Retrain for audit verification scenarios. The framework’s unlearning module enables Article 17 compliance at 5–37% of the computational cost of full model retraining.

TABLE S-LV

CLIENT SELECTION STRATEGY COMPARISON: AVERAGE ACCURACY, JAIN FAIRNESS INDEX, AND GINI COEFFICIENT

Strategy	Cardiovascular			PTB-XL			Breast Cancer			Overall
	Acc	Jain	Gini	Acc	Jain	Gini	Acc	Jain	Gini	Avg Acc
random	65.76	0.904	0.166	87.48	0.981	0.055	56.15	0.688	0.359	72.52
loss_based	70.04	0.958	0.096	88.79	0.983	0.044	72.70	0.799	0.243	78.07
importance_sampling	66.61	0.922	0.150	88.09	0.984	0.049	74.16	0.808	0.230	76.71
resource_aware	70.37	0.964	0.093	86.18	0.980	0.066	69.81	0.764	0.273	76.58
fairness_aware	62.82	0.869	0.203	87.41	0.983	0.054	57.09	0.639	0.397	71.51
oort	70.85	0.963	0.096	88.99	0.983	0.044	72.70	0.799	0.243	78.47

Key findings: (1) **Oort achieves the best overall performance**: 78.47% average accuracy with Jain fairness index 0.938 and Gini coefficient 0.105, achieving the highest accuracy on both Cardiovascular (70.85%) and PTB-XL (88.99%). (2) **Loss-based selection is a close second**: 78.07% accuracy with nearly identical fairness (Jain 0.936, Gini 0.105), confirming that selecting clients by training loss jointly optimizes accuracy and equity. (3) **Fairness-aware selection is paradoxically the worst**: despite targeting balanced participation, it achieves the lowest accuracy (71.51%) and worst Jain index (0.868), driven by catastrophic K=10 behavior on Cardiovascular (Jain 0.658). This reveals that *participation* fairness $\neq$ *outcome* fairness—a critical distinction for EHDS deployment. (4) **Importance sampling achieves the best Breast Cancer accuracy** (74.16%, Jain 0.808), suggesting that variance-based selection benefits small, imbalanced datasets. *EHDS implication*: Oort or loss-based selection should be the default for cross-border federations, as they jointly optimize accuracy and equity without requiring explicit fairness constraints.

TABLE S-LVI
FEDERATION SCALE IMPACT: $K=5$ vs. $K=10$ (CARDIOVASCULAR + PTB-XL)

Scale	Avg Accuracy (%)	Avg Jain Index
$K=5$	79.27	0.978
$K=10$	76.30	0.934
Δ	-2.97pp	-0.044

Increasing federation size from 5 to 10 clients reduces average accuracy by 2.97pp and Jain fairness by 0.044 points. The degradation is strategy-dependent: fairness_aware suffers most (-4.85 pp accuracy, -0.085 Jain) while random is nearly unaffected (-0.98 pp). This is consistent with the client scalability findings in Section E-D.11: larger federations require proportionally more communication rounds for convergence, a critical consideration for EHDS deployments spanning many hospitals.

TABLE S-LVII
GRADIENT INVERSION ATTACK RESULTS BY DATASET AND DIFFERENTIAL PRIVACY LEVEL

Dataset	DP Level	Avg MSE	Avg Cosine Sim	Avg Correlation	Result	MSE Factor
Cardiovascular (dim=11)	no DP	0.63	0.369	0.200	3/3 PARTIAL	1.0×
	$\epsilon=10$	1,401.52	0.027	0.031	3/3 FAILED	2,215×
	$\epsilon=1$	63,374	0.029	-0.000	3/3 FAILED	100,133×
PTB-XL (dim=9)	no DP	0.10	0.816	0.683	2/3 SUCCESSFUL [†]	1.0×
	$\epsilon=10$	1,600.20	-0.035	0.001	3/3 FAILED	16,163×
	$\epsilon=1$	39,163 [‡]	-0.003	-0.002	3/3 FAILED	395,586×
Breast Cancer (dim=30)	no DP	1.07	0.389	0.160	1/3 SUCCESSFUL [†]	1.0×
	$\epsilon=10$	468.02	0.072	0.040	3/3 FAILED	438×
	$\epsilon=1$	131,662 [‡]	0.038	-0.023	3/3 FAILED	123,197×
Aggregated (all datasets)						
	no DP	0.60	0.525	0.348	3S/4P/2F	1.0×
	$\epsilon=10$	1,156.58	0.021	0.024	9/9 FAILED	1,927×
	$\epsilon=1$	56,761 [‡]	0.021	-0.008	9/9 FAILED	94,569×

[†]SUCCESSFUL: MSE < 0.1; [‡]Median reported (2 outlier runs with numerical divergence excluded from eps=1 average). Key findings: (1) **Without DP, gradient inversion achieves successful reconstruction on low-dimensional data:** PTB-XL (9 features) yields 2/3 successful attacks (MSE 0.10, cosine similarity 0.816), with Breast Cancer also partially vulnerable (1/3 successful). This confirms that gradient sharing without privacy protection poses a real data leakage risk in EHDS deployments [62]. (2) **PTB-XL is the most vulnerable:** its lower dimensionality yields the highest reconstruction fidelity (correlation 0.683), recovering clinically meaningful signal from shared gradients. (3) **Even moderate DP ($\epsilon=10$) completely prevents gradient inversion:** 100% attack failure rate (9/9 failed) with MSE increasing 1,927× and cosine similarity dropping from 0.525 to 0.021 (indistinguishable from random). This is the same $\epsilon=10$ that imposes <2pp accuracy cost—confirming that practical EHDS deployments get privacy “for free.” (4) **Strong DP ($\epsilon=1$) provides 94,569× MSE amplification** and causes optimizer divergence in 2/9 runs (MSE > 10^9), rendering the attack not merely ineffective but numerically catastrophic. (5) **Dimensionality inversely correlates with vulnerability:** PTB-XL (9 dim, MSE 0.10) > Cardiovascular (11 dim, MSE 0.63) > Breast Cancer (30 dim, MSE 1.07), suggesting that EHDS datasets with rich feature spaces have inherent structural protection against gradient-based reconstruction. *EHDS implication:* DP at $\epsilon=10$ should be mandatory for all cross-border gradient sharing—it eliminates gradient inversion risk at negligible accuracy cost.

TABLE S-LVIII
DATASET COVERAGE EXPANSION: MEAN ACCURACY (%) ACROSS SEEDS

Dataset	Condition	FedAvg	Ditto	HPFL	FedProx	SCAFFOLD
Stroke (5,110)	IID	94.9	94.9	94.9	94.9	94.9
	NonIID	95.3	88.5	91.3	95.3	95.3
	NonIID+DP	90.6	83.0	90.3	90.5	89.8
CKD (399)	IID	59.5	59.5	64.1	59.5	59.5
	NonIID	77.2	74.9	90.7	77.2	79.1
	NonIID+DP	70.3	88.8	91.7	70.9	71.5
Cirrhosis (418)	IID	62.2	62.1	62.1	62.2	62.2
	NonIID	64.6	62.2	83.6	64.6	64.0
	NonIID+DP	65.9	79.2	88.9	65.9	65.3

HPFL dominates on CKD and Cirrhosis across all conditions, with only 1.1pp DP degradation on Stroke (vs. 4–5pp for other algorithms). On CKD and Cirrhosis, DP paradoxically *improves* accuracy over NonIID alone for Ditto (+13.9pp on CKD, +17.0pp on Cirrhosis) and HPFL (+1.1pp on CKD, +5.4pp on Cirrhosis), consistent with DP noise acting as regularization on small, imbalanced datasets—an effect previously observed only on PTB-XL under compound stress. Under IID, FedAvg/FedProx/SCAFFOLD converge to identical solutions, confirming the convex loss landscape observation from main experiments. 11 instances of client accuracy = 0.0% occur under NonIID partitioning on CKD and Cirrhosis (Client 3), driven by pathological Dirichlet splits on small datasets. *EHDS implication*: HPFL’s personalized architecture is essential for small-sample clinical datasets (<500 samples), where standard algorithms fail to adapt to heterogeneous client distributions.

TABLE S-LIX
HIERARCHICAL FL: MEAN ACCURACY (%) BY TOPOLOGY

Dataset	Flat	2-Tier	3-Tier
Cardiovascular	72.52	72.52	72.56
PTB-XL	92.86	92.86	92.87
Breast Cancer	74.63	74.63	75.06

No meaningful difference between topologies (max range 0.44pp on Breast Cancer). Hierarchical aggregation layers are functionally equivalent to flat aggregation for compact MLP models, suggesting that EHDS multi-level governance can be implemented without accuracy penalty. Seed variance (Breast Cancer: 64.3% vs. 85.0% across seeds) dominates any topology effect, confirming that deployment decisions about HDAB aggregation hierarchy can be driven by governance requirements rather than accuracy considerations.

TABLE S-LX
ASYNCHRONOUS FL: ACCURACY (%) BY STRATEGY AND STALENESS

Strategy	Cardiovascular			PTB-XL			Breast Cancer		
	$s=1$	$s=3$	$s=5$	$s=1$	$s=3$	$s=5$	$s=1$	$s=3$	$s=5$
FedAsync	72.68	72.69	72.76	92.88	92.88	92.72	64.29	64.29	64.29
FedBuff	72.71	72.67	72.59	89.84	89.84	89.82	64.29	64.29	64.29
AsyncSGD	72.66	72.67	71.79	92.88	92.88	92.77	64.29	64.29	64.29
FedAT	72.73	72.74	72.73	92.88	92.88	92.88	64.29	64.29	64.29

FedAT (attentive aggregation) is the most staleness-resilient strategy, with 0.00pp degradation on PTB-XL and Breast Cancer across all staleness levels. FedBuff shows a persistent 3pp accuracy deficit on PTB-XL compared to other strategies (89.84% vs. 92.88%), though staleness degradation within FedBuff is minimal (0.02pp). AsyncSGD shows the largest degradation at high staleness (0.87pp on Cardiovascular). All strategies converge to identical results on Breast Cancer (64.29%), consistent with the dataset’s convergence behavior in other experiments. *EHDS implication*: FedAT should be the default for cross-border federations where hospital update latencies vary widely; FedBuff should be avoided on multi-class datasets (PTB-XL) where its buffered aggregation consistently underperforms.

TABLE S-LXI
MULTI-TASK FL: PRIMARY AND SECONDARY TASK ACCURACY (%)

Sharing	Task Combo	Cardiovascular		PTB-XL		Breast Cancer	
		Primary	Secondary	Primary	Secondary	Primary	Secondary
Hard	diagnosis+severity	72.50	50.30	92.85	49.89	74.63	55.09
Hard	diagnosis+readmission	72.55	60.40	92.80	68.19	74.18	64.85
Soft	diagnosis+severity	72.56	50.02	92.65	49.93	65.33	51.09
Soft	diagnosis+readmission	72.59	60.47	92.65	68.53	65.33	64.85

Hard and soft sharing produce equivalent primary task accuracy on Cardiovascular (<0.1 pp difference) and PTB-XL (<0.2 pp). On Breast Cancer, hard sharing outperforms soft by 9.3pp in primary accuracy, suggesting that small datasets benefit from the stronger inductive bias of fully shared representations. The severity secondary task performs near random ($\sim 50\%$) across all datasets—consistent with synthetic label generation—while the readmission secondary task achieves 60–68%. *EHDS implication*: Hard parameter sharing is preferable for multi-task EHDS analytics, providing equivalent or better accuracy with simpler implementation. Secondary task performance depends critically on task-specific signal quality, and EHDS data permits should specify minimum label quality requirements for auxiliary tasks.

TABLE S-LXII
BYZANTINE ROBUSTNESS: ACCURACY (%) BY ATTACK AND DEFENSE

Attack	None	Cardiovascular			None	PTB-XL			None	Breast Cancer		
		Krum	TrimMean	Median		Krum	TrimMean	Median		Krum	TrimMean	Median
Scale	49.84	72.78	72.74	72.75	91.10	89.16	92.51	92.53	64.29	35.71	35.71	35.71
Label flip	71.23	72.78	72.60	72.74	89.96	89.16	89.89	89.33	64.29	29.46	35.71	35.71
Sign flip	72.66	72.78	72.74	72.76	89.86	89.16	89.86	89.79	64.29	35.71	35.71	35.71

Byzantine defense effectiveness is strongly dataset-dependent. On Cardiovascular, all three defenses effectively neutralize all attacks (recovering to 72.7% from scale-attack collapse at 49.8%). On PTB-XL, Trimmed Mean is the most reliable defense (best or near-best on all attacks), while Krum returns a uniform 89.16% regardless of attack type—suggesting it deterministically selects the same client update every round. On Breast Cancer, **all defenses are catastrophically counterproductive**: accuracy drops from 64.29% (no defense) to 29–36% (with any defense), because robust aggregation rules exclude correct client updates on this small, imbalanced dataset. Scale attack is the most damaging without defense (Cardiovascular: -22.8 pp to near-random), while label and sign flips have minimal impact on well-trained models. *EHDS implication*: Trimmed Mean should be the default Byzantine defense for multi-hospital federations, but small-dataset deployments (<500 samples) should disable Byzantine defenses entirely, as the robust aggregation assumption (majority honest) breaks down with few clients.

TABLE S-LXIII
SHAPLEY VALUES: MEAN CLIENT CONTRIBUTIONS (5 SEEDS, $\pm$ STD)

Dataset	Condition	Client 0	Client 1	Client 2	Client 3	Client 4	Coalition Acc
Cardiovascular	IID	0.045 $\pm$.001	0.046 $\pm$.001	0.045 $\pm$.001	0.045 $\pm$.001	0.045 $\pm$.001	72.66%
	NonIID	0.035 $\pm$.027	0.044 $\pm$.036	0.020 $\pm$.031	0.048 $\pm$.028	0.035 $\pm$.019	68.22%
PTB-XL	IID	0.175 $\pm$.008	0.178 $\pm$.010	0.161 $\pm$.012	0.107 $\pm$.018	0.078 $\pm$.018	89.90%
	NonIID	0.155 $\pm$.040	0.108 $\pm$.059	0.121 $\pm$.049	0.172 $\pm$.066	0.142 $\pm$.049	89.85%
Breast Cancer	IID	0.097 $\pm$.115	0.077 $\pm$.085	0.051 $\pm$.084	0.023 $\pm$.083	—	74.82%
	NonIID	0.059 $\pm$.142	0.025 $\pm$.098	0.053 $\pm$.120	0.068 $\pm$.127	—	70.44%

The IID/NonIID comparison reveals three key patterns. (1) **IID produces symmetric Shapley values**: on Cardiovascular with equal data shares (20% each), all clients contribute ~ 0.045 (range 0.0005), exactly as game-theoretic symmetry predicts. On PTB-XL, IID Shapley values are proportional to data quantity (Client 0 holds 40.6% of data and has the highest SV at 0.175). (2) **NonIID introduces large variance and negative contributions**: Cardiovascular NonIID SV dispersion is $59\times$ larger than IID (range 0.028 vs. 0.0005). 6/25 client-seed combinations produce negative Shapley values on Cardiovascular NonIID, and 9/20 on Breast Cancer NonIID—indicating clients whose biased data actively hurts the coalition. Breast Cancer IID also produces 1/20 negative SVs due to dataset noise ($n=112$). (3) **NonIID degrades coalition accuracy**: Cardiovascular drops 4.4pp (72.66% $\rightarrow$ 68.22%), Breast Cancer drops 4.4pp (74.82% $\rightarrow$ 70.44%), while PTB-XL is resilient (-0.05 pp). Shapley efficiency gaps are $< 2 \times 10^{-6}$ (exact enumeration). *EHDS implication*: Shapley-based incentive mechanisms should be mandatory for cross-border federations to detect and penalize clients with harmful data distributions. The high negative-SV rate under NonIID (up to 45% on small datasets) suggests that client data quality certification should precede federation participation.

7) *Cascade 7: Trustworthiness Detailed Results:* The following tables present the quantitative results from Cascade 7’s trustworthiness evaluation (243 experiments across 6 blocks, Cardiovascular/PTB_XL/Breast Cancer datasets, 2 seeds each).

Model Calibration under DP (Block M, 54 experiments). Table S-LXIV shows Expected Calibration Error (ECE) before and after temperature scaling across 3 DP levels. Without DP, models are well-calibrated ($ECE \leq 0.03$ on CV and PTB_XL). DP noise induces systematic overconfidence: at $\varepsilon=10$, ECE rises to 0.07–0.42, but temperature scaling recovers calibration effectively (e.g., HPFL on CV: 0.215→0.030). At $\varepsilon=1$, overconfidence saturates (ratio ≈ 1.0) and temperature scaling provides negligible improvement.

Conformal Prediction under FL and DP (Block N, 36 experiments). Table S-LXV reports conformal prediction coverage and prediction set sizes. At target 90% coverage, noDP and $\varepsilon=10$ models maintain valid coverage with compact prediction sets (1.0–1.6 average set size). However, $\varepsilon=1$ produces *vacuous predictions*: 100% coverage achieved by including *all* classes (set size=2.0 for binary tasks, 5.0 for 5-class PTB_XL), rendering predictions clinically useless.

Feature Attribution Stability (Block O, 36 experiments).

Table S-LXVI reports cross-client Spearman rank correlation of feature importance vectors. Under IID, clients agree strongly on feature rankings ($\rho > 0.80$ on CV and PTB_XL). NonIID + DP destroys this agreement for Cardiovascular ($\rho \approx 0$), while PTB_XL retains moderate consistency ($\rho \approx 0.60$), suggesting that feature salience is more robust in high-dimensional signal spaces.

Demographic Fairness (Block P, 36 experiments). Table S-LXVII reports equalized odds gaps (TPR and FPR differences between demographic groups) on Cardiovascular disease prediction. Gender fairness is generally acceptable under IID (< 0.03 gap), but NonIID + DP amplifies disparities (FedAvg NonIID+ $\varepsilon=10$: TPR gap 0.19). Age fairness is consistently more problematic: demographic parity gaps reach 0.33–0.45 across all conditions, driven by inherent prevalence differences between age groups.

Concept Drift Robustness (Block Q, 36 experiments). Table S-LXVIII reports accuracy before and after simulated distribution shifts on Cardiovascular. Mild drift causes negligible degradation ($< 0.6pp$), while severe drift induces a uniform $\sim 13pp$ drop across all algorithms. Critically, *neither retraining nor EMA adaptation provides meaningful recovery* (max recovery $< 0.3pp$), indicating that single-round post-hoc adaptation is insufficient for severe covariate shift in FL—multi-round continual adaptation strategies (cf. Cascade 4 Block B) are required.

DP Composition for Sequential Studies (Block R, 45 experiments). Table S-LXIX compares naïve, advanced, and Rényi DP (RDP) composition across 1–20 sequential studies. Naïve and advanced composition yield identical (linear) budgets in all tested configurations. RDP provides strictly sub-linear composition, with the improvement ratio increasing with the number of studies and with tighter per-study budgets. At $\varepsilon=1$ per study, RDP enables 20 sequential studies at a total budget of 4.86 (vs. naïve 20.0), a $4.12\times$ improvement. This result is dataset-independent (composition depends only on noise parameters).

8) *Cascade 10: Detailed Results:* The following tables present the key quantitative results from Cascade 10’s clinical imbalance deep-dive (384 experiments, 0 errors, 22 min 16 sec).

DP as Anti-Collapse Regularizer (Block AC). Table S-LXX shows the collapse rate (fraction of experiments producing $F1=0.0$) across all 6 experimental conditions. The Block AA conditions (no DP) are from Cascade 9; the Block AC conditions (with DP) are from Cascade 10.

Mitigation Strategy Comparison (Block AD). Table S-LXXI reports the best F1 score achieved per dataset under each loss function, compared to the standard (unweighted) baseline from Cascade 9.

Ditto Local Epochs Effect (Block AE). Table S-LXXII shows the monotonic improvement of Ditto’s minority-class F1 as local epochs increase from 3 (Cascade 9 baseline) to 10.

Threshold Rescue (Block TH). Of the 20 collapsed models ($F1=0.0$) from Cascade 9 Block AA, all 20 produce threshold-

TABLE S-LXIV
MODEL CALIBRATION: ECE BEFORE AND AFTER TEMPERATURE SCALING (BLOCK M, 54 EXPERIMENTS, 2 SEEDS). LOWER ECE IS BETTER.
OVERCONF. RATIO: FRACTION OF BINS WHERE PREDICTED CONFIDENCE EXCEEDS TRUE ACCURACY.

Dataset	Algorithm	No DP			$\epsilon=10$			$\epsilon=1$		
		ECE _{pre}	ECE _{post}	Overconf.	ECE _{pre}	ECE _{post}	Overconf.	ECE _{pre}	ECE _{post}	Overconf.
Cardiovascular	FedAvg	0.024	0.025	0.15	0.416	0.270	0.92	0.441	0.438	1.00
	Ditto	0.026	0.027	0.15	0.299	0.120	0.82	0.363	0.363	1.00
	HPFL	0.024	0.024	0.15	0.215	0.030	0.70	0.330	0.328	1.00
PTB_XL	FedAvg	0.028	0.012	0.11	0.129	0.053	0.88	0.151	0.154	1.00
	Ditto	0.026	0.013	0.10	0.070	0.019	0.75	0.095	0.093	1.00
	HPFL	0.028	0.011	0.11	0.119	0.043	0.90	0.140	0.142	1.00
Breast Cancer	FedAvg	0.169	0.098	0.00	0.156	0.163	0.94	0.178	0.213	1.00
	Ditto	0.168	0.101	0.00	0.157	0.124	0.88	0.188	0.223	1.00
	HPFL	0.168	0.087	0.00	0.162	0.146	0.81	0.166	0.192	1.00

Model Calibration: ECE Before and After Temperature Scaling




Fig. S-41. ECE before and after temperature scaling across DP levels on Cardiovascular and PTB-XL. All ECE values are near zero (<0.05), indicating that both original and temperature-scaled models are well calibrated under No-DP and $\epsilon=10$. The visually flat bars reflect this near-perfect calibration. Under $\epsilon=1$, models degrade to trivial predictions (see Table S-LXIV), where ECE becomes undefined due to single-class collapse.

TABLE S-LXV
CONFORMAL PREDICTION: COVERAGE AND AVERAGE PREDICTION SET SIZE AT TARGET 90% COVERAGE (BLOCK N, 36 EXPERIMENTS, 2 SEEDS).
SINGLETON FRACTION: PROPORTION OF PREDICTIONS CONTAINING EXACTLY ONE CLASS.

Dataset	Condition	No DP			$\epsilon=10$			$\epsilon=1$		
		Cov.	Set Size	Single.	Cov.	Set Size	Single.	Cov.	Set Size	Single.
Cardiovascular	IID	0.896	1.460	0.540	0.901	1.620	0.380	1.000	2.000	0.000
	NonIID	0.896	1.457	0.543	0.897	1.550	0.451	1.000	2.000	0.000
PTB_XL	IID	0.910	0.961	0.961	0.895	1.168	0.850	1.000	5.000	0.000
	NonIID	0.902	0.983	0.955	0.903	1.094	0.909	1.000	5.000	0.000
Breast Cancer	IID	0.911	1.277	0.723	0.866	1.429	0.429	0.902	1.500	0.500
	NonIID	0.929	1.364	0.637	0.921	1.248	0.752	1.000	2.000	0.000

Red-highlighted cells indicate vacuous predictions where the conformal set includes all possible classes, providing zero diagnostic value.

TABLE S-LXVI
CROSS-CLIENT FEATURE IMPORTANCE RANK CORRELATION (SPEARMAN ρ , BLOCK O, 36 EXPERIMENTS, 2 SEEDS). HIGHER IS BETTER.

Dataset	Algorithm	IID	NonIID+ $\varepsilon=10$
Cardiovascular	FedAvg	0.845	0.020
	Ditto	0.804	−0.041
	HPFL	0.850	−0.038
PTB_XL	FedAvg	0.851	0.606
	Ditto	0.837	0.560
	HPFL	0.861	0.616
Breast Cancer	FedAvg	0.234	0.204
	Ditto	0.255	0.170
	HPFL	0.243	0.141

TABLE S-LXVII
DEMOGRAPHIC FAIRNESS GAPS ON CARDIOVASCULAR (BLOCK P, 36 EXPERIMENTS, 2 SEEDS). GENDER: MALE/FEMALE. AGE: YOUNG/OLD. LOWER GAPS INDICATE FAIRER MODELS.

Algorithm	Cond.	DP	Gender Fairness				Age Fairness			
			Acc Gap	TPR Gap	FPR Gap	DP Gap	Acc Gap	TPR Gap	FPR Gap	DP Gap
FedAvg	IID	noDP	0.005	0.008	0.011	0.004	0.084	0.243	0.249	0.335
	IID	$\varepsilon=10$	0.014	0.008	0.026	0.010	0.050	0.097	0.270	0.222
	NonIID	noDP	0.011	0.040	0.022	0.029	0.075	0.242	0.216	0.314
	NonIID	$\varepsilon=10$	0.052	0.190	0.267	0.232	0.049	0.201	0.188	0.263
Ditto	IID	noDP	0.007	0.007	0.013	0.004	0.084	0.242	0.248	0.334
	IID	$\varepsilon=10$	0.010	0.017	0.024	0.015	0.089	0.181	0.347	0.335
	NonIID	noDP	0.013	0.016	0.010	0.004	0.051	0.181	0.202	0.322
	NonIID	$\varepsilon=10$	0.010	0.030	0.018	0.022	0.029	0.114	0.105	0.207
HPFL	IID	noDP	0.005	0.009	0.011	0.005	0.084	0.243	0.249	0.335
	IID	$\varepsilon=10$	0.019	0.035	0.045	0.035	0.111	0.247	0.367	0.394
	NonIID	noDP	0.011	0.039	0.020	0.028	0.059	0.238	0.196	0.315
	NonIID	$\varepsilon=10$	0.016	0.120	0.087	0.110	0.034	0.190	0.129	0.267

TABLE S-LXVIII
CONCEPT DRIFT ROBUSTNESS ON CARDIOVASCULAR (BLOCK Q, 36 EXPERIMENTS, 2 SEEDS). DROP: ACCURACY DECREASE FROM PRE-DRIFT. RECOVERY: ACCURACY GAIN FROM ADAPTATION.

Algorithm	Severity	Pre-Drift	Post-Drift	Drop	Best Recovery
FedAvg	Mild	72.5%	72.1%	0.4pp	<0.1pp (ema)
	Severe	72.5%	59.4%	13.1pp	0.1pp (ema)
Ditto	Mild	72.6%	72.1%	0.6pp	0.1pp (retrain)
	Severe	72.6%	59.4%	13.2pp	0.1pp (ema)
HPFL	Mild	72.5%	72.1%	0.4pp	<0.1pp (ema)
	Severe	72.5%	59.5%	13.1pp	0.1pp (ema)




Fig. S-42. Conformal prediction set size inflation under DP. Under No-DP and $\epsilon=10$, prediction sets are near-ideal (size ≈ 1.0), appearing at the baseline. Under $\epsilon=1$, sets should expand to 2.0 (binary) or 5.0 (PTB-XL), but the visualization saturates at the reference lines. See Table S-LXV for exact numerical values.




Fig. S-43. Majority-class collapse rate across experimental conditions. No-DP conditions (left) show 75–83% collapse; any DP level (right) reduces collapse to 0–3.3%, supporting the hypothesis of DP noise as implicit anti-collapse regularizer.




Fig. S-44. Loss function mitigation strategies for class imbalance. Weighted CE and Focal loss ($\gamma=2$) eliminate majority-class collapse across all three datasets; Weighted CE achieves the highest F1 on Stroke (0.39) and Cirrhosis (0.72).




Fig. S-45. Effect of local epochs on Ditto minority-class F1. Cirrhosis shows monotonic improvement (0.36 $\rightarrow$ 0.57 $\rightarrow$ 0.78 for NonIID+noDP); Stroke improves dramatically at ep=10 (0.00 $\rightarrow$ 0.36), suggesting more local personalization steps help overcome extreme imbalance.

TABLE S-LXIX
DP COMPOSITION: TOTAL ε AFTER k SEQUENTIAL STUDIES (BLOCK R, 45 EXPERIMENTS). RDP RATIO: NAÏVE/RDP IMPROVEMENT FACTOR.

Per-Study ε	Number of Sequential Studies (k)				
	$k=1$	$k=2$	$k=5$	$k=10$	$k=20$
<i>Naïve / Advanced Composition ($\varepsilon_{total} = k \cdot \varepsilon$)</i>					
$\varepsilon=1.0$	1.0	2.0	5.0	10.0	20.0
$\varepsilon=5.0$	5.0	10.0	25.0	50.0	100.0
$\varepsilon=10.0$	10.0	20.0	50.0	100.0	200.0
<i>RDP Composition (ε_{total})</i>					
$\varepsilon=1.0$	1.0	1.44	2.32	3.35	4.86
$\varepsilon=5.0$	5.0	8.13	13.74	21.74	32.79
$\varepsilon=10.0$	10.0	18.52	32.89	52.63	86.96
<i>RDP Improvement Ratio (naïve / RDP)</i>					
$\varepsilon=1.0$	1.00	1.39	2.15	2.99	4.12×
$\varepsilon=5.0$	1.00	1.23	1.82	2.30	3.05×
$\varepsilon=10.0$	1.00	1.08	1.52	1.90	2.30×

tuned F1>0 after sweeping the classification threshold from 0.05 to 0.95. For Cirrhosis, the rescue is clinically meaningful (mean threshold-tuned F1=0.648, optimal thresholds 0.30–0.45); for Stroke, the rescue is marginal (mean threshold-tuned F1=0.190, optimal thresholds concentrated at the minimum 0.05), indicating that the models have not learned sufficient minority-class signal.

9) *Main Paper Extended Content (v3)*: The following content provides extended details from the main paper that were condensed for the 8-page conference submission.

FL-EHDS Algorithm Catalogue (17 Algorithms). Table S-LXXIII lists all 17 algorithms implemented in the framework with their venues and key properties.

Threat Model. We consider three adversary classes. (i) *Honest-but-curious server*: the aggregation server follows the protocol but may attempt to infer patient information from model updates; central DP (Gaussian mechanism, $\varepsilon \in \{1,5,10,50\}$, $\delta=10^{-5}$) with Rényi composition bounds per-round leakage, while secure aggregation ensures the server

observes only the noised aggregate. (ii) *Malicious clients*: up to $f < n/3$ compromised institutions may submit poisoned updates; Byzantine-resilient aggregation (Krum, Trimmed Mean, Bulyan) provides robustness. (iii) *External attacker*: membership or attribute inference against published outputs is mitigated by Article 71 output filtering and HDAB permit-based access control. **Protected assets**: raw patient data (never leaves institutional boundaries), gradient updates (DP + secure aggregation), and the global model (HDAB permit-controlled). **Out of scope**: server-client collusion, side-channel attacks, and covert-channel leakage through model architecture.

Reference Implementation Details. A modular Python implementation is available as open-source software, designed following FAIR principles (findable via GitHub with DOI, accessible under MIT license, interoperable via PyTorch and FHIR R4 interfaces, reusable with comprehensive documentation). The codebase (~40K lines, 159 modules) provides: (1) orchestration modules implementing all 17 algorithms with RDP accounting and secure aggregation; (2) six Byzan-

TABLE S-LXX
MAJORITY-CLASS COLLAPSE RATE BY CONDITION (STROKE + CIRRHOSIS)

IID/NonIID	DP Level	Collapses	Total	Rate
IID	No DP	10	12	83.3%
NonIID	No DP	9	12	75.0%
NonIID	$\varepsilon=10$	1	12	8.3%
IID	$\varepsilon=10$	0	30	0.0%
IID	$\varepsilon=1$	0	30	0.0%
NonIID	$\varepsilon=1$	1	30	3.3%
No DP (any partition)		19	24	79.2%
Any DP (any partition)		2	102	2.0%

TABLE S-LXXI
BEST F1 BY LOSS TYPE AND DATASET (BLOCK AD, MEAN $\pm$ STD)

Dataset	Loss Type	Best Config	Mean F1	Collapses
Stroke	Standard (AA)	—	0.000	55.6%
	Weighted CE	Ditto NonIID+noDP	0.391 ± 0.072	0/45
	Focal	Ditto NonIID+eps10	0.334 ± 0.072	0/45
Cirrhosis	Standard (AA)	—	0.000–0.693	55.6%
	Weighted CE	Ditto NonIID+eps10	0.715 ± 0.102	2/45
	Focal	Ditto NonIID+eps10	0.712 ± 0.094	0/45
CDC Diabetes	Weighted CE	Ditto NonIID+noDP	0.572 ± 0.137	0/27
	Focal	HPFL NonIID+noDP	0.488 ± 0.206	0/27

TABLE S-LXXII
DITTO F1 BY LOCAL EPOCHS (NONIID, 5 SEEDS)

Dataset	Condition	ep=3 (baseline)	ep=5	ep=10
Cirrhosis	NonIID+noDP	0.364 ± 0.364	0.571 ± 0.213	0.785 ± 0.073
Cirrhosis	NonIID+eps10	0.693 ± 0.100	0.678 ± 0.078	0.729 ± 0.059
Stroke	NonIID+noDP	0.000 ± 0.000	0.025 ± 0.049	0.357 ± 0.247
Stroke	NonIID+eps10	0.239 ± 0.188	0.353 ± 0.130	0.335 ± 0.148




Fig. S-46. Post-hoc threshold rescue of F1=0.0 collapsed models. Cirrhosis models (blue) achieve clinically meaningful rescue (mean threshold-tuned F1=0.648); Stroke models (red) show marginal improvement (mean threshold-tuned F1=0.190), indicating insufficient minority-class signal.

TABLE S-LXXIII
FL-EHDS ALGORITHM CATALOGUE

Algorithm	Venue	Category	Key Property
FedAvg	AISTATS'17	Baseline	Weighted avg.
FedProx	MLSys'20	Non-IID	Proximal reg.
SCAFFOLD	ICML'20	Non-IID	Variance red.
FedNova	NeurIPS'20	Non-IID	Normalized avg.
FedDyn	ICLR'21	Non-IID	Dynamic reg.
FedAdam	ICLR'21	Adaptive	Server momentum
FedYogi	ICLR'21	Adaptive	Sparse stability
FedAdagrad	ICLR'21	Adaptive	Grad. accum.
Ditto	ICML'21	Personal.	Dual models
Per-FedAvg	NeurIPS'20	Personal.	MAML-based
FedLC	ICML'22	Label skew	Logit calibration
FedSAM	ICML'22	Generalize	Flat minima
FedDecorr	ICLR'23	Represent.	Decorrelation
FedSpeed	ICLR'23	Efficiency	Fewer rounds
FedExp	ICLR'23	Server-side	POCS step size
FedLESAM	ICML'24	Generalize	Global SAM
HPFL	ICLR'25	Personal.	Local classif.

Bold: newly added algorithms (2024–2025). All 17 implemented in the open-source reference implementation.

tine resilience methods; (3) data holder utilities for adaptive training and FHIR R4 preprocessing; (4) a Streamlit-based dashboard for interactive FL training, EHDS governance workflow, and real-time monitoring; (5) a professional terminal UI with 11 specialized screens; (6) reproducible benchmark suite generating all experimental results. HDAB integration, data permit lifecycle, and Article 71 opt-out compliance are implemented as fully functional simulation backends demonstrating the complete workflow (OAuth2/mTLS authentication, permit CRUD, cross-border coordination, LRU-cached registry lookups, scope-granular filtering). These modules are architecturally designed for production binding—requiring only configuration changes (endpoint URLs, mTLS certificates), not

architectural modifications.

Extended Algorithm Comparison (Heart Disease, Diabetes). Table S-LXXIV provides supplementary validation on Heart Disease and Diabetes with 5 foundational algorithms (3 seeds). The main paper Table VII extends Heart Disease to 10 seeds with 5 approaches including HPFL and Local-Only baselines.

Deployment Implications (Extended). Sixteen deployment implications emerge from our experimental evaluation. *First*, the up to 18.7pp accuracy gap demonstrates that EHDS SPE configurations cannot treat FL as a black box; algorithm selection must be part of data permit specification. *Sec-*

TABLE S-LXXIV
FL ALGORITHM COMPARISON: HEART DISEASE AND DIABETES

Algo.	Heart Disease (4 hosp.)			Diabetes (5 hosp.)		
	Acc.	F1	AUC	Acc.	F1	AUC
FedAvg	62.5±8.0	.736±.06	.834±.03	68.1±4.2	.259±.01	.643±.00
FedProx	61.7±8.0	.732±.05	.834±.03	71.0±6.3	.254±.01	.638±.00
SCAFFOLD	66.3±5.1	.667±.02	.791±.05	11.2±0.0	.201±.00	.514±.00
FedNova	56.4±5.4	.711±.04	.831±.03	13.0±0.9	.203±.00	.636±.00
Ditto	75.1±2.0	.761±.03	.826±.01	71.7±0.2	.262±.00	.643±.00

20 rounds, 3 local epochs. Heart Disease: natural non-IID. Diabetes: Dirichlet $\alpha=0.5$. Mean $\pm$ std over 3 seeds.

ond, SCAFFOLD and FedNova fail catastrophically on class-imbalanced tabular data (known failure modes exacerbated by clinical-scale class ratios); EHDS delegated acts should mandate algorithm validation protocols. *Third*, the success of personalized FL is confirmed with strong statistical evidence ($p < 0.001$, pooled Wilcoxon). *Fourth*, the convergence of 12+ algorithms confirms deployment simplification. *Fifth*, aggregate accuracy is insufficient for clinical evaluation (DEI reveals catastrophic per-class failures). *Sixth*, Byzantine defenses and DP are synergistic. *Seventh*, the privacy–personalization tension (HPFL MIA vulnerability) requires explicit EHDS policy. *Eighth*, personalization is necessary under strict privacy ($\epsilon=2$: FedAvg 41%, HPFL 81%). *Ninth*, Non-IID $\times$ DP interaction is sub-additive for HPFL but super-additive for global algorithms. *Tenth*, Top- k and DP synergy benefits HPFL uniquely. *Eleventh*, optimal personalization depth is task-complexity-dependent. *Twelfth*, HPFL’s statistical advantage strengthens under DP stress. *Thirteenth*, HPFL is Pareto-optimal across all tabular datasets. *Fourteenth*, DP has asymmetric equity effects across dataset sizes. *Fifteenth*, HPFL uniquely survives compound stress (worst-case 86.2% exceeds FedAvg best-case 72.6%). *Sixteenth*, FL with 50% data achieves within 1% of full performance.

Practical Deployment Scenario. A cardiology consortium of 50 hospitals across three Member States (DE, IT, NL) would: (1) apply for data permits through respective HDABs specifying “Ditto” or “HPFL” as the FL algorithm with $\epsilon=10$; (2) deploy FL-EHDS data holder components with FHIR R4 preprocessing; (3) execute 25 rounds of federated training within the SPE (~ 1 MB total communication for tabular models); (4) each hospital retains its personalized local model for clinical decision support. The process requires no raw data transfer, achieves 82.3% accuracy on cardiovascular risk prediction under full $\epsilon=10$ privacy, and preserves cross-site fairness (Jain 0.951).

10) Extended Standalone Experiment Results: The following tables present detailed results from 787 additional standalone experiments spanning five independent evaluation tracks: centralized baselines (9 exp.), server-side optimizers

(108 exp.), 10-seed statistical robustness (240 exp.), CDC Diabetes full evaluation (214 exp.), and vertical FL extended (216 exp.). All experiments use the same atomic checkpointing infrastructure and deterministic seeding described in Section D-A.

Centralized Baselines (RQ#37, 9 experiments). Table S-LXXV provides the centralized upper-bound reference for three core datasets, trained with standard PyTorch (Adam optimizer, $\text{lr}=0.01$, 50 epochs, batch size 128). The centralized-FL gap is remarkably small: $\leq 0.5\text{pp}$ on PTB-XL and Cardiovascular, confirming that FL achieves near-optimal performance while preserving data sovereignty. CDC Diabetes exhibits low F1 (0.116) even under centralized training, confirming that the 86% majority-class imbalance is a data-level challenge, not an FL artifact.

Server-Side Optimizer Comparison (RQ#36, 108 experiments). Table S-LXXVI evaluates six server-side adaptive optimizers across three datasets under IID and NonIID conditions ($K=5$ clients, 20 rounds, 3 local epochs, 3 seeds). FedAdam, FedYogi, and FedAdagrad [31] produce statistically indistinguishable results on all three datasets ($p>0.9$, Kruskal–Wallis on IID Cardiovascular), confirming that the nearly convex loss landscape of compact clinical models ($\sim 10\text{K}$ parameters) renders server-side optimizer choice irrelevant. FedDyn achieves the highest PTB-XL accuracy (92.5%), matching the best cascade results. SCAFFOLD consistently underperforms: 62.8% on Cardiovascular IID (-10.2pp vs. FedAdagrad) and 78.2% on PTB-XL (-13.9pp), consistent with its known failure mode on heterogeneous tabular data. FedNova is competitive under IID but collapses under NonIID on CDC Diabetes (63.3%).

CDC Diabetes Population-Scale Evaluation (RQ#38, 214 experiments). Tables S-LXXVII–S-LXXIX present the complete evaluation on CDC Diabetes (253,680 samples, 21 features, binary classification). This is the largest dataset in our evaluation, testing FL scalability to population-health analytics.

Algorithm comparison (Block A, 70 experiments). Table S-LXXVII shows 7 algorithms under IID and NonIID ($\alpha=0.5$).

TABLE S-LXXV
CENTRALIZED TRAINING BASELINES (UPPER-BOUND REFERENCE). MEAN $\pm$ STD OVER 3 SEEDS (42, 123, 456). ADAM OPTIMIZER, LR=0.01, 50 EPOCHS.

Dataset	Acc. (%)	F1	Precision	Recall
PTB-XL (5-class, 21,799)	92.6 $\pm$ 0.4	.906 $\pm$.006	.928 $\pm$.010	.895 $\pm$.010
Cardiovascular (70K)	73.5 $\pm$ 0.3	.720 $\pm$.006	.767 $\pm$.001	.678 $\pm$.011
CDC Diabetes (253K)	86.5 $\pm$ 0.1	.116 $\pm$.034	.639 $\pm$.023	.065 $\pm$.021

CDC Diabetes F1 reflects 86% majority-class imbalance—centralized training also fails to learn the minority class, confirming this is a data-level challenge.

TABLE S-LXXVI
SERVER-SIDE OPTIMIZER COMPARISON (108 EXPERIMENTS). $K=5$ CLIENTS, 20 ROUNDS, 3 LOCAL EPOCHS. MEAN $\pm$ STD OVER 3 SEEDS $\times$ IID/NonIID. BEST IID ACCURACY PER DATASET IN **BOLD**.

Algorithm	PTB-XL (5-class)				Cardiovascular				CDC Diabetes (253K)			
	IID		NonIID		IID		NonIID		IID		NonIID	
	Acc	F1	Acc	F1	Acc	F1	Acc	F1	Acc	F1	Acc	F1
FedAdam	92.1 $\pm$ 0.4	.894	92.1 $\pm$ 0.4	.894	72.5 $\pm$ 0.5	.717	69.0 $\pm$ 5.3	.708	73.1 $\pm$ 0.7	.445	79.4 $\pm$ 4.9	.301
FedYogi	92.1 $\pm$ 0.4	.893	92.1 $\pm$ 0.4	.893	72.7 $\pm$ 0.4	.721	68.8 $\pm$ 5.3	.700	73.2 $\pm$ 0.4	.446	79.3 $\pm$ 4.9	.301
FedAdagrad	91.9 $\pm$ 0.2	.892	91.9 $\pm$ 0.2	.892	73.0$\pm$0.2	.709	68.7 $\pm$ 4.4	.720	73.1 $\pm$ 2.9	.445	80.1 $\pm$ 4.4	.303
FedDyn	92.5$\pm$0.2	.898	92.5$\pm$0.2	.898	72.6 $\pm$ 0.1	.716	69.4 $\pm$ 2.7	.700	70.7 $\pm$ 2.5	.432	79.2 $\pm$ 4.9	.299
FedNova	92.2 $\pm$ 0.2	.894	92.2 $\pm$ 0.2	.894	72.8 $\pm$ 0.2	.718	64.0 $\pm$ 9.9	.673	71.3 $\pm$ 0.9	.438	63.3 $\pm$ 18.6	.369
SCAFFOLD	78.2 $\pm$ 6.6	.562	78.2 $\pm$ 6.6	.562	62.8 $\pm$ 3.5	.710	58.7 $\pm$ 7.2	.357	58.1 $\pm$ 31.5	.226	86.2$\pm$0.0	.005

SCAFFOLD’s 86.2% on CDC NonIID reflects majority-class prediction (F1=0.005). Server lr=0.1, $\beta_1=0.9$, $\beta_2=0.99$, $\tau=1e-3$ for adaptive methods.

Five algorithms (FedAvg, FedProx, Ditto, FedExp, FedLE-SAM) converge to identical IID performance ($\sim 70.1\%$), consistent with Finding #5 on compact models. HPFL significantly underperforms (57.2% IID), as its classifier-head architecture is counterproductive when class imbalance (86% majority) dominates data heterogeneity—the personalized heads overfit to the imbalanced local distributions rather than learning useful decision boundaries. Under NonIID, all algorithms improve markedly (70–81%), as Dirichlet partitioning creates clients with less extreme imbalance ratios than the IID case.

DP ablation (Block B, 72 experiments). Table S-LXXVIII reveals a striking DP regularization effect on CDC Diabetes. FedAvg accuracy *improves* from 70.1% (no DP, Block A) to 85.8% at $\varepsilon=5$ (+15.7pp IID), consistent with the DP-as-regularizer phenomenon observed on Breast Cancer. However, this “improvement” is illusory: F1 drops from 0.429 to 0.013, meaning DP noise pushes the model toward majority-class prediction. Ditto collapses under strict DP ($\varepsilon=1$: 57.9% IID, 38.4% NonIID), while HPFL is catastrophically affected (18.0% at $\varepsilon=1$ IID).

Client scalability (Block C, 72 experiments). Table S-LXXIX evaluates scaling from $K=3$ to $K=20$ hospitals. FedAvg and Ditto maintain stable performance (≤ 1 pp variation under IID), confirming that CDC Diabetes scales gracefully. Under NonIID, increasing K paradoxically *improves* accuracy (FedAvg: 78.0% at $K=3 \rightarrow 84.2\%$ at $K=20$), as more clients dilute the per-client label skew. HPFL degrades significantly with scale (63.5% at $K=3 \rightarrow 55.6\%$ at $K=10$ under IID), as more classifier heads amplify the overfitting problem on imbalanced

data.

Vertical FL Extended Evaluation (RQ#39, 216 experiments). Table S-LXXX presents the extended vertical FL evaluation across three datasets, 2–4 parties, three overlap rates, and DP ablation. PTB-XL (binarized: NORM vs. AB-NORMAL, 43%/57% class balance) achieves near-perfect accuracy (99.3%) with 2 parties and no DP, validating VFL as viable for European ECG data. However, DP substantially degrades VFL performance: at $\varepsilon=1$, PTB-XL accuracy drops to 50.7–71.3% depending on party count, as the split learning architecture amplifies noise through the concatenation layer. Counterintuitively, 3-party configurations are more DP-robust than 2-party on PTB-XL (71.3% vs. 61.4% at $\varepsilon=1$, 60% overlap), as distributing features across more parties reduces per-party gradient sensitivity. Cardiovascular achieves ~ 60 –62% without DP (below the HFL baseline of 73%), indicating that feature splitting reduces model capacity on small feature sets (11 features). CDC Diabetes exhibits $F1 \approx 0$ across all VFL configurations, confirming that the 86% majority-class imbalance causes collapse regardless of learning paradigm.

10-Seed Statistical Robustness (RQ#40, 240 experiments). Table S-LXXXI provides enhanced statistical validation with 10 seeds (42, 123, 456, 789, 999, 1111, 2222, 3333, 4444, 5555) for 4 algorithms across 3 datasets under IID and NonIID conditions. The 95% confidence intervals confirm several key findings: (1) FedAvg, FedProx, and Ditto are statistically indistinguishable on Cardiovascular IID (overlapping CIs: 72.7–73.3, 72.6–73.1, 72.7–73.2), reinforcing Finding #5. (2) HPFL is significantly lower than FedAvg on CDC Diabetes NonIID

TABLE S-LXXVII
CDC DIABETES: ALGORITHM COMPARISON (BLOCK A, 70 EXP.). $K=5$, 20 ROUNDS, 3 LOCAL EPOCHS, NO DP. MEAN $\pm$ STD OVER 5 SEEDS.

Algorithm	IID		NonIID ($\alpha=0.5$)	
	Acc (%)	F1	Acc (%)	F1
FedAvg	70.1 $\pm$ 0.5	.429	80.8 $\pm$ 6.6	.218
FedProx	69.5 $\pm$ 0.7	.425	80.6 $\pm$ 6.7	.230
Ditto	70.1 $\pm$ 0.6	.429	80.6 $\pm$ 6.8	.212
FedExP	70.1 $\pm$ 0.5	.429	80.8 $\pm$ 6.6	.218
FedLC	69.7 $\pm$ 0.5	.428	80.7 $\pm$ 6.8	.228
FedLESAM	70.1 $\pm$ 0.5	.429	80.8 $\pm$ 6.6	.218
HPFL	57.2 $\pm$ 0.7	.375	72.2 $\pm$ 8.3	.415

HPFL achieves higher F1 under NonIID (0.415 vs. 0.218 for FedAvg) despite lower accuracy, as its personalized heads partially mitigate class collapse at the cost of overall accuracy.

TABLE S-LXXVIII
CDC DIABETES: DP ABLATION (BLOCK B, 72 EXP.). $K=5$, IID AND NonIID. MEAN $\pm$ STD OVER 3 SEEDS. ACC (%) / F1.

	$\epsilon=1$	$\epsilon=5$	$\epsilon=10$	$\epsilon=50$
<i>IID</i>				
FedAvg	83.8 / .071	85.8 / .013	83.1 / .255	68.7 / .423
Ditto	57.9 / .353	62.8 / .211	81.4 / .310	71.6 / .433
HPFL	18.0 / .235	41.4 / .159	64.8 / .092	56.3 / .372
<i>NonIID</i> ($\alpha=0.5$)				
FedAvg	81.6 / .093	81.4 / .121	81.7 / .196	78.1 / .295
Ditto	38.4 / .166	40.1 / .208	67.3 / .245	76.3 / .287
HPFL	49.3 / .171	41.4 / .201	40.2 / .265	58.6 / .373

FedAvg’s high accuracy under strict DP ($\epsilon=1$: 83.8%) masks majority-class collapse (F1=0.071). The DP “regularization” effect is a statistical artifact on imbalanced data.

TABLE S-LXXIX
CDC DIABETES: CLIENT SCALABILITY (BLOCK C, 72 EXP.). NO DP, IID AND NonIID. MEAN $\pm$ STD OVER 3 SEEDS. ACC (%) / F1.

	$K=3$	$K=10$	$K=20$
<i>IID</i>			
FedAvg	70.2 $\pm$ 0.6 / .433	69.2 $\pm$ 0.4 / .418	69.3 $\pm$ 0.1 / .413
Ditto	70.3 $\pm$ 0.2 / .433	69.3 $\pm$ 0.2 / .418	69.4 $\pm$ 0.2 / .413
HPFL	63.5 $\pm$ 3.5 / .405	55.6 $\pm$ 0.5 / .364	57.1 $\pm$ 0.2 / .367
<i>NonIID</i> ($\alpha=0.5$)			
FedAvg	78.0 $\pm$ 5.9 / .300	81.8 $\pm$ 3.8 / .310	84.2$\pm$0.3 / .413
Ditto	76.8 $\pm$ 7.0 / .296	81.8 $\pm$ 4.0 / .311	84.0 $\pm$ 0.3 / .424
HPFL	74.8 $\pm$ 8.3 / .342	69.8 $\pm$ 8.9 / .409	68.6 $\pm$ 0.3 / .417

Baseline $K=5$ results in Table S-LXXVII. NonIID improvement with increasing K reflects dilution of per-client label skew.

(−6.8pp, non-overlapping CIs: 73.9–78.3 vs. 79.3–86.5), confirming its counterproductive behavior on imbalanced large-scale data. (3) All algorithms achieve $>92\%$ on PTB-XL regardless of IID/NonIID condition (CIs entirely above 91.8% for FedAvg/FedProx/Ditto), confirming PTB-XL’s robustness as an FL benchmark. (4) NonIID variance is substantially higher than IID on CDC Diabetes (std 3.0–5.2% vs. 0.8–1.6%), as Dirichlet partitioning creates heterogeneous client-level class distributions.

I. Computational Cost

Table S-LXXXII reports mean training time per algorithm on the three main datasets (20 rounds $\times$ 3 local epochs, 5

seeds, single Apple M-series CPU). The overhead column shows the ratio relative to FedAvg.

J. Thesis Robustness Validation

To stress-test the central thesis that personalization architecture—not aggregation algorithm—is the critical deployment decision, 516 additional experiments were conducted across three dimensions: hyperparameter sensitivity (75 experiments), data efficiency (225 experiments), and compound EHDS stress (216 experiments). All experiments use 5 seeds except compound stress (3 seeds), with results reported as mean $\pm$ standard deviation.

TABLE S-LXXX

VERTICAL FL EXTENDED RESULTS (216 EXPERIMENTS). SPLITNN WITH BINARY CLASSIFICATION. MEAN $\pm$ STD OVER 3 SEEDS (42, 123, 456). ACC (%) / F1. BEST NO-DP ACCURACY PER DATASET IN **BOLD**.

Dataset	Parties	No DP			$\epsilon=5$			$\epsilon=1$		
		ovl=0.6	ovl=0.8	ovl=1.0	ovl=0.6	ovl=0.8	ovl=1.0	ovl=0.6	ovl=0.8	ovl=1.0
PTB-XL (bin.)	2	99.3/ 994	99.3/994	85.1/904	96.3/967	95.4/957	70.8/813	61.4/713	57.2/647	50.7/668
	3	94.2/944	94.8/950	76.2/820	91.3/917	92.4/928	77.1/824	71.3/760	70.8/766	70.1/758
	4	—	—	—	—	—	—	—	—	—
Cardiovascular	2	61.8/595	62.0/590	60.6/544	48.2/166	49.5/203	51.8/457	50.0/173	50.9/198	55.1/378
	3	61.0/588	62.4/609	61.8/594	52.7/459	54.2/497	58.1/497	51.6/325	51.0/323	53.7/459
	4	59.8/556	60.6/569	62.2/619	54.0/428	53.7/484	50.6/362	54.7/424	55.0/432	47.9/240
CDC Diabetes	2	86.2/036	86.2/043	86.1/045	86.1/000	86.1/000	86.1/000	86.1/000	86.1/000	86.1/000
	3	86.0/029	86.1/049	86.1/088	86.1/000	86.1/000	86.1/000	86.1/000	86.1/000	86.1/000
	4	86.1/047	86.1/067	86.1/075	86.1/000	86.1/000	86.1/000	86.1/000	86.1/000	86.1/000

PTB-XL binarized: NORM (43%) vs. ABNORMAL (57%). CDC Diabetes F1 $\approx$ 0 indicates majority-class collapse across all VFL configurations. PTB-XL 4-party not tested (9 features insufficient for 4-way split). Std omitted for space; all <5% except PTB-XL under DP.

TABLE S-LXXXI

10-SEED STATISTICAL ROBUSTNESS (240 EXPERIMENTS). 4 ALGORITHMS $\times$ 3 DATASETS $\times$ IID/NONIID $\times$ 10 SEEDS. $K=5$ CLIENTS, 20 ROUNDS, 3 LOCAL EPOCHS, NO DP. BEST IID ACCURACY IN **BOLD**. 95% CI VIA t -DISTRIBUTION.

Algo.	PTB-XL (5-class)				Cardiovascular				CDC Diabetes (253K)			
	IID		NonIID		IID		NonIID		IID		NonIID	
	Mean $\pm$ Std	95% CI	Mean $\pm$ Std	95% CI	Mean $\pm$ Std	95% CI	Mean $\pm$ Std	95% CI	Mean $\pm$ Std	95% CI	Mean $\pm$ Std	95% CI
FedAvg	92.2 $\pm$ 0.5	[91.8, 92.5]	92.2 $\pm$ 0.5	[91.8, 92.5]	73.0$\pm$0.4	[72.7, 73.3]	69.5 $\pm$ 4.8	[66.0, 72.9]	71.6 $\pm$ 0.8	[71.0, 72.2]	82.9 $\pm$ 5.0	[79.3, 86.5]
FedProx	92.2 $\pm$ 0.5	[91.8, 92.5]	92.2 $\pm$ 0.5	[91.8, 92.5]	72.9 $\pm$ 0.3	[72.6, 73.1]	70.2 $\pm$ 3.4	[67.8, 72.7]	70.8 $\pm$ 0.8	[70.2, 71.4]	82.7 $\pm$ 5.1	[79.0, 86.3]
Ditto	92.4$\pm$0.5	[92.1, 92.8]	92.4$\pm$0.5	[92.1, 92.8]	72.9 $\pm$ 0.4	[72.7, 73.2]	69.2 $\pm$ 4.7	[65.9, 72.6]	71.7$\pm$1.0	[71.0, 72.4]	82.5 $\pm$ 5.2	[78.7, 86.2]
HPFL	87.7 $\pm$ 0.5	[87.4, 88.1]	87.7 $\pm$ 0.5	[87.4, 88.1]	68.6 $\pm$ 0.6	[68.2, 69.0]	66.7 $\pm$ 5.5	[62.8, 70.7]	71.6 $\pm$ 1.6	[70.4, 72.7]	76.1 $\pm$ 3.0	[73.9, 78.3]

HPFL's 95% CI on CDC NonIID (73.9–78.3) does not overlap with FedAvg (79.3–86.5), confirming significant underperformance ($p<0.01$, paired t -test). HPFL's lower PTB-XL accuracy (87.7% vs. 92.4%) reflects the 5-class evaluation where backbone sharing is critical.

TABLE S-LXXXII

TRAINING TIME (SECONDS) AND RELATIVE OVERHEAD VS. FEDAVG. 20 ROUNDS, 3 LOCAL EPOCHS, 5 CLIENTS, SINGLE CPU.

Algorithm	Runtime (s)			Overhead vs. FedAvg		
	PX	CV	BC	PX	CV	BC
FedAvg	4.9	16.3	0.1	1.0 $\times$	1.0 $\times$	1.0 $\times$
FedProx	6.4	23.4	0.2	1.3 $\times$	1.4 $\times$	2.0 $\times$
FedLC	5.6	14.3	0.1	1.1 $\times$	0.9 $\times$	1.0 $\times$
FedExP	5.3	16.5	0.1	1.1 $\times$	1.0 $\times$	1.0 $\times$
FedLESAM	5.7	19.1	0.2	1.2 $\times$	1.2 $\times$	2.0 $\times$
Ditto	15.7	38.2	0.4	3.2 $\times$	2.3 $\times$	4.0 $\times$
HPFL	6.1	13.4	0.2	1.2 $\times$	0.8 $\times$	2.0 $\times$

Ditto's higher cost (2–4 $\times$) reflects per-client personalized model training. HPFL's split-head adds negligible overhead ($\leq 1.2\times$). On compact MLP models, all algorithms complete within seconds to minutes, making computational cost a non-factor for EHDS deployment decisions.

1) *Phase A: Ditto Hyperparameter Sensitivity*: Ditto's L2 regularization strength λ was varied across a $100\times$ range ($\lambda \in \{0.01, 0.05, 0.1, 0.5, 1.0\}$) on three datasets (CV, HD, PTB-XL), 5 seeds each (75 experiments total). Table S-LXXXIII shows that accuracy is remarkably insensitive to λ : maximum range is only 1.44pp on PTB-XL and 0.23pp on Cardiovascular. This confirms that the personalization advantage reported in the main paper is not an artifact of hyperparameter tuning.

2) *Phase B: Personalization vs. Data Fraction*: To determine whether personalization requires large datasets, we evaluated FedAvg, Ditto, and HPFL at five data fractions ($\{10\%, 25\%, 50\%, 75\%, 100\%\}$) across three datasets and 5

seeds (225 experiments total). Table S-LXXXIV reveals that personalization advantage *emerges* at $\geq 25\%$ data fraction: on Cardiovascular, Ditto reaches +11.2pp over FedAvg at 25% data and stabilizes at +12.0pp with full data. At 10% data, personalization provides no benefit (insufficient local data for fine-tuning). On PTB-XL, FedAvg leads at 10% fraction due to sufficient global data, while personalized methods catch up at 25%.

This finding has direct EHDS implications: phased hospital rollout (2025–2031) reaches near-optimal analytics at just 25% participation, and personalized algorithms are the optimal choice as soon as minimum data thresholds are met.

TABLE S-LXXXIII

DITTO λ SENSITIVITY: ACCURACY (%) ACROSS $100 \times \lambda$ RANGE. MEAN $\pm$ STD OVER 5 SEEDS. MAX RANGE PER DATASET SHOWN IN LAST COLUMN.

λ	0.01	0.05	0.1	0.5	1.0	Range
CV	82.5 $\pm$ 4.8	82.5 $\pm$ 4.7	82.5 $\pm$ 4.8	82.4 $\pm$ 4.8	82.3 $\pm$ 4.8	0.23pp
HD	74.0 $\pm$ 2.9	74.2 $\pm$ 2.8	74.3 $\pm$ 2.8	74.8 $\pm$ 3.4	75.3 $\pm$ 3.3	1.32pp
PX	91.1 $\pm$ 0.5	91.5 $\pm$ 0.4	92.0 $\pm$ 0.4	92.5 $\pm$ 0.3	92.5 $\pm$ 0.3	1.44pp

TABLE S-LXXXIV

PERSONALIZATION VS. DATA FRACTION: ACCURACY (%) AT VARYING DATASET SIZES. MEAN $\pm$ STD OVER 5 SEEDS. PER-DATASET DITTO–FEDAVG GAP (Δ) SHOWN. BOLD = BEST PER CONDITION.

Frac.	Cardiovascular (CV)				Heart Disease (HD)				PTB-XL (PX)			
	FedAvg	Ditto	HPFL	Δ	FedAvg	Ditto	HPFL	Δ	FedAvg	Ditto	HPFL	Δ
10%	67.1 $\pm$ 1.7	78.6$\pm$6.4	77.3 $\pm$ 6.8	+11.5	55.2 $\pm$ 3.8	61.8$\pm$9.4	59.3 $\pm$ 7.9	+6.6	80.5$\pm$4.8	59.8 $\pm$ 11.7	61.5 $\pm$ 11.2	−20.7
25%	70.4 $\pm$ 2.8	81.6$\pm$5.1	81.6 $\pm$ 4.9	+11.2	55.7 $\pm$ 3.9	67.0$\pm$4.4	64.0 $\pm$ 9.5	+11.3	90.1$\pm$0.6	86.9 $\pm$ 1.1	89.8 $\pm$ 1.0	−3.2
50%	70.8 $\pm$ 2.1	82.2$\pm$4.9	82.2 $\pm$ 4.6	+11.4	55.7 $\pm$ 3.9	72.1 $\pm$ 3.2	72.6$\pm$2.8	+16.4	91.8 $\pm$ 0.4	90.9 $\pm$ 0.8	92.1$\pm$0.4	−0.9
75%	71.0 $\pm$ 2.1	82.5$\pm$4.8	82.4 $\pm$ 4.7	+11.5	58.8 $\pm$ 6.7	71.2$\pm$2.8	70.9 $\pm$ 1.7	+12.4	91.8 $\pm$ 0.2	91.6 $\pm$ 0.3	92.4$\pm$0.2	−0.2
100%	70.5 $\pm$ 2.7	82.5$\pm$4.8	82.5 $\pm$ 4.8	+12.0	59.9 $\pm$ 6.9	74.3$\pm$2.8	73.3 $\pm$ 3.0	+14.4	92.0 $\pm$ 0.3	92.0 $\pm$ 0.4	92.5$\pm$0.3	0.0

On CV and HD, personalized methods dominate at all fractions. On PTB-XL, FedAvg leads at 10% (global data suffices) but the gap closes rapidly; at $\geq 50\%$, all three algorithms converge. Δ = Ditto − FedAvg. Negative Δ on PTB-XL reflects PTB-XL’s near-IID 52-site distribution where global aggregation is already optimal.

3) *Phase C: Compound EHDS Stress Frontier*: To evaluate robustness under realistic EHDS deployment conditions, 216 experiments on Cardiovascular tested all combinations of: data distribution (IID, non-IID $\alpha=0.5$), differential privacy ($\epsilon \in \{\infty, 10, 1\}$), client participation (100%, 60%), and Article 71 opt-out (0%, 20%), across FedAvg, Ditto, and HPFL with 3 seeds each (24 conditions $\times$ 3 algorithms $\times$ 3 seeds = 216 experiments).

Table S-LXXXV presents key conditions. The main findings are:

- 1) **Ditto dominance**: Ditto outperforms FedAvg in 81% of conditions (58/72 comparisons), with a mean advantage of +9.6pp $\pm$ 8.7pp. The maximum advantage reaches +29.3pp under specific stress conditions.
- 2) **DP is the dominant stressor**: Under non-IID conditions with no other stressors, DP at $\epsilon=1$ causes −21.0pp for FedAvg and −23.7pp for Ditto, while participation (100% $\rightarrow$ 60%) causes only −3.0pp for FedAvg and <0.1pp for Ditto, and opt-out (0% $\rightarrow$ 20%) causes <0.1pp for both. This hierarchy (DP $\gg$ participation $>$ opt-out) directly informs EHDS privacy budget allocation.
- 3) **HPFL excels under strict DP**: Under non-IID + $\epsilon=1$ (the most challenging privacy regime), HPFL retains 70.4% vs. Ditto 58.0% vs. FedAvg 49.0%—confirming that HPFL’s split-head architecture absorbs DP noise while preserving personalization capacity.
- 4) **Worst-case compound stress**: Under simultaneous non-IID + $\epsilon=1$ + 60% participation + 20% opt-out, Ditto retains 72.6% vs. FedAvg 52.4% (+20.2pp). Even HPFL (66.7%) substantially outperforms FedAvg. The personalization advantage is *amplified* under stress.

4) *Summary of Thesis Robustness Claims*: The 516 thesis robustness experiments support eight citable claims:

- 1) Ditto’s personalization advantage is hyperparameter-insensitive: ≤ 1.44 pp accuracy variation across $100 \times \lambda$ range on three datasets (75 experiments).
- 2) Personalization advantage emerges at $\geq 25\%$ data availability and stabilizes, with Ditto reaching +11.2pp over FedAvg at 25% on Cardiovascular (225 experiments).
- 3) At 10% data fraction, personalization provides no benefit—insufficient local data for fine-tuning nullifies the personalization mechanism.
- 4) On PTB-XL (near-IID, 52 European sites), global aggregation matches personalized methods at all fractions $\geq 25\%$, confirming the modality-dependent thesis from the main paper.
- 5) Under compound EHDS stress, Ditto outperforms FedAvg in 81% of conditions (+9.6pp mean advantage, 216 experiments).
- 6) DP at $\epsilon=1$ is the dominant EHDS stressor (−21pp FedAvg, −23.7pp Ditto), dwarfing participation (<3pp) and opt-out (<1pp) effects.
- 7) HPFL’s split-head architecture is uniquely resilient to strict DP: retains 70.4% at $\epsilon=1$ vs. Ditto 58.0%, consistent with the main paper’s Non-IID $\times$ DP interaction finding.
- 8) Under worst-case compound stress (non-IID + $\epsilon=1$ + 60% participation + 20% opt-out), personalized methods retain substantial advantages: Ditto 72.6%, HPFL 66.7% vs. FedAvg 52.4%.

APPENDIX F MEDICAL IMAGING RESULTS

A. Clinical Imaging: Extended Details

1) *Datasets*: Three clinical imaging datasets cover representative EHDS scenarios:

TABLE S-LXXXV

COMPOUND EHDS STRESS FRONTIER (CARDIOVASCULAR, NON-IID $\alpha=0.5$): ACCURACY (%) UNDER COMBINED STRESSORS. MEAN $\pm$ STD OVER 3 SEEDS. BOLD = BEST PER CONDITION. Δ_D = DITTO – FEDAVG GAP.

Distribution	DP	Participation	Opt-out	FedAvg	Ditto	HPFL	Δ_D
Non-IID	No DP	100%	0%	70.0 $\pm$ 3.0	81.7$\pm$1.1	81.7 $\pm$ 1.3	+11.7
Non-IID	$\varepsilon=10$	100%	0%	68.0 $\pm$ 3.0	80.9$\pm$1.1	80.7 $\pm$ 1.2	+12.9
Non-IID	$\varepsilon=1$	100%	0%	49.0 $\pm$ 7.0	58.0 $\pm$ 5.2	70.4$\pm$3.4	+9.0
Non-IID	No DP	60%	0%	67.0 $\pm$ 6.1	81.6$\pm$1.2	81.6 $\pm$ 1.3	+14.6
Non-IID	No DP	100%	20%	70.0 $\pm$ 3.2	81.6 $\pm$ 1.1	81.8$\pm$1.1	+11.6
Non-IID	$\varepsilon=10$	60%	20%	63.1 $\pm$ 4.0	80.4$\pm$1.4	79.9 $\pm$ 1.4	+17.3
Non-IID	$\varepsilon=10$	100%	20%	56.4 $\pm$ 2.1	80.8$\pm$1.4	78.8 $\pm$ 2.9	+24.4
Non-IID	$\varepsilon=1$	60%	20%	52.4 $\pm$ 3.8	72.6$\pm$3.8	66.7 $\pm$ 5.5	+20.2

Complete results for all 24 conditions (including IID) available in the checkpoint file. Key finding: DP at $\varepsilon=1$ is the dominant stressor (-21 pp FedAvg), while participation (100% $\rightarrow$ 60%) and opt-out (0% $\rightarrow$ 20%) are negligible (<3 pp). Personalized methods maintain substantial advantages under all stress combinations.

- **Chest X-ray** [50]: 5,856 pediatric radiographs (NORMAL/PNEUMONIA, 2.7:1 imbalance)
- **Brain Tumor MRI**: 7,023 T1-weighted CE MRI slices (4-class: glioma, healthy, meningioma, pituitary)
- **Skin Cancer**: 3,297 dermoscopy images (binary benign/malignant)

2) *Model Architectures*: **HealthcareResNet**: ResNet-18 [49] pretrained on ImageNet, GroupNorm replacing BatchNorm for FL stability. FedBN [48] skips normalization during aggregation. Partial backbone freeze (level 1). ~ 11.2 M parameters.

HealthcareCNN: 5-block CNN with GroupNorm, progressive channels (32 $\rightarrow$ 512), graduated Dropout (0.15 $\rightarrow$ 0.3). Classifier: Flatten $\rightarrow$ FC(512) $\rightarrow$ FC(128) $\rightarrow$ FC(K). ~ 12 M parameters.

Data augmentation: random horizontal flip, rotation ($\pm 15^\circ$), brightness jitter ($\pm 10\%$). ImageNet normalization.

3) *Experimental Configuration*: All three imaging datasets use 4 algorithms (FedAvg, FedLESAM, Ditto, HPFL) $\times$ 3 seeds (42, 123, 456) = 36 experiments total, enabling full statistical comparison across algorithms and datasets.

- 5 clients, Dirichlet $\alpha=0.5$, Adam optimizer
- ResNet-18, FedBN, class-weighted loss, mixed precision
- freeze_level=2 (conv1+bn1+layer1 frozen), local_epochs=2
- Chest X-ray: 20 rounds, early stopping (patience=6), lr=0.001
- Brain Tumor: 20 rounds, early stopping (patience=4, min_rounds=8), lr=0.0005
- Skin Cancer: 20 rounds, early stopping (patience=4, min_rounds=8), lr=0.001

4) *Reproducibility*: All experiments are fully reproducible:

```
cd fl-ehds-framework
# Chest X-ray (4 algos x 3 seeds, ~4h)
python -m benchmarks.run_imaging_extended
# FedLESAM+HPFL (18 exp, ~12h)
python -m benchmarks.run_imaging_delta
```

```
# FedAvg+Ditto completion (8 exp, ~17h)
python -m benchmarks.run_imaging_completion
# Confusion matrices
python -m benchmarks.run_confusion_matrix_chest
python -m benchmarks.run_confusion_matrix_bc
```

Results, checkpoints, and logs are auto-saved to `benchmarks/paper_results_delta/`.

Repository: <https://github.com/FabioLiberti/FL-EHDS-FLICS2026>

5) *Chest X-ray Results*: Table S-LXXXVI reports Chest X-ray results for four FL algorithms (ResNet-18, 5 clients, Dirichlet $\alpha=0.5$, 20 rounds max with early stopping).

Key findings. *First*, federated learning on Chest X-ray achieves **87.3% accuracy** (FedAvg) for binary pneumonia detection with ResNet-18, confirming FL viability for medical imaging under EHDS-relevant non-IID conditions. *Second*, **FedLESAM is identical to FedAvg on imaging**: despite SAM’s theoretical sharpness-aware optimization, FedLESAM produces per-seed accuracies indistinguishable from FedAvg (87.8% vs. 87.3%), a pattern confirmed on all three imaging datasets (see Tables S-LXXXVII). *Third*, **on Chest X-ray, personalization is counterproductive**: Ditto achieves 80.0% (-7.3 pp vs. FedAvg), with substantially higher variance ($\sigma=10.6$ pp vs. 3.6pp). HPFL *fails catastrophically* (69.1%, -18.2 pp), early-stopping at round 8 on all seeds with peak accuracy at rounds 1–2, indicating that the split-head architecture cannot learn shared representations with Dirichlet-partitioned data. However, as shown in Table S-LXXXVII, this Chest X-ray finding does *not* generalize: Ditto dominates on Brain Tumor (+23.5pp) and Skin Cancer (+25.5pp). *Fourth*, Ditto exhibits a **fairness failure mode** on seed 456: per-client accuracies of [99.7%, 100.0%, 48.5%, 39.5%, 100.0%] (Jain 0.888) indicate that personalized local models overfit to majority-class clients while failing on minority-class clients.

TABLE S-LXXXVI

CHEST X-RAY PNEUMONIA DETECTION: FEDERATED ACCURACY (%) WITH RESNET-18, 5 CLIENTS, DIRICHLET $\alpha=0.5$, 20 ROUNDS, ADAM LR=0.001, EARLY STOPPING (PATIENCE=6). PER-SEED RESULTS WITH JAIN FAIRNESS INDEX.

Algo	s42	s123	s456	Mean	Jain
FedAvg	91.6	87.4	82.8	87.3	0.984
FedLESAM	91.6	87.8	84.0	87.8 [†]	0.984
Ditto	84.8	65.1	90.1	80.0	0.958
HPFL	62.1	73.4	71.8	69.1	0.762

Jain index averaged over 3 seeds. [†]FedLESAM produces results identical to FedAvg (SAM perturbation ineffective on this architecture). HPFL early-stops at round 8 on all seeds with best accuracy at rounds 1–2.

B. Brain Tumor and Skin Cancer Results

Table S-LXXXVII presents the complete imaging evaluation for Brain Tumor MRI (4-class) and Skin Cancer (binary), using the same 4-algorithm $\times$ 3-seed design as Chest X-ray (Table S-LXXXVI). This 24-experiment extension reveals a **method-dependent personalization effect** that fundamentally revises the Chest X-ray-only conclusion.

Cross-dataset imaging synthesis. The complete 36-experiment imaging evaluation (Tables S-LXXXVI and S-LXXXVII) reveals three findings:

First, personalization effectiveness is method-dependent, not modality-dependent. Ditto’s L2-regularized personal models dominate on Brain Tumor (77.3% vs. FedAvg 53.8%, +23.5pp) and Skin Cancer (90.5% vs. 65.0%, +25.5pp), while HPFL’s split-head architecture fails on all three datasets (−3.8pp to −18.2pp). The critical difference is architectural: Ditto maintains a *complete* personal model copy with L2 regularization toward the global model, preserving full CNN representation capacity; HPFL decouples the classifier head, which cannot learn robust shared representations under non-IID partitioning.

Second, Chest X-ray is the outlier, not the rule. On Chest X-ray, Ditto underperforms FedAvg (−7.3pp) and exhibits high variance ($\sigma=10.6$ pp). We hypothesize this reflects Chest X-ray’s specific characteristics: binary task with 2.7:1 class imbalance and the largest sample count (5,856) among imaging datasets, where global aggregation provides sufficient signal and local personalization introduces noise. On the smaller or multi-class datasets, personalization provides essential adaptation.

Third, FedLESAM $\equiv$ FedAvg across all imaging datasets. SAM’s sharpness-aware perturbation has no measurable effect on ResNet-18 under federated aggregation—per-seed accuracies are identical to four decimal places on Skin Cancer and Brain Tumor (seeds 123, 456). This confirms that the ResNet-18 + GroupNorm + FedBN loss landscape is sufficiently flat that SAM perturbation adds no benefit, consistent with the Chest X-ray finding.

Implications for EHDS: These results strengthen the practical recommendation that EHDS data permits should specify not only the FL algorithm but also the personalization mechanism. For imaging analytics, Ditto (L2-regularized personal models) should be the default recommendation, while HPFL

should be reserved for tabular analytics where its split-head architecture excels.

C. Statistical Significance and Effect Sizes

Table S-LXXXVIII reports paired *t*-test results and Cohen’s *d* effect sizes for Ditto vs. FedAvg across all three imaging datasets (3 seeds each). Despite the limited sample size ($n=3$), the large effect sizes provide strong evidence.

The Cohen’s *d* values for Ditto on Brain Tumor ($d=4.05$) and Skin Cancer ($d=3.62$) far exceed the conventional “large effect” threshold ($d>0.8$), indicating that the observed accuracy improvements are not artifacts of seed selection. HPFL shows no significant difference from FedAvg on any imaging dataset (Brain Tumor: $d=0.10$, $p=0.88$; Skin Cancer: $d=-0.45$, $p=0.73$; Chest X-ray: $d=-3.66$, $p=0.08$).

D. 10-Seed Wilcoxon Signed-Rank Tests on Imaging

To strengthen the 3-seed *t*-test results (Table S-LXXXVIII), we extend the imaging evaluation to 10 seeds (42, 123, 456, 789, 999, 7, 13, 31, 67, 101), yielding 90 additional experiments (3 algorithms $\times$ 3 datasets $\times$ 10 seeds). This enables Wilcoxon signed-rank tests—the non-parametric analog used for the tabular validation (Table S-XVI)—with sufficient power ($n=10$, minimum achievable $p=0.00195$).

Key findings. The 10-seed Wilcoxon analysis confirms and extends the 3-seed *t*-test results with higher statistical power:

First, Ditto significantly outperforms FedAvg on Brain Tumor ($p=0.004$, $r=0.96$, $\Delta=+9.1$ pp), confirming the large effect observed in the 3-seed analysis ($d=5.86$). The pooled test across all datasets is also highly significant ($p=0.004$, $r=0.60$, $\Delta=+5.6$ pp), establishing that Ditto’s advantage is not dataset-specific.

Second, HPFL is significantly worse than FedAvg on imaging—the inverse of the tabular finding. On Brain Tumor, HPFL loses 11.5pp ($p=0.027$); the pooled test yields $p=0.002$ ($r=0.68$, $\Delta=-12.7$ pp). This represents a **modality-dependent reversal**: the same algorithm that significantly outperforms FedAvg on all three tabular datasets (Table S-XVI) significantly *underperforms* on imaging. HPFL’s split-head architecture cannot learn robust shared CNN representations under non-IID partitioning; its detached classifier head fails to benefit from the rich feature hierarchies that ResNet-18 produces.

TABLE S-LXXXVII

BRAIN TUMOR MRI (4-CLASS) AND SKIN CANCER (BINARY): COMPLETE 4-ALGORITHM EVALUATION WITH RESNET-18. 5 CLIENTS, DIRICHLET $\alpha=0.5$, 20 ROUNDS MAX, EARLY STOPPING (PATIENCE=4, MIN_ROUNDS=8), FREEZE_LEVEL=2, LOCAL_EPOCHS=2. PER-SEED BEST ACCURACY (%) WITH MEAN $\pm$ STD AND JAIN FAIRNESS INDEX. BEST ACCURACY PER DATASET IN **BOLD**.

Dataset	Algorithm	s42	s123	s456	Mean $\pm$ std	Jain	Avg Rnds	Δ vs FedAvg
Brain Tumor	FedAvg [†]	59.9	48.9	52.6	53.8 $\pm$ 4.6	0.830	13.0	—
	FedLESAM [†]	59.9	48.9	52.6	53.8 $\pm$ 4.6	0.830	13.0	+0.0
	Ditto	71.9	83.8	76.3	77.3$\pm$4.9	0.894	16.0	+23.5
	HPFL	45.0	43.3	61.8	50.0 $\pm$ 8.5	—	—	−3.8
Skin Cancer	FedAvg [†]	69.0	72.1	53.8	65.0 $\pm$ 8.0	—	—	—
	FedLESAM [†]	69.0	72.1	53.8	65.0 $\pm$ 8.0	—	—	+0.0
	Ditto	88.6	91.2	91.6	90.5$\pm$1.3	0.863	8.0	+25.5
	HPFL	56.8	55.4	70.4	60.9 $\pm$ 6.9	—	—	−4.1

[†]FedLESAM produces results identical to FedAvg on all three imaging datasets (SAM perturbation ineffective on ResNet-18 under federated aggregation), confirming the pattern observed on Chest X-ray. Brain Tumor: $\text{lr}=0.0005$; Skin Cancer: $\text{lr}=0.001$. Ditto on Skin Cancer achieves peak accuracy at rounds 1–2 and early-stops at round 8, indicating rapid personalization convergence on binary tasks.

TABLE S-LXXXVIII

STATISTICAL SIGNIFICANCE OF DITTO VS. FEDAVG ON IMAGING DATASETS. PAIRED t -TEST ($n=3$ SEEDS). COHEN’S d CONVENTIONS: $|d|>0.8$ = LARGE EFFECT.

Dataset	Δ (pp)	Cohen’s d	t -stat	p -value
Brain Tumor	+23.5	4.05	3.56	0.071
Skin Cancer	+25.5	3.62	4.13	0.054
Chest X-ray	−0.9	−0.22	−0.25	0.828

Brain Tumor Ditto achieves $d=4.05$ (very large effect); with only 3 seeds, the paired t -test approaches but does not reach $\alpha=0.05$ ($p=0.071$); the 10-seed Wilcoxon test (Table S-LXXXIX) confirms significance ($p=0.004$). Skin Cancer Ditto ($d=3.62$) narrowly misses significance ($p=0.054$) due to high FedAvg variance ($\sigma=8.0\text{pp}$ vs. Ditto $\sigma=1.3\text{pp}$). On Chest X-ray, Ditto is statistically indistinguishable from FedAvg.

TABLE S-LXXXIX

IMAGING WILCOXON SIGNED-RANK p -VALUES VS. FEDAVG (10 SEEDS, PAIRED BY SEED). $\dagger$: $p < 0.05$; $\ddagger$: $p < 0.01$. POOLED: TEST ACROSS 30 PAIRED OBSERVATIONS (3 DATASETS $\times$ 10 SEEDS). r : RANK-BISERIAL CORRELATION ($|r| > 0.5$ = LARGE EFFECT). $\bar{\Delta}$: MEAN ACCURACY DIFFERENCE (PP).

vs FedAvg	Dataset	p	r	$\bar{\Delta}$ (pp)
Ditto	Brain Tumor	0.004[‡]	0.96	+9.1
	Skin Cancer	0.084	0.64	+6.2
	Chest X-ray	0.557	0.24	+1.5
	Pooled (30 obs.)	0.004[‡]	0.60	+5.6
HPFL	Brain Tumor	0.027[†]	0.78	−11.5
	Skin Cancer	0.066	0.75	−8.7
	Chest X-ray	0.160	0.53	−17.8
	Pooled (30 obs.)	0.002[‡]	0.68	−12.7

Third, **Chest X-ray remains the exception**: neither Ditto ($p=0.557$) nor HPFL ($p=0.160$) differs significantly from FedAvg, consistent with the hypothesis that large, binary datasets with class imbalance provide sufficient signal for global aggregation alone.

EHDS implication: The 10-seed validation produces an actionable recommendation for EHDS imaging analytics: **Ditto should be the default algorithm** (significant improvement on complex tasks, never worse than FedAvg), while **HPFL should be excluded** from imaging deployments despite its superiority on tabular data.

E. Extended Algorithm Comparison on Imaging (5 Seeds)

To verify whether the personalization gap generalises beyond the four algorithms in the main imaging comparison, we evaluate five additional FL algorithms—FedSAM [38], FedLC [39], FedDecorr [40], FedExP [42], and FedAvg as baseline—on all three imaging datasets with 5 seeds each (75 additional experiments). These algorithms employ server-side or regularisation-based strategies without maintaining separate personal models. A lightweight training configuration was used (12 rounds max, frozen backbone, local_epochs=1, early stopping patience=4, min_rounds=6) to isolate the effect of aggregation strategy from training depth.

Three observations emerge from this extended comparison.

TABLE S-XC

EXTENDED IMAGING ALGORITHM COMPARISON: 5 NON-PERSONALISATION FL ALGORITHMS WITH RESNET-18 (FROZEN BACKBONE). 5 CLIENTS, DIRICHLET $\alpha=0.5$, 12 ROUNDS MAX, EARLY STOPPING (PATIENCE=4, MIN_ROUNDS=6), LOCAL_EPOCHS=1. BEST ACCURACY (%) OVER 5 SEEDS WITH MEAN $\pm$ STD, JAIN FAIRNESS INDEX, AND AVERAGE ROUNDS TO EARLY STOPPING. BEST MEAN PER DATASET IN **BOLD**. FOR COMPARISON, DITTO (TABLE S-LXXXVII) ACHIEVES 77.3 $\pm$ 4.9% ON BRAIN TUMOR AND 90.5 $\pm$ 1.3% ON SKIN CANCER—FAR ABOVE ALL ALGORITHMS IN THIS TABLE.

Dataset	Algorithm	s42	s123	s456	s789	s999	Mean $\pm$ std	Jain	Avg Rnd
Chest X-ray	FedSAM	90.5	88.5	71.8	74.3	87.1	82.4$\pm$8.7	0.956	9.0
	FedAvg	84.7	89.7	71.8	74.3	72.6	78.6 $\pm$ 8.1	0.947	7.8
	FedExP	85.3	87.3	71.8	74.3	72.6	78.3 $\pm$ 7.4	0.918	7.2
	FedLC	73.9	85.3	71.8	74.3	75.2	76.1 $\pm$ 5.3	0.880	7.2
	FedDecorr	73.9	87.5	71.8	74.3	72.6	76.0 $\pm$ 6.5	0.903	7.2
Brain Tumor	FedAvg	53.7	53.6	61.2	49.3	56.5	54.9$\pm$4.4	0.892	12.0
	FedSAM	52.8	50.3	53.9	50.8	50.0	51.6 $\pm$ 1.7	0.876	12.0
	FedLC	50.4	44.1	47.3	48.2	54.2	48.8 $\pm$ 3.8	0.843	10.8
	FedDecorr	47.3	50.0	48.4	48.6	46.9	48.2 $\pm$ 1.3	0.830	10.8
	FedExP	46.9	50.1	45.7	31.7	56.5	46.2 $\pm$ 9.1	0.811	10.0
Skin Cancer	FedLC	55.6	62.4	53.8	76.6	77.0	65.1$\pm$11.2	0.878	9.2
	FedDecorr	63.4	76.9	53.8	75.5	54.5	64.8 $\pm$ 11.1	0.805	7.6
	FedExP	45.4	65.8	53.8	75.7	54.5	59.0 $\pm$ 11.8	0.803	8.2
	FedAvg	44.4	55.4	53.8	75.2	54.5	56.7 $\pm$ 11.3	0.801	7.2
	FedSAM	44.4	74.9	53.8	53.6	54.5	56.2 $\pm$ 11.2	0.711	6.4

All algorithms use frozen ResNet-18 backbone, Adam optimiser (lr=0.0005; Brain Tumor lr=0.0003), class-weighted loss, mixed precision. Kruskal–Wallis tests: Chest X-ray $H=1.89$, $p=0.76$; Brain Tumor $H=8.93$, $p=0.063$; Skin Cancer $H=4.85$, $p=0.30$. Only Brain Tumor approaches significance (Friedman $\chi^2=11.35$, $p=0.023$). Within-algorithm std (5–11pp) exceeds between-algorithm spread (6–9pp) on Chest X-ray and Skin Cancer, confirming that non-personalisation algorithms are **statistically indistinguishable** on imaging. For reference, Ditto (Table S-LXXXVII) achieves 77.3%, 90.5%, and 80.0% on Brain Tumor, Skin Cancer, and Chest X-ray respectively—20–35pp above any algorithm in this table on the first two datasets.

First, all five non-personalisation algorithms cluster within a narrow accuracy band (6–9pp spread per dataset), with no single algorithm consistently dominating: FedSAM leads on Chest X-ray, FedAvg on Brain Tumor, and FedLC on Skin Cancer. The Kruskal–Wallis tests confirm no statistically significant differences on Chest X-ray ($p=0.76$) or Skin Cancer ($p=0.30$). *Second*, this clustering mirrors the pattern observed on tabular models (main paper, Table VII), where server-side strategies also converge to FedAvg-equivalent solutions. On imaging, even SAM-based gradient perturbation (FedSAM) and feature decorrelation (FedDecorr) fail to break the baseline. *Third*, the 3-to-5 seed expansion confirms ranking stability on Chest X-ray (Kendall $\tau=0.80$) and Brain Tumor ($\tau=0.60$) but reveals instability on Skin Cancer ($\tau=0.20$), where within-algorithm variance ($\sigma\approx 11$ pp) exceeds the between-algorithm spread—underscoring the need for ≥ 5 seeds in imaging FL evaluations.

These results reinforce the central finding: on medical imaging, the **only intervention that produces a statistically significant accuracy improvement is architectural personalisation** (Ditto’s complete personal models), not the choice among non-personalisation aggregation strategies. This simplifies EHDS deployment decisions for imaging workloads to a single binary choice: personalised vs. global architecture.

F. Imaging Convergence Analysis

Figure S-47 shows training convergence curves (mean $\pm$ std over 3 seeds) for all four algorithms across the three imaging datasets.

Three convergence patterns emerge: (1) On **Chest X-ray**, FedAvg converges steadily to $\sim 88\%$ while Ditto plateaus at $\sim 87\%$ with higher variance—personalization adds noise rather than signal. (2) On **Brain Tumor**, Ditto separates from FedAvg after round 5 and continues improving until round 14–16, indicating that L2-regularized local models progressively adapt to client-specific tumour distributions. (3) On **Skin Cancer**, Ditto achieves peak accuracy at rounds 1–2 and early-stops at round 8, suggesting that even minimal local adaptation suffices for binary dermoscopy classification. HPFL collapses early on all datasets (best accuracy at rounds 1–4), confirming that the split-head architecture fundamentally fails on CNN-based imaging tasks.

G. Per-Client Fairness on Imaging

Figure S-48 presents per-client accuracy heatmaps for FedAvg vs. Ditto across all imaging datasets and seeds.

The per-client analysis reveals a nuanced fairness picture. On **Brain Tumor**, Ditto achieves higher Jain index (0.894 vs. 0.830), indicating that the accuracy improvement is distributed across clients rather than concentrated in a single client. On **Skin Cancer**, Ditto exhibits a fairness anomaly on seed 456: client 0 achieves only 9% accuracy while clients 2 and 4 achieve 100%—yielding Jain 0.704. This suggests that Ditto’s local models can overfit to majority-class distributions when a client has extreme class imbalance under Dirichlet partitioning. For EHDS deployment, per-client monitoring and minimum performance thresholds should be enforced.




Fig. S-47. Training convergence on imaging datasets (4 algorithms, mean $\pm$ std over 3 seeds). Left: Chest X-ray (FedAvg dominates throughout). Centre: Brain Tumor (Ditto separates after round 5, reaching 77%). Right: Skin Cancer (Ditto peaks at round 1–2 and maintains 90%+). HPFL collapses early on all datasets. FedLESAM is indistinguishable from FedAvg.

Per-Client Accuracy Heatmap: FedAvg vs Ditto (Imaging Datasets)




Fig. S-48. Per-client accuracy (%) for FedAvg (top) and Ditto (bottom) across Chest X-ray, Brain Tumor, and Skin Cancer. Each cell shows accuracy for one client on one seed. Ditto improves worst-client performance on Brain Tumor (min 44% $\rightarrow$ 22% is seed-dependent) but introduces fairness variance on Skin Cancer seed 456 (client 0: 9%).

H. Diagnostic Equity Analysis

Aggregate accuracy—the standard FL evaluation metric—can mask severe diagnostic disparities. A model achieving 57.3% accuracy on Breast Cancer might appear only moderately underperforming, but confusion matrix analysis (Table S-XXIV) reveals that FedAvg correctly identifies Malignant tumors only 21.5% of the time—clinically unacceptable for cancer screening. Similarly, FedAvg achieves 50.4% on Brain Tumor but completely misses Pituitary tumors on individual seeds (0% recall, Table S-XXVI). To quantify this diagnostic imbalance, we introduce the Diagnostic Equity Index.

1) *Definition:* The **Diagnostic Equity Index (DEI)** is defined as:

$$\text{DEI} = \min_c (R_c) \cdot (1 - \text{CV}(\mathbf{R})) \quad (1)$$

where $R_c \in [0, 1]$ is the recall (sensitivity) for class c , $\mathbf{R} = \{R_1, \dots, R_C\}$ is the recall vector across all C diagnostic classes, and $\text{CV}(\mathbf{R}) = \sigma(\mathbf{R})/\mu(\mathbf{R})$ is the coefficient of variation.

Properties:

- 1) $\text{DEI} \in [0, 1]$.
- 2) $\text{DEI} = 0$ if *any* class has 0% recall—the system fails to detect at least one disease category, which is clinically unacceptable regardless of aggregate performance.
- 3) $\text{DEI} = 1$ only if all classes have perfect and equal recall ($R_c = 1 \forall c$).
- 4) The CV term penalizes high inter-class variance even when the worst-case recall is non-zero, capturing “lop-sided” models that excel on one class at the expense of

others.

- 5) The $\min_c$ term ensures that DEI is dominated by the worst-performing class, aligning with clinical risk management where the missed diagnosis carries the highest cost.

Relationship to existing metrics: Unlike Jain’s fairness index (which measures inter-*client* performance equity), DEI measures inter-*class* diagnostic equity—a distinct and complementary concern. A model can have high Jain index (all clients perform similarly) but low DEI (all clients consistently miss the same disease class). For binary tasks, DEI reduces to $\min(R_0, R_1) \cdot (1 - |R_0 - R_1|/(R_0 + R_1))$; for multiclass tasks with $C \geq 3$, the CV term captures the full distribution shape.

2) *DEI Results:* Table S-XCI reports DEI values for all algorithm–dataset combinations where per-class confusion matrix data is available (Breast Cancer: 10 seeds, Table S-XXIV; Chest X-ray: 3 seeds, Table S-XXV; Brain Tumor: FedAvg and Ditto each 3 seeds, Table S-XXVI).

3) *Key Findings:* **Finding 1—DEI exposes hidden disparities:** On Breast Cancer, HPFL achieves $8\times$ higher DEI than FedAvg/FedProx/Ditto (0.740 vs. 0.092). All three non-personalized algorithms exhibit nearly identical DEI (0.092), confirming the algorithm collapse phenomenon (Table S-XXIV): FedAvg, FedProx, and Ditto converge to functionally identical predictions on the compact tabular MLP. Only HPFL’s per-client classifier heads escape single-class collapse, achieving 78.7% Malignant recall compared to 21.5%.

Finding 2—Accuracy and DEI can diverge: On Chest X-ray, HPFL achieves *higher* DEI (0.652) than FedAvg (0.528) **despite lower accuracy** (76.7% vs. 87.3%). FedAvg’s accuracy advantage comes from over-predicting the majority PNEUMONIA class (95.5% recall) while under-detecting NORMAL cases (65.1% recall). HPFL provides more balanced class detection (88.5% NORMAL, 72.4% PNEUMONIA). For a clinical deployment where both pneumonia detection *and* ruling out false positives matter, HPFL may be preferable despite its lower aggregate accuracy—a conclusion invisible to accuracy-only evaluation.

Finding 3—Ditto stabilizes multiclass imaging: On Brain Tumor (4 classes, 3 seeds), Ditto achieves $6.9\times$ higher DEI than FedAvg (0.435 vs. 0.063), with improvements across all under-detected classes: Pituitary +40.3pp (14.2%→54.5%), Glioma +20.1pp (36.6%→56.7%), and Meningioma +22.1pp (58.0%→80.1%), with only a minor Healthy trade-off (−2.1pp). The per-seed analysis confirms systematic improvement: Ditto DEI ranges 0.118–0.686 across seeds, while FedAvg DEI ranges 0.000–0.026—including a complete Pituitary recall failure (0%) at seed 123.

Finding 4—DEI reveals algorithm-specific failure modes: The low DEI values for FedAvg on Brain Tumor (0.063) and Breast Cancer (0.092) identify specific failure modes: class-level prediction collapse (Breast Cancer) and class-level instability across seeds (Brain Tumor, where Pituitary recall ranges 0–37%). These failures are invisible to aggregate accuracy and Jain fairness but have direct clinical consequences.

4) *EHDS Policy Implications:* The DEI results yield three specific recommendations for EHDS delegated acts and data permit design:

- 1) **Minimum DEI thresholds:** Data permits for diagnostic FL models should specify minimum DEI values alongside accuracy targets, particularly for cancer screening and multiclass diagnostic tasks. A threshold of $\text{DEI} \geq 0.5$ would exclude models exhibiting single-class collapse or severe class imbalance. For tasks with irreversible clinical consequences (e.g., missed cancer diagnosis), higher thresholds may be warranted.
- 2) **Algorithm–DEI conditioning:** The relationship between algorithm choice and DEI is task-specific. HPFL maximizes DEI on tabular class-imbalanced data (Breast Cancer: 0.740), while Ditto maximizes DEI on multiclass imaging (Brain Tumor: 0.435). The evidence-based recommendation matrix (Table S-VIII) incorporates DEI as a selection criterion alongside accuracy.
- 3) **Audit trail integration:** DEI should be computed and logged automatically during FL training as part of GDPR Article 30 audit trails, enabling post-hoc verification that deployed models meet per-class diagnostic standards. The FL-EHDS framework’s compliance logging module (Algorithm S1) can be extended to compute and record DEI after each evaluation round.

I. Byzantine Robustness on Medical Imaging

While the preceding section evaluated Byzantine resilience on tabular clinical data (Table S-LXII), medical imaging presents fundamentally different challenges: larger model capacity (ResNet-18 with 11.2M parameters vs. 3-layer MLP), richer gradient signals, and heterogeneous visual features across institutions. We evaluate three attack types—*label flipping* (adversary inverts class labels), *gradient noise injection* (Gaussian noise $\sigma=10$ added to model updates), and *sign flipping* (gradients multiplied by -5)—with Krum defense ($f=1$ Byzantine out of $K=5$ clients) across 3 imaging datasets, 3 algorithms, and 3 seeds (108 experiments total, ~ 27 hours on NVIDIA A40).

Key findings. The 108-experiment imaging Byzantine evaluation reveals three findings that complement the tabular analysis:

First, label flipping bypasses Krum defense on imaging. Unlike noise and sign flipping—which Krum filters effectively (average drop $< 3\text{pp}$)—label flipping causes significant accuracy degradation across all algorithms: FedAvg -7.1pp , Ditto -6.6pp , HPFL -16.1pp (averaged across datasets). This occurs because label flipping corrupts the *training data* rather than the gradient signal: Krum operates on model updates and cannot detect that a client’s update was trained on corrupted labels, since the update itself has normal magnitude and direction.

Second, Ditto is the most Byzantine-resilient algorithm. Ditto achieves the highest resilience scores on Chest X-ray (99.1%) and Skin Cancer (96.1%), and the highest absolute accuracy under attack on 2/3 datasets. Its L2-regularized

TABLE S-XCI

DIAGNOSTIC EQUITY INDEX (DEI) ACROSS ALGORITHMS AND DATASETS. HIGHER IS BETTER. $DEI = \min_c R_c \cdot (1 - CV(\mathbf{R}))$. BEST DEI PER DATASET IN **BOLD**.

Dataset	Algorithm	Acc (%)	$\min_c R_c$	CV	DEI
Breast Cancer	FedAvg	57.3	0.215	0.573	0.092
	FedProx	57.2	0.215	0.573	0.092
	Ditto	57.4	0.215	0.573	0.092
	HPFL	84.8	0.787	0.059	0.740
Chest X-ray	FedAvg	87.3	0.651	0.189	0.528
	HPFL	76.7	0.724	0.100	0.652
Brain Tumor	FedAvg	50.4	0.142	0.557	0.063
	Ditto	70.0	0.545	0.203	0.435

Breast Cancer: aggregated over 10 seeds (Table S-XXIV). Chest X-ray: aggregated over 3 seeds (Table S-XXV). Brain Tumor: FedAvg and Ditto both aggregated over 3 seeds (Table S-XXVI).

TABLE S-XCII

DEI STABILITY UNDER DIFFERENTIAL PRIVACY ON BREAST CANCER ($K=3$, $\alpha=0.5$, NON-IID). PER-CLASS RECALL AND DEI AVERAGED OVER 5 SEEDS. DP: GAUSSIAN MECHANISM, $C=1.0$, $\delta=10^{-5}$.

	No DP	$\epsilon=10$	$\epsilon=5$	$\epsilon=1$
<i>FedAvg</i>				
Accuracy (%)	56.0	61.5	59.8	59.9
Benign recall	0.626	0.846	0.814	0.812
Malignant recall	0.400	0.200	0.205	0.214
DEI	0.006	0.017	0.003	0.003
<i>HPFL</i>				
Accuracy (%)	88.1	91.0	90.8	90.1
Benign recall	0.855	0.958	0.960	0.960
Malignant recall	0.919	0.824	0.817	0.799
DEI	0.748	0.756	0.748	0.724

FedAvg exhibits single-class collapse (Malignant recall ≤ 0.40) regardless of DP level, yielding DEI ≈ 0 . HPFL maintains balanced per-class recall and stable DEI ≈ 0.72 – 0.76 across all DP budgets, confirming that privacy regulation does not exacerbate diagnostic disparities when personalized methods are used.

TABLE S-XCIII

BYZANTINE ROBUSTNESS ON MEDICAL IMAGING: BEST ACCURACY (%), MEAN $\pm$ STD OVER 3 SEEDS) UNDER KRUM DEFENSE ($f=1/5$ BYZANTINE). RESNET-18, 5 CLIENTS, DIRICHLET $\alpha=0.5$, 20 ROUNDS, EARLY STOPPING (PATIENCE=4). Δ : ACCURACY CHANGE FROM NO_ATTACK BASELINE. BEST PER-DATASET IN **BOLD**. RESILIENCE: RATIO OF MEAN ATTACKED ACCURACY TO NO_ATTACK BASELINE.

Dataset	Algorithm	Accuracy (%)				Δ vs. no_attack (pp)				Resil.
		no_attack	label_flip	noise	sign_flip	label_flip	noise	sign_flip	avg	
Chest X-ray	FedAvg	88.6 $\pm$ 2.2	84.8 $\pm$ 6.1	85.8 $\pm$ 2.6	86.4 $\pm$ 0.9	−3.8	−2.8	−2.2	−2.9	96.7%
	Ditto	85.4 $\pm$ 2.1	82.5 $\pm$ 3.9	83.3 $\pm$ 4.4	88.1$\pm$1.9	−2.9	−2.1	+2.7	−0.8	99.1%
	HPFL	67.8 $\pm$ 29.5	42.7 $\pm$ 23.5	66.1 $\pm$ 28.4	59.7 $\pm$ 25.6	−25.1	−1.7	−8.1	−11.6	82.9%
Brain Tumor	FedAvg	51.0 $\pm$ 17.7	44.1 $\pm$ 13.8	61.9 $\pm$ 2.1	61.8 $\pm$ 2.1	−6.9	+10.9	+10.8	+4.9	109.6%
	Ditto	57.7 $\pm$ 12.3	50.2 $\pm$ 13.1	66.3$\pm$7.2	70.1$\pm$5.0	−7.5	+8.6	+12.4	+4.5	107.8%
	HPFL	41.1 $\pm$ 12.1	28.1 $\pm$ 4.9	42.9 $\pm$ 9.6	40.6 $\pm$ 8.1	−13.0	+1.8	−0.5	−3.9	90.4%
Skin Cancer	FedAvg	71.5 $\pm$ 6.5	61.0 $\pm$ 12.5	69.5 $\pm$ 4.3	72.2 $\pm$ 3.0	−10.5	−2.0	+0.7	−3.9	94.5%
	Ditto	77.4$\pm$2.0	68.1$\pm$3.8	77.5$\pm$1.3	77.6$\pm$0.9	−9.3	+0.1	+0.2	−3.0	96.1%
	HPFL	64.0 $\pm$ 8.0	53.7 $\pm$ 7.5	62.6 $\pm$ 8.4	63.4 $\pm$ 7.4	−10.3	−1.4	−0.6	−4.1	93.7%

Label flipping is the most damaging attack across all algorithms and datasets (average -9.9 pp), while noise and sign flipping are largely neutralized by Krum (average -0.3 pp and $+1.5$ pp respectively). HPFL exhibits extreme variance on Chest X-ray ($\sigma=29.5\%$ at baseline, 23.5% under label_flip), confirming architectural instability. Brain Tumor no_attack baselines are lower than the non-Byzantine results (Table S-LXXXXVII) because Krum’s client filtering reduces effective participation even without adversaries—the filtering itself acts as a regularizer that *improves* accuracy under noise/sign_flip attacks (see text). Ditto achieves the highest absolute accuracy and resilience on 2/3 datasets, with remarkably low variance under attack ($\sigma \leq 5.0\%$ on Brain Tumor noise/sign_flip vs. FedAvg $\sigma=17.7\%$ at baseline).

personal model provides an implicit defense: even when the global model is degraded by Byzantine clients, the personal model’s L2 penalty toward the (partially corrupted) global acts as a conservative anchor, limiting the propagation of Byzantine influence to local models.

Third, noise and sign flipping anomalously improve Brain Tumor accuracy for FedAvg (+10.9pp, +10.8pp) and Ditto (+8.6pp, +12.4pp). This “noise-as-regularization” effect occurs because Krum’s filtering of the adversarial client also excludes that client’s potentially harmful non-IID update—on this small, 4-class dataset with high heterogeneity, removing one client’s contribution can actually improve the global model. This mirrors the finding from the tabular analysis (Table S-LXII) where Breast Cancer showed counterproductive defense effects.

EHDS implication: For cross-border imaging federations, Ditto with Krum defense provides the best resilience against Byzantine participants. However, **data poisoning via label corruption remains a critical threat** that aggregation-level defenses cannot address—EHDS implementations should complement Byzantine-resilient aggregation with data quality auditing mechanisms at the governance layer (Layer 1) to detect corrupted training labels before they enter the federation.

APPENDIX G

EHDS GOVERNANCE HYPOTHESIS TESTING

This section presents 720 dedicated governance validation experiments organized as three pre-registered hypotheses, each tested on two datasets of contrasting scale: Heart Disease (HD, 920 samples, 4 natural hospital sites, 13 features) and Cardiovascular (CV, 70,000 samples, 5 Dirichlet-partitioned clients, 11 features). The cross-dataset design distinguishes genuine governance effects from small-sample regularization artifacts. All experiments use $K=4$ (HD) or $K=5$ (CV) clients, 20 rounds, 3 local epochs, and atomic checkpointing with deterministic seeding.

A. Extended Governance Validation (E1–E4)

Prior to hypothesis testing, 140 extended governance experiments (E1–E4) established baselines for data minimization and opt-out impact:

- **E1: Data Minimization on HD** (40 exp.): Article 44 purpose-based feature selection across 4 EHDS purposes $\times$ 10 seeds. Scientific research retains 10/13 features; official statistics reduces to 2/13 (age, sex only).
- **E2: Data Minimization on CV** (40 exp.): Same protocol on CV. Mutual information filtering (threshold 0.01) retains 5/11 features for scientific research, 2/11 for official statistics. Six features (gender, height, glucose, smoke, alcohol, activity) have $MI < 0.01$ and are correctly removed.
- **E3: Governance Overhead** (10 exp.): Computational overhead of governance layer $< 3\%$ of total training time—negligible for deployment.
- **E4: Progressive Opt-Out** (50 exp.): Opt-out rates 0–30% on HD with 10 seeds. Cost $< 1\text{pp}$ at $\leq 10\%$ opt-out; -2.4pp at 30%, consistent with main paper findings.

B. Hypothesis H1: Compound EHDS Compliance

Hypothesis: Full EHDS compliance (data minimization + 10% opt-out + differential privacy) imposes an acceptable accuracy cost ($< 5\text{pp}$) for personalized FL algorithms at $\epsilon=10$.

Design: 3 algorithms (FedAvg, Ditto, HPFL) $\times$ 4 conditions (no governance, governance without DP, governance + $\epsilon=10$, governance + $\epsilon=1$) $\times$ 2 datasets. HD uses 5 seeds per condition; CV uses 10 seeds. Total: 240 experiments.

Key Findings (H1):

- 1) **Full compliance at $\epsilon=10$ is essentially free for personalized FL.** On CV (70K), Ditto loses only -0.7pp ($p=0.0007$) and HPFL -1.7pp ($p=0.004$) under simultaneous data minimization (11 $\rightarrow$ 4 features), 10% opt-out, and DP at $\epsilon=10$. FedAvg loses -3.8pp .
- 2) **Data minimization alone is negligible.** gov_noDP costs $< 0.5\text{pp}$ for all algorithms on CV despite removing 64% of features (11 $\rightarrow$ 4). The retained features (age, weight, systolic/diastolic BP) carry $> 95\%$ of the mutual information for cardiovascular prediction.
- 3) **$\epsilon=1$ is universally destructive.** All algorithms lose $> 10\text{pp}$ at $\epsilon=1$ on CV. FedAvg collapses to 53.4% (near random). This confirms the main paper finding: EHDS deployments should target $\epsilon \geq 10$.
- 4) **HD paradoxes are regularization artifacts.** On HD, Ditto *gains* $+5.2\text{pp}$ with $\epsilon=1$ and FedAvg gains $+9.6\text{pp}$ from minimization alone. These paradoxes disappear on CV, confirming they are small-sample phenomena (920 vs. 70K samples).

C. Hypothesis H2: Non-Uniform Article 71 Opt-Out Fairness

Hypothesis: Non-uniform opt-out patterns (where some hospitals have higher opt-out rates than others) do not significantly degrade cross-hospital fairness compared to uniform or no opt-out.

Design: 5 opt-out patterns $\times$ 10 seeds on each dataset using FedAvg as the worst-case algorithm. HD uses 4 hospitals; CV uses 5 clients. Total: 100 experiments.

Key Findings (H2):

- 1) **Non-uniform opt-out does not degrade fairness.** No pattern produces a statistically significant difference from the uniform_0 baseline on either dataset (all $p > 0.14$).
- 2) **On CV, all fairness metrics remain excellent.** Jain index ≥ 0.96 and accuracy range $< 29\text{pp}$ across all patterns. The CV dataset’s large sample size (14K per client) provides sufficient redundancy to absorb opt-out.
- 3) **Paradoxical fairness improvement on HD.** The half_30 and gradient patterns *improve* Jain fairness (0.944 vs. 0.922) and reduce the accuracy range (29.1pp vs. 47.5pp). This occurs because removing samples from the dominant Swiss hospital reduces its over-representation, acting as a natural re-balancing mechanism on the small, naturally heterogeneous HD dataset.
- 4) **EHDS Article 71 is operationally safe.** Even extreme opt-out (50% from a single hospital) does not signifi-

TABLE S-XCIV

H1: COMPOUND EHDS COMPLIANCE STRESS TEST. “ Δ ” SHOWS ACCURACY CHANGE FROM NO_GOVERNANCE BASELINE. FEATURES RETAINED AFTER ARTICLE 44 MINIMIZATION SHOWN AS FRACTION OF TOTAL. HD: 5 SEEDS; CV: 10 SEEDS. STATISTICALLY SIGNIFICANT DIFFERENCES ($p < 0.05$, WILCOXON) MARKED WITH *.

Algorithm	Condition	Heart Disease (920 samples, 13 feat.)				Cardiovascular (70K samples, 11 feat.)			
		Features	Acc. (%)	Δ (pp)	AUC	Features	Acc. (%)	Δ (pp)	AUC
FedAvg	no_governance	13/13	59.9 ± 7.7	—	.811	11/11	69.4 ± 3.6	—	.789
	gov_noDP	7/13	69.5 ± 8.7	$+9.6^\dagger$	.832	4/11	69.5 ± 2.6	$+0.1$	.785
	gov_ε=10	7/13	55.8 ± 16.7	-4.1	.551	4/11	65.6 ± 5.4	-3.8^*	.777
	gov_ε=1	7/13	57.0 ± 13.8	-2.9	.476	4/11	53.4 ± 5.3	-16.0^*	.493
Ditto	no_governance	13/13	74.3 ± 3.1	—	.798	11/11	83.6 ± 5.1	—	.911
	gov_noDP	7/13	73.2 ± 5.4	-1.1	.817	4/11	83.2 ± 5.3	-0.3	.908
	gov_ε=10	7/13	74.0 ± 5.9	$-0.3^\dagger$	.815	4/11	82.9 ± 5.4	-0.7^*	.906
	gov_ε=1	7/13	79.5 ± 1.6	$+5.2^\dagger$	.842	4/11	70.5 ± 11.9	-13.1^*	.823
HPFL	no_governance	13/13	73.3 ± 3.3	—	.793	11/11	83.6 ± 5.1	—	.911
	gov_noDP	7/13	71.5 ± 4.8	-1.8	.808	4/11	83.1 ± 5.3	-0.4	.908
	gov_ε=10	7/13	77.1 ± 5.6	$+3.8^\dagger$	.851	4/11	81.9 ± 5.9	-1.7^*	.904
	gov_ε=1	7/13	68.1 ± 6.1	-5.2	.794	4/11	73.2 ± 11.5	-10.3^*	.825

[†]HD paradoxes: positive Δ under governance indicates small-sample regularization artifact (920 samples); feature reduction acts as implicit regularization. CV (70K) shows expected monotonically negative Δ , confirming true governance cost. *Significant at $p < 0.05$.

TABLE S-XCV

H2: NON-UNIFORM ARTICLE 71 OPT-OUT IMPACT ON FAIRNESS. FEDAVG, 10 SEEDS PER PATTERN. JAIN INDEX = 1.0 INDICATES PERFECT EQUITY; DEI RATIO = MIN/MAX CLIENT ACCURACY. p -VALUES FROM WILCOXON RANK-SUM VS. UNIFORM_0 BASELINE.

Dataset	Pattern	Acc. (%)	Jain	DEI	Range (pp)	p
HD (4 hosp.)	uniform_0 [0,0,0,0]	61.5 ± 7.9	.922	.470	47.5	—
	uniform_15 [.15,.15,.15,.15]	60.8 ± 6.4	.917	.455	48.9	.84
	half_30 [0,0,.3,.3]	69.5 ± 10.6	.944	.616	29.1	.14
	single_50 [0,0,0,.5]	68.5 ± 10.7	.943	.591	32.6	.23
	gradient [0,.1,.2,.3]	66.8 ± 10.3	.940	.597	33.8	.30
CV (5 cl.)	uniform_0 [0,0,0,0,0]	69.4 ± 3.6	.971	.693	27.0	—
	uniform_13 [.13,.13,.13,.13,.13]	69.3 ± 4.1	.966	.699	26.6	.90
	half_30 [0,0,0,.3,.3]	69.7 ± 4.9	.959	.742	22.8	.81
	single_50 [0,0,0,0,.5]	69.2 ± 3.9	.963	.676	28.5	.73
	gradient [0,.05,.1,.2,.3]	69.9 ± 4.3	.966	.737	23.1	.76

No pattern produces a statistically significant fairness degradation (all $p > 0.14$ on HD, $p > 0.73$ on CV). On CV, all Jain indices remain ≥ 0.96 , confirming that Article 71 opt-out at realistic rates is fairness-neutral.

cantly harm global accuracy or cross-hospital equity on adequately-sized datasets.

D. Hypothesis H3: Algorithm Governance Resilience

Hypothesis: Personalized FL algorithms (Ditto, HPFL) maintain higher accuracy under progressively restrictive EHDS purpose-based data minimization than global FL (FedAvg), as measured by governance resilience $R = \text{acc}(\text{purpose})/\text{acc}(\text{scientific_research}) \times 100\%$.

Design: 3 algorithms $\times$ 4 EHDS purposes $\times$ 10 seeds on each dataset. The purposes follow Article 53(1) of the EHDS Regulation: scientific research, public health surveil-

lance, patient safety, and official statistics. Each purpose triggers different feature retention via the DataMinimizer. Total: 240 experiments.

Key Findings (H3):

- 1) **Personalized algorithms are dramatically more governance-resilient.** On CV, Ditto/HPFL maintain 94.7–94.8% resilience at official_statistics (2/11 features, 81.8% reduction), while FedAvg drops to 84.7%—a 10pp gap in governance tolerance.
- 2) **Feature importance hierarchy is consistent.** Both datasets show: scientific research $>$ public health $>$ patient safety $>$ official statistics, matching the EHDS

TABLE S-XCVI

H3: ALGORITHM GOVERNANCE RESILIENCE ACROSS EHDS PURPOSES. FEATURES RETAINED AFTER ARTICLE 44 MINIMIZATION (PURPOSE-BASED + MI IMPORTANCE THRESHOLD 0.01). RESILIENCE R = ACCURACY RELATIVE TO SCIENTIFIC_RESEARCH BASELINE (%). 10 SEEDS PER CONDITION. BEST RESILIENCE PER PURPOSE IN **BOLD**.

Algorithm	EHDS Purpose	Heart Disease (920 samples)				Cardiovascular (70K samples)			
		Feat.	Acc. (%)	R (%)	AUC	Feat.	Acc. (%)	R (%)	AUC
Ditto	Scientific research	10/13	75.2 $\pm$ 3.5	100.0	.807	5/11	83.6 $\pm$ 5.2	100.0	.911
	Public health surv.	7/13	74.9 $\pm$ 3.1	99.6	.819	4/11	83.2 $\pm$ 5.3	99.5	.908
	Patient safety	10/13	73.9 $\pm$ 4.1	98.3	.808	3/11	83.0 $\pm$ 5.3	99.3	.903
	Official statistics	2/13	67.2 $\pm$ 3.0	89.4	.764	2/11	79.2 $\pm$ 7.7	94.7	.865
HPFL	Scientific research	10/13	74.6 $\pm$ 3.5	100.0	.805	5/11	83.6 $\pm$ 5.2	100.0	.911
	Public health surv.	7/13	73.3 $\pm$ 3.2	98.3	.812	4/11	83.1 $\pm$ 5.3	99.4	.908
	Patient safety	10/13	73.3 $\pm$ 4.3	98.3	.806	3/11	83.0 $\pm$ 5.3	99.3	.903
	Official statistics	2/13	68.0 $\pm$ 3.4	91.2	.763	2/11	79.3 $\pm$ 7.6	94.8	.865
FedAvg	Scientific research	10/13	61.3 $\pm$ 7.8	100.0	.834	5/11	70.8 $\pm$ 2.2	100.0	.793
	Public health surv.	7/13	66.8 $\pm$ 8.2	109.0 [†]	.842	4/11	69.7 $\pm$ 2.8	98.4	.786
	Patient safety	10/13	60.4 $\pm$ 6.9	98.5	.829	3/11	70.2 $\pm$ 4.0	99.2	.765
	Official statistics	2/13	65.0 $\pm$ 3.6	105.9 [†]	.708	2/11	60.0 $\pm$ 3.2	84.7	.669

[†]HD paradoxes: FedAvg resilience >100% indicates that feature reduction acts as regularization on the 920-sample dataset. On CV (70K), FedAvg correctly shows the lowest resilience (84.7% at official_statistics), confirming that personalized algorithms are genuinely more governance-resilient.

Article 53 data minimization intent (more restrictive purposes retain fewer features).

- 3) **CV confirms no paradoxes.** On CV, all resilience values are $\leq 100\%$, all governance costs are non-negative, and the algorithm ranking (Ditto $\approx$ HPFL $\gg$ FedAvg) is consistent—validating the governance resilience claim on a statistically adequate dataset.
- 4) **Mutual information filtering is effective.** On CV, only 5/11 features have $MI \geq 0.01$ for cardiovascular prediction. The 6 removed features (gender, height, glucose, smoking, alcohol, physical activity) contribute negligible predictive power in this dataset’s representation, enabling aggressive minimization without proportional accuracy loss.

E. Cross-Dataset Comparison and Small-Sample Analysis

The cross-dataset design reveals a systematic pattern: HD (920 samples) produces paradoxical results that CV (70K samples) resolves. Table S-XCVII summarizes the divergences.

Implication for reviewers: The cross-dataset validation provides three forms of evidence that personalized FL algorithms are governance-resilient:

- 1) **CV confirms the thesis directly** with statistically significant results ($p < 0.01$) and no paradoxes.
- 2) **HD paradoxes have a known explanation** (small-sample regularization) that is independently verifiable.
- 3) **The algorithm ranking is consistent across datasets:** Ditto $\approx$ HPFL $\gg$ FedAvg for governance resilience, with the gap widening under more restrictive purposes.

F. Governance Validation Summary

Table S-XCVIII synthesizes the 720 governance experiments into actionable deployment guidance.

TABLE S-XCVII
CROSS-DATASET DIVERGENCES: HEART DISEASE VS. CARDIOVASCULAR. “EXPECTED” INDICATES MONOTONICALLY NEGATIVE GOVERNANCE COST; “PARADOX” INDICATES POSITIVE COST (GOVERNANCE IMPROVES ACCURACY).

Finding	HD	CV	Explanation
FedAvg gov_noDP Δ	+9.6pp	+0.1pp	Regularization
Ditto gov_ε=1 Δ	+5.2pp	−13.1pp	DP as regularizer
HPFL gov_ε=10 Δ	+3.8pp	−1.7pp	Feat. reduction reg.
FedAvg H3 resilience	105.9%	84.7%	Spurious improvement
Ditto H3 resilience	89.4%	94.7%	True resilience
H2 Jain improvement	Yes	No	Sample rebalancing

All HD paradoxes follow a single mechanism: on 920 samples with 4 naturally heterogeneous hospitals (Cleveland 303, Hungarian 294, Switzerland 123, VA Long Beach 200), any form of data reduction—feature minimization, opt-out, DP noise—acts as implicit regularization, preventing overfitting to site-specific patterns. On CV (70K, Dirichlet $\alpha=0.5$), sufficient sample size eliminates this artifact.

TABLE S-XCVIII
GOVERNANCE VALIDATION: SUMMARY OF 720 EXPERIMENTS

Claim	Evidence (CV, 70K)
Full compliance cost <2pp	Ditto: −0.7pp, HPFL: −1.7pp at $\epsilon=10$ ($p < 0.01$)
$\epsilon=1$ is destructive	All algorithms lose >10pp; FedAvg collapses to 53%
Minimization alone is free	11→4 features costs <0.5pp (MI-based selection)
Opt-out is fairness-neutral	Jain ≥ 0.96 across all patterns ($p > 0.73$)
Personalization = resilience	Ditto/HPFL: 94.8% vs. FedAvg: 84.7% at 2 features
HD paradoxes are artifacts	All resolve on CV; explained by regularization

These results support the main paper’s thesis that *personalization architecture—not aggregation algorithm—is the critical EHDS deployment decision*. The governance layer imposes minimal cost when paired with personalized algorithms, enabling full Article 44 (data minimization), Article 71 (opt-out), and Recital 57 (differential privacy) compliance simultaneously.

REFERENCES

- [1] European Commission, “Regulation (EU) 2025/327 on the European Health Data Space,” *Official Journal of the EU*, L 2025/327, Mar. 2025.
- [2] C. Staunton *et al.*, “Ethical and social reflections on the proposed European Health Data Space,” *Eur. J. Human Genetics*, vol. 32, no. 5, pp. 498–505, 2024.
- [3] P. Quinn, E. Ellyne, and C. Yao, “Will the GDPR restrain health data access bodies under the EHDS?” *Computer Law & Security Review*, vol. 54, art. 105993, 2024.
- [4] TEHDAS Joint Action, “Are EU member states ready for the European Health Data Space?” *Eur. J. Public Health*, vol. 34, no. 6, pp. 1102–1108, 2024.
- [5] H. Fröhlich *et al.*, “Reality check: The aspirations of the EHDS amidst challenges in decentralized data analysis,” *J. Med. Internet Res.*, vol. 27, art. e76491, 2025.
- [6] S. van Drumpt *et al.*, “Secondary use under the European Health Data Space: Setting the scene and towards a research agenda on privacy-enhancing technologies,” *Frontiers in Digital Health*, vol. 7, art. 1602101, 2025.
- [7] R. Hussein *et al.*, “Interoperability framework of the EHDS for secondary use,” *J. Med. Internet Res.*, vol. 27, art. e69813, 2025.
- [8] R. Forster *et al.*, “User journeys in cross-European secondary use of health data,” *Eur. J. Public Health*, vol. 35, Suppl. 3, pp. iii18–iii24, 2025.
- [9] L. Svingel *et al.*, “Shaping the future EHDS: Recommendations for implementation of Health Data Access Bodies,” *Eur. J. Public Health*, vol. 35, Suppl. 3, pp. iii32–iii38, 2025.
- [10] C. Christiansen *et al.*, “Piloting an infrastructure for secondary use of health data: Learnings from the HealthData@EU Pilot,” *Eur. J. Public Health*, vol. 35, Suppl. 3, pp. iii3–iii4, 2025.
- [11] M. Shabani and P. Borry, “The European Health Data Space: Challenges and opportunities for health data governance,” *Eur. J. Human Genetics*, vol. 32, no. 8, pp. 891–897, 2024.
- [12] A. Ganna, E. Ingelsson, and D. Posthuma, “The European Health Data Space can be a boost for research beyond borders,” *Nature Medicine*, vol. 30, pp. 3053–3056, 2024.
- [13] B. McMahan *et al.*, “Communication-efficient learning of deep networks from decentralized data,” in *Proc. AISTATS*, pp. 1273–1282, 2017.
- [14] T. Li *et al.*, “Federated optimization in heterogeneous networks,” in *Proc. MLSys*, vol. 2, pp. 429–450, 2020.
- [15] P. Kairouz *et al.*, “Advances and open problems in federated learning,” *Found. Trends Mach. Learn.*, vol. 14, no. 1–2, pp. 1–210, 2021.
- [16] N. Rieke *et al.*, “The future of digital health with federated learning,” *npj Digital Medicine*, vol. 3, art. 119, 2020.
- [17] K. Bonawitz *et al.*, “Towards federated learning at scale: A system design,” in *Proc. MLSys*, pp. 374–388, 2019.
- [18] M. Chavero-Diez *et al.*, “Federated learning frameworks: Quality and interoperability for biomedical research,” *NAR Genomics Bioinformatics*, vol. 8, no. 1, art. lqag010, 2026.
- [19] Z. L. Teo *et al.*, “Federated machine learning in healthcare: A systematic review,” *Cell Reports Medicine*, vol. 5, no. 2, art. 101419, 2024.
- [20] L. Peng *et al.*, “Federated machine learning in healthcare: A systematic review on clinical applications and technical architecture,” *Comput. Methods Programs Biomed.*, vol. 247, art. 108066, 2024.
- [21] L. Zhu, Z. Liu, and S. Han, “Deep leakage from gradients,” in *Proc. NeurIPS*, vol. 32, pp. 14774–14784, 2019.
- [22] R. Shokri *et al.*, “Membership inference attacks against machine learning models,” in *Proc. IEEE S&P*, pp. 3–18, 2017.
- [23] N. Carlini *et al.*, “Membership inference attacks from first principles,” in *Proc. IEEE S&P*, pp. 1897–1914, 2022.
- [24] C. Dwork and A. Roth, “The algorithmic foundations of differential privacy,” *Found. Trends Theor. Comput. Sci.*, vol. 9, no. 3–4, pp. 211–407, 2014.
- [25] M. Abadi *et al.*, “Deep learning with differential privacy,” in *Proc. ACM CCS*, pp. 308–318, 2016.
- [26] I. Mironov, “Rényi differential privacy,” in *Proc. IEEE CSF*, pp. 263–275, 2017.
- [27] I. Dayan *et al.*, “Federated learning for predicting clinical outcomes in patients with COVID-19,” *Nature Medicine*, vol. 27, no. 10, pp. 1735–1743, 2021.
- [28] M. J. Sheller *et al.*, “Federated learning in medicine: Facilitating multi-institutional collaborations without sharing patient data,” *Scientific Reports*, vol. 10, art. 12598, 2020.
- [29] S. P. Karimireddy *et al.*, “SCAFFOLD: Stochastic controlled averaging for federated learning,” in *Proc. ICML*, pp. 5132–5143, 2020.
- [30] J. Wang *et al.*, “Tackling the objective inconsistency problem in heterogeneous federated optimization,” in *Proc. NeurIPS*, vol. 33, pp. 7611–7623, 2020.
- [31] S. Reddi *et al.*, “Adaptive federated optimization,” in *Proc. ICLR*, 2021.
- [32] D. A. E. Acar, Y. Zhao, R. M. Navarro, M. Mattina, P. N. Whatmough, and V. Saligrama, “Federated learning based on dynamic regularization,” in *Proc. ICLR*, 2021.
- [33] T. Li, S. Hu, A. Beirami, and V. Smith, “Ditto: Fair and robust federated learning through personalization,” in *Proc. ICML*, PMLR 139, pp. 6357–6368, 2021.
- [34] A. Fallah, A. Mokhtari, and A. Ozdaglar, “Personalized federated learning with theoretical guarantees: A model-agnostic meta-learning approach,” in *Proc. NeurIPS*, vol. 33, pp. 3557–3568, 2020.
- [35] T. Dinh, N. Tran, and J. Nguyen, “Personalized federated learning with Moreau envelopes,” in *Proc. NeurIPS*, vol. 33, pp. 21394–21405, 2020.
- [36] K. Wei *et al.*, “Federated learning with differential privacy: Algorithms and performance analysis,” *IEEE Trans. Inf. Forensics Secur.*, vol. 15, pp. 3454–3469, 2020.
- [37] J. Jordon *et al.*, “Synthetic data—A privacy mirage?” *J. Mach. Learn. Res.*, vol. 23, no. 1, art. 298, 2022.
- [38] Z. Qu *et al.*, “Generalized federated learning via sharpness aware minimization,” in *Proc. ICML*, PMLR 162, pp. 18250–18280, 2022.
- [39] J. Zhang *et al.*, “Federated learning with label distribution skew via logits calibration,” in *Proc. ICML*, PMLR 162, pp. 26311–26329, 2022.
- [40] Y. Shi *et al.*, “Towards understanding and mitigating dimensional collapse in heterogeneous federated learning,” in *Proc. ICLR*, 2023.
- [41] Y. Sun *et al.*, “FedSpeed: Larger local interval, less communication round, and higher generalization accuracy,” in *Proc. ICLR*, 2023.
- [42] D. Jhunjhunwala, S. Wang, and G. Joshi, “FedExp: Speeding up federated averaging via extrapolation,” in *Proc. ICLR*, 2023.
- [43] Z. Qu *et al.*, “FedLESAM: Federated learning with locally estimated sharpness-aware minimization,” in *Proc. ICML*, PMLR 235, 2024. (Spotlight)
- [44] Y. Chen, X. Cao, and L. Sun, “HPFL: Hot-pluggable federated learning with shared backbone and personalized classifiers,” in *Proc. ICLR*, 2025.
- [45] D. J. Beutel *et al.*, “Flower: A friendly federated learning research framework,” *arXiv:2007.14390*, 2023.
- [46] NVIDIA, “NVIDIA FLARE: An open-source federated learning platform,” *GitHub Repository*, 2023.
- [47] Google, “TensorFlow Federated: Machine learning on decentralized data,” 2019.
- [48] X. Li *et al.*, “FedBN: Federated learning on non-IID features via local batch normalization,” in *Proc. ICLR*, 2021.
- [49] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proc. IEEE CVPR*, pp. 770–778, 2016.
- [50] D. S. Kermay *et al.*, “Identifying medical diagnoses and treatable diseases by image-based deep learning,” *Cell*, vol. 172, no. 5, pp. 1122–1131, 2018.
- [51] P. Wagner *et al.*, “PTB-XL, a large publicly available electrocardiography dataset,” *Scientific Data*, vol. 7, art. 154, 2020.
- [52] B. Strack *et al.*, “Impact of HbA1c measurement on hospital readmission rates: Analysis of 70,000 clinical database patient records,” *BioMed Research International*, vol. 2014, art. 781670, 2014.
- [53] T. Li *et al.*, “Fair resource allocation in federated learning,” in *Proc. ICLR*, 2020.
- [54] M. Mohri, G. Sivek, and A. T. Suresh, “Agnostic federated learning,” in *Proc. ICML*, PMLR 97, pp. 4615–4625, 2019.
- [55] Z. Hu, K. Shaloudegi, G. Zhang, and Y. Yu, “Federated learning meets multi-objective optimization,” *IEEE Trans. Network Science and Engineering*, vol. 9, no. 4, pp. 2039–2051, 2022.

- [56] J. Kirkpatrick *et al.*, “Overcoming catastrophic forgetting in neural networks,” *Proc. National Academy of Sciences*, vol. 114, no. 13, pp. 3521–3526, 2017.
- [57] Z. Li and D. Hoiem, “Learning without forgetting,” *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 40, no. 12, pp. 2935–2947, 2017.
- [58] G. Liu, X. Ma, Y. Yang, C. Wang, and J. Liu, “FedEraser: Enabling efficient client-level data removal from federated learning,” in *Proc. IEEE/ACM IWQoS*, pp. 1–10, 2021.
- [59] C. Wu, S. Zhu, and P. Mitra, “Federated unlearning with knowledge distillation,” *arXiv:2201.09441*, 2022.
- [60] F. Lai, X. Zhu, H. V. Madhyastha, and M. Chowdhury, “Oort: Efficient federated learning via guided participant selection,” in *Proc. OSDI*, pp. 19–35, 2021.
- [61] Y. J. Cho, J. Wang, and G. Joshi, “Towards understanding biased client selection in federated learning,” in *Proc. AISTATS*, PMLR 151, pp. 10351–10375, 2022.
- [62] J. Geiping, H. Bauermeister, H. Dröge, and M. Moeller, “Inverting gradients—How easy is it to break privacy in federated learning?” in *Proc. NeurIPS*, vol. 33, pp. 16937–16947, 2020.